

Package ‘brms’

September 9, 2025

Encoding UTF-8

Type Package

Title Bayesian Regression Models using 'Stan'

Version 2.23.0

Date 2025-09-08

Depends R (>= 3.6.0), Rcpp (>= 0.12.0), methods

Imports rstan (>= 2.29.0), ggplot2 (>= 2.0.0), loo (>= 2.8.0),
posterior (>= 1.6.0), Matrix (>= 1.1.1), mgcv (>= 1.8-13),
rstantools (>= 2.1.1), bayesplot (>= 1.5.0), bridgesampling (>= 0.3-0), glue (>= 1.3.0), rlang (>= 1.0.0), future (>= 1.19.0),
future.apply (>= 1.0.0), matrixStats, nleqslv, nlme, coda,
abind, stats, utils, parallel, grDevices, backports

Suggests testthat (>= 3.1.9), emmeans (>= 1.4.2), cmdstanr (>= 0.5.0),
projpred (>= 2.0.0), priorsense (>= 1.0.0), shinystan (>= 2.4.0), splines2 (>= 0.5.0), RWiener, rtdists, extraDistr, processx, mice, spdep, mnormt, lme4, MCMCglmm, ape, arm, statmod, digest, diffobj, betareg, R.rsp, gtable, shiny, knitr, rmarkdown, ragg, colorspace, mirai, future.mirai

Description Fit Bayesian generalized (non-)linear multivariate multilevel models using 'Stan' for full Bayesian inference. A wide range of distributions and link functions are supported, allowing users to fit -- among others -- linear, robust linear, count data, survival, response times, ordinal, zero-inflated, hurdle, and even self-defined mixture models all in a multilevel context. Further modeling options include both theory-driven and data-driven non-linear terms, auto-correlation structures, censoring and truncation, meta-analytic standard errors, and quite a few more.

In addition, all parameters of the response distribution can be predicted in order to perform distributional regression. Prior specifications are flexible and explicitly encourage users to apply prior distributions that actually reflect their prior knowledge. Models can easily be evaluated and compared using several methods assessing posterior or prior predictions.

References: Bürkner (2017) <[doi:10.18637/jss.v080.i01](https://doi.org/10.18637/jss.v080.i01)>;

Bürkner (2018) <[doi:10.32614/RJ-2018-017](https://doi.org/10.32614/RJ-2018-017)>;

Bürkner (2021) <[doi:10.18637/jss.v100.i05](https://doi.org/10.18637/jss.v100.i05)>; Carpenter et al. (2017)

<doi:10.18637/jss.v076.i01>.

LazyData true

NeedsCompilation no

License GPL-2

URL <https://github.com/paul-buerkner/brms>,
<https://discourse.mc-stan.org/>, <https://paulbuerkner.com/brms/>

BugReports <https://github.com/paul-buerkner/brms/issues>

Additional_repositories <https://stan-dev.r-universe.dev/>

SystemRequirements pngquant

VignetteBuilder knitr, R.rsp

RoxygenNote 7.3.2

Author Paul-Christian Bürkner [aut, cre],
Jonah Gabry [ctb],
Sebastian Weber [ctb],
Andrew Johnson [ctb],
Martin Modrak [ctb],
Hamada S. Badr [ctb],
Frank Weber [ctb],
Aki Vehtari [ctb],
Mattan S. Ben-Shachar [ctb],
Hayden Rabel [ctb],
Simon C. Mills [ctb],
Stephen Wild [ctb],
Ven Popov [ctb],
Ioannis Kosmidis [ctb],
Ben Schneider [ctb],
Noa Kallioinen [ctb],
Wellington J. Silva [ctb],
Luiz Carvalho [ctb],
Sermet Pekin [ctb]

Maintainer Paul-Christian Bürkner <paul.buerkner@gmail.com>

Repository CRAN

Date/Publication 2025-09-09 10:50:02 UTC

Contents

brms-package	6
addition-terms	8
add_criterion	11
add_loo	12
add_rstan_model	13
ar	13

arma	14
as.brmsprior	15
as.data.frame.brmsfit	16
as.mcmc.brmsfit	17
AsymLaplace	18
autocor-terms	19
autocor.brmsfit	20
bayes_factor.brmsfit	20
bayes_R2.brmsfit	21
BetaBinomial	23
bridge_sampler.brmsfit	24
brm	25
brmsfamily	34
brmsfit-class	40
brmsformula	41
brmsformula-helpers	51
brmshypothesis	53
brmsterms	54
brm_multiple	56
car	60
coef.brmsfit	61
combine_models	63
compare_ic	63
conditional_effects.brmsfit	64
conditional_smooths.brmsfit	69
constant	72
control_params	73
cor_ar	73
cor_arma	74
cor_brms	75
cor_car	76
cor_cosy	77
cor_fixed	78
cor_ma	79
cor_sar	80
cosy	81
create_priorsense_data.brmsfit	82
cs	83
custom_family	83
default_prior	86
default_prior.default	87
density_ratio	89
diagnostic-quantities	90
Dirichlet	91
draws-brms	92
draws-index-brms	93
emmeans-brms-helpers	94
epilepsy	96

ExGaussian	97
expose_functions.brmsfit	98
expp1	98
family.brmsfit	99
fcor	99
fitted.brmsfit	100
fixef.brmsfit	102
Frechet	103
GenExtremeValue	104
get_dpar	105
get_refmodel.brmsfit	106
gp	108
gr	110
horseshoe	112
Hurdle	114
hypothesis.brmsfit	115
inhaler	118
inits.brmsfit	119
InvGaussian	120
inv_logit_scaled	121
is.brmsfit	121
is.brmsfit_multiple	122
is.brmsformula	122
is.brmsprior	122
is.brmsterms	123
is.cor_brms	123
is.mvbrmsformula	124
is.mvbrmsterms	124
kfold.brmsfit	125
kfold_predict	128
kidney	130
lasso	131
launch_shinystan.brmsfit	132
LogisticNormal	133
logit_scaled	133
logm1	134
log_lik.brmsfit	134
loo.brmsfit	136
loo_compare.brmsfit	138
loo_model_weights.brmsfit	139
loo_moment_match.brmsfit	140
loo_predict.brmsfit	142
loo_R2.brmsfit	144
loo_subsample.brmsfit	145
loss	146
ma	147
make_conditions	148
mcmc_plot.brmsfit	149

me	151
mi	152
mixture	153
mm	155
mmc	157
mo	158
model_weights.brmsfit	159
MultiNormal	160
MultiStudentT	161
mvbind	162
mvbrmsformula	162
ngrps.brmsfit	163
nsamples.brmsfit	164
opencl	164
pairs.brmsfit	165
parnames	166
plot.brmsfit	166
posterior_average.brmsfit	168
posterior_epred.brmsfit	170
posterior_interval.brmsfit	172
posterior_linpred.brmsfit	173
posterior_predict.brmsfit	174
posterior_samples.brmsfit	177
posterior_smooths.brmsfit	178
posterior_summary	180
posterior_table	181
post_prob.brmsfit	182
pp_average.brmsfit	183
pp_check.brmsfit	185
pp_mixture.brmsfit	187
predict.brmsfit	189
predictive_error.brmsfit	192
predictive_interval.brmsfit	193
prepare_predictions.brmsfit	194
print.brmsfit	196
print.brmsprior	197
prior_draws.brmsfit	197
prior_summary.brmsfit	198
psis.brmsfit	199
R2D2	201
ranef.brmsfit	202
read_csv_as_stanfit	203
recompile_model	205
reloo.brmsfit	205
rename_pars	207
residuals.brmsfit	208
restructure	210
restructure.brmsfit	211

rows2labels	212
s	212
sar	213
save_pars	215
set_prior	216
Shifted_Lognormal	222
SkewNormal	223
stancode	224
stancode.brmsfit	225
stancode.default	226
standata	228
standata.brmsfit	229
standata.default	230
stanvar	232
StudentT	234
summary.brmsfit	235
theme_black	236
theme_default	237
threading	237
unstr	238
update.brmsfit	239
update.brmsfit_multiple	240
update_adterms	241
validate_newdata	242
validate_prior	243
VarCorr.brmsfit	245
vcov.brmsfit	246
VonMises	247
waic.brmsfit	248
Wiener	249
ZeroInflated	251

Index	253
--------------	------------

Description

The **brms** package provides an interface to fit Bayesian generalized multivariate (non-)linear multilevel models using **Stan**, which is a C++ package for obtaining full Bayesian inference (see <https://mc-stan.org/>). The formula syntax is an extended version of the syntax applied in the **lme4** package to provide a familiar and simple interface for performing regression analyses.

Details

The main function of **brms** is `brm`, which uses formula syntax to specify a wide range of complex Bayesian models (see `brmsformula` for details). Based on the supplied formulas, data, and additional information, it writes the Stan code on the fly via `stancode`, prepares the data via `standata` and fits the model using `Stan`.

Subsequently, a large number of post-processing methods can be applied: To get an overview on the estimated parameters, `summary` or `conditional_effects` are perfectly suited. Detailed visual analyses can be performed by applying the `pp_check` and `stanplot` methods, which both rely on the `bayesplot` package. Model comparisons can be done via `loo` and `waic`, which make use of the `loo` package as well as via `bayes_factor` which relies on the `bridgesampling` package. For a full list of methods to apply, type `methods(class = "brmsfit")`.

Because **brms** is based on `Stan`, a C++ compiler is required. The program Rtools (available on <https://cran.r-project.org/bin/windows/Rtools/>) comes with a C++ compiler for Windows. On Mac, you should use Xcode. For further instructions on how to get the compilers running, see the prerequisites section at the [RStan-Getting-Started](#) page.

When comparing other packages fitting multilevel models to **brms**, keep in mind that the latter needs to compile models before actually fitting them, which will require between 20 and 40 seconds depending on your machine, operating system and overall model complexity.

Thus, fitting smaller models may be relatively slow as compilation time makes up the majority of the whole running time. For larger / more complex models however, fitting may take several minutes or even hours, so that the compilation time won't make much of a difference for these models.

See `vignette("brms_overview")` and `vignette("brms_multilevel")` for a general introduction and overview of **brms**. For a full list of available vignettes, type `vignette(package = "brms")`.

Author(s)

Maintainer: Paul-Christian Bürkner <paul.buerkner@gmail.com>

Other contributors:

- Jonah Gabry [contributor]
- Sebastian Weber [contributor]
- Andrew Johnson [contributor]
- Martin Modrak [contributor]
- Hamada S. Badr [contributor]
- Frank Weber [contributor]
- Aki Vehtari [contributor]
- Mattan S. Ben-Shachar [contributor]
- Hayden Rabel [contributor]
- Simon C. Mills [contributor]
- Stephen Wild [contributor]
- Ven Popov [contributor]
- Ioannis Kosmidis [contributor]
- Ben Schneider [contributor]

- Noa Kallioinen [contributor]
- Wellington J. Silva [contributor]
- Luiz Carvalho [contributor]
- Sermet Pekin [contributor]

References

- Paul-Christian Buerkner (2017). brms: An R Package for Bayesian Multilevel Models Using Stan. *Journal of Statistical Software*, 80(1), 1-28. doi:10.18637/jss.v080.i01
- Paul-Christian Buerkner (2018). Advanced Bayesian Multilevel Modeling with the R Package brms. *The R Journal*. 10(1), 395–411. doi:10.32614/RJ-2018-017
- The Stan Development Team. *Stan Modeling Language User's Guide and Reference Manual*. <https://mc-stan.org/users/documentation/>.
- Stan Development Team (2020). RStan: the R interface to Stan. R package version 2.21.2. <https://mc-stan.org/>

See Also

[brm](#), [brmsformula](#), [brmsfamily](#), [brmsfit](#)

addition-terms

Additional Response Information

Description

Provide additional information on the response variable in **brms** models, such as censoring, truncation, or known measurement error. Detailed documentation on the use of each of these functions can be found in the Details section of [brmsformula](#) (under "Additional response information").

Usage

```
resp_se(x, sigma = FALSE)

resp_weights(x, scale = FALSE)

resp_trials(x)

resp_thres(x, gr = NA)

resp_cat(x)

resp_dec(x)

resp_bhaz(gr = NA, df = 5, ...)
```

```

resp_cens(x, y2 = NA)

resp_trunc(lb = -Inf, ub = Inf)

resp_mi(sdy = NA)

resp_index(x)

resp_rate(denom)

resp_subset(x)

resp_vreal(...)

resp_vint(...)

```

Arguments

x	A vector; Ideally a single variable defined in the data (see Details). Allowed values depend on the function: <code>resp_se</code> and <code>resp_weights</code> require positive numeric values. <code>resp_trials</code> , <code>resp_thres</code> , and <code>resp_cat</code> require positive integers. <code>resp_dec</code> requires 0 and 1, or alternatively 'lower' and 'upper'. <code>resp_subset</code> requires 0 and 1, or alternatively FALSE and TRUE. <code>resp_cens</code> requires 'left', 'none', 'right', and 'interval' (or equivalently -1, 0, 1, and 2) to indicate left, no, right, or interval censoring. <code>resp_index</code> does not make any requirements other than the value being unique for each observation.
sigma	Logical; Indicates whether the residual standard deviation parameter <code>sigma</code> should be included in addition to the known measurement error. Defaults to FALSE for backwards compatibility, but setting it to TRUE is usually the better choice.
scale	Logical; Indicates whether weights should be scaled so that the average weight equals one. Defaults to FALSE.
gr	A vector of grouping indicators.
df	Degrees of freedom of baseline hazard splines for Cox models.
...	For <code>resp_vreal</code> , vectors of real values. For <code>resp_vint</code> , vectors of integer values. In Stan, these variables will be named <code>vreal1</code> , <code>vreal2</code> , ..., and <code>vint1</code> , <code>vint2</code> , ..., respectively.
y2	A vector specifying the upper bounds in interval censoring. Will be ignored for non-interval censored observations. However, it should NOT be NA even for non-interval censored observations to avoid accidental exclusion of these observations.
lb	A numeric vector or single numeric value specifying the lower truncation bound.
ub	A numeric vector or single numeric value specifying the upper truncation bound.
sdy	Optional known measurement error of the response treated as standard deviation. If specified, handles measurement error and (completely) missing values at the same time using the plausible-values-technique.
denom	A vector of positive numeric values specifying the denominator values from which the response rates are computed.

Details

These functions are almost solely useful when called in formulas passed to the **brms** package. Within formulas, the `resp_` prefix may be omitted. More information is given in the 'Details' section of **brmsformula** (under "Additional response information").

It is highly recommended to use a single data variable as input for `x` (instead of a more complicated expression) to make sure all post-processing functions work as expected.

Value

A list of additional response information to be processed further by **brms**.

See Also

[brm](#), [brmsformula](#)

Examples

```
## Not run:
## Random effects meta-analysis
nstudies <- 20
true_effects <- rnorm(nstudies, 0.5, 0.2)
sei <- runif(nstudies, 0.05, 0.3)
outcomes <- rnorm(nstudies, true_effects, sei)
data1 <- data.frame(outcomes, sei)
fit1 <- brm(outcomes | se(sei, sigma = TRUE) ~ 1,
            data = data1)
summary(fit1)

## Probit regression using the binomial family
n <- sample(1:10, 100, TRUE) # number of trials
success <- rbinom(100, size = n, prob = 0.4)
x <- rnorm(100)
data2 <- data.frame(n, success, x)
fit2 <- brm(success | trials(n) ~ x, data = data2,
            family = binomial("probit"))
summary(fit2)

## Survival regression modeling the time between the first
## and second recurrence of an infection in kidney patients.
fit3 <- brm(time | cens(censored) ~ age * sex + disease + (1|patient),
            data = kidney, family = lognormal())
summary(fit3)

## Poisson model with truncated counts
fit4 <- brm(count | trunc(ub = 104) ~ zBase * Trt,
            data = epilepsy, family = poisson())
summary(fit4)

## End(Not run)
```

add_criterion	<i>Add model fit criteria to model objects</i>
---------------	--

Description

Add model fit criteria to model objects

Usage

```
add_criterion(x, ...)

## S3 method for class 'brmsfit'
add_criterion(
  x,
  criterion,
  model_name = NULL,
  overwrite = FALSE,
  file = NULL,
  force_save = FALSE,
  ...
)
```

Arguments

x	An R object typically of class <code>brmsfit</code> .
...	Further arguments passed to the underlying functions computing the model fit criteria. If you are recomputing an already stored criterion with other ... arguments, make sure to set <code>overwrite = TRUE</code> .
criterion	Names of model fit criteria to compute. Currently supported are "loo", "waic", "kfold", "loo_subsample", "bayes_R2" (Bayesian R-squared), "loo_R2" (LOO-adjusted R-squared), and "marglik" (log marginal likelihood).
model_name	Optional name of the model. If <code>NULL</code> (the default) the name is taken from the call to <code>x</code> .
overwrite	Logical; Indicates if already stored fit indices should be overwritten. Defaults to <code>FALSE</code> . Setting it to <code>TRUE</code> is useful for example when changing additional arguments of an already stored criterion.
file	Either <code>NULL</code> or a character string. In the latter case, the fitted model object including the newly added criterion values is saved via <code>saveRDS</code> in a file named after the string supplied in <code>file</code> . The <code>.rds</code> extension is added automatically. If <code>x</code> was already stored in a file before, the file name will be reused automatically (with a message) unless overwritten by <code>file</code> . In any case, <code>file</code> only applies if new criteria were actually added via <code>add_criterion</code> or if <code>force_save</code> was set to <code>TRUE</code> .
force_save	Logical; only relevant if <code>file</code> is specified and ignored otherwise. If <code>TRUE</code> , the fitted model object will be saved regardless of whether new criteria were added via <code>add_criterion</code> .

Details

Functions `add_loo` and `add_waic` are aliases of `add_criterion` with fixed values for the `criterion` argument.

Value

An object of the same class as `x`, but with model fit criteria added for later usage.

Examples

```
## Not run:
fit <- brm(count ~ Trt, data = epilepsy)
# add both LOO and WAIC at once
fit <- add_criterion(fit, c("loo", "waic"))
print(fit$criteria$loo)
print(fit$criteria$waic)

## End(Not run)
```

`add_loo`

Add model fit criteria to model objects

Description

Deprecated aliases of [add_criterion](#).

Usage

```
add_loo(x, model_name = NULL, ...)
add_waic(x, model_name = NULL, ...)
add_ic(x, ...)

## S3 method for class 'brmsfit'
add_ic(x, ic = "loo", model_name = NULL, ...)
add_ic(x, ...) <- value
```

Arguments

- `x` An R object typically of class `brmsfit`.
- `model_name` Optional name of the model. If `NULL` (the default) the name is taken from the call to `x`.
- `...` Further arguments passed to the underlying functions computing the model fit criteria. If you are recomputing an already stored criterion with other `...` arguments, make sure to set `overwrite = TRUE`.

<code>ic, value</code>	Names of model fit criteria to compute. Currently supported are "loo", "waic", "kfold", "R2" (R-squared), and "marglik" (log marginal likelihood).
------------------------	--

Value

An object of the same class as `x`, but with model fit criteria added for later usage. Previously computed criterion objects will be overwritten.

<code>add_rstan_model</code>	<i>Add compiled rstan models to brmsfit objects</i>
------------------------------	---

Description

Compile a `stanmodel` and add it to a `brmsfit` object. This enables some advanced functionality of `rstan`, most notably `log_prob` and friends, to be used with `brms` models fitted with other Stan backends.

Usage

```
add_rstan_model(x, overwrite = FALSE)
```

Arguments

<code>x</code>	A <code>brmsfit</code> object to be updated.
<code>overwrite</code>	Logical. If TRUE, overwrite any existing <code>stanmodel</code> . Defaults to FALSE.

Value

A (possibly updated) `brmsfit` object.

<code>ar</code>	<i>Set up AR(p) correlation structures</i>
-----------------	---

Description

Set up an autoregressive (AR) term of order p in `brms`. The function does not evaluate its arguments – it exists purely to help set up a model with AR terms.

Usage

```
ar(time = NA, gr = NA, p = 1, cov = FALSE)
```

Arguments

time	An optional time variable specifying the time ordering of the observations. By default, the existing order of the observations in the data is used.
gr	An optional grouping variable. If specified, the correlation structure is assumed to apply only to observations within the same grouping level.
p	A non-negative integer specifying the autoregressive (AR) order of the ARMA structure. Default is 1.
cov	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If FALSE (the default), a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Value

An object of class 'arma_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#), [arma](#), [ma](#)

Examples

```
## Not run:
data("LakeHuron")
LakeHuron <- as.data.frame(LakeHuron)
fit <- brm(x ~ ar(p = 2), data = LakeHuron)
summary(fit)

## End(Not run)
```

arma

Set up ARMA(p,q) correlation structures

Description

Set up an autoregressive moving average (ARMA) term of order (p, q) in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with ARMA terms.

Usage

```
arma(time = NA, gr = NA, p = 1, q = 1, cov = FALSE)
```

Arguments

time	An optional time variable specifying the time ordering of the observations. By default, the existing order of the observations in the data is used.
gr	An optional grouping variable. If specified, the correlation structure is assumed to apply only to observations within the same grouping level.
p	A non-negative integer specifying the autoregressive (AR) order of the ARMA structure. Default is 1.
q	A non-negative integer specifying the moving average (MA) order of the ARMA structure. Default is 1.
cov	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If FALSE (the default), a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Value

An object of class 'arma_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#), [ar](#), [ma](#),

Examples

```
## Not run:
data("LakeHuron")
LakeHuron <- as.data.frame(LakeHuron)
fit <- brm(x ~ arma(p = 2, q = 1), data = LakeHuron)
summary(fit)

## End(Not run)
```

as.brmsprior

Transform into a brmsprior object

Description

Try to transform an object into a **brmsprior** object.

Usage

`as.brmsprior(x)`

Arguments

- x An object to be transformed.

Value

A `brmsprior` object if the transformation was possible.

`as.data.frame.brmsfit` *Extract Posterior Draws*

Description

Extract posterior draws in conventional formats as `data.frames`, `matrices`, or `arrays`.

Usage

```
## S3 method for class 'brmsfit'
as.data.frame(
  x,
  row.names = NULL,
  optional = TRUE,
  pars = NA,
  variable = NULL,
  draw = NULL,
  subset = NULL,
  ...
)

## S3 method for class 'brmsfit'
as.matrix(x, pars = NA, variable = NULL, draw = NULL, subset = NULL, ...)

## S3 method for class 'brmsfit'
as.array(x, pars = NA, variable = NULL, draw = NULL, subset = NULL, ...)
```

Arguments

- x A `brmsfit` object or another R object for which the methods are defined.
- row.names, optional
Unused and only added for consistency with the `as.data.frame` generic.
- pars
Deprecated alias of `variable`. For reasons of backwards compatibility, `pars` is interpreted as a vector of regular expressions by default unless `fixed = TRUE` is specified.
- variable
A character vector providing the variables to extract. By default, all variables are extracted.
- draw
The draw indices to be select. Subsetting draw indices will lead to an automatic merging of chains.

subset	Deprecated alias of <code>draw</code> .
...	Further arguments to be passed to the corresponding <code>as_draws_*</code> methods as well as to <code>subset_draws</code> .

Value

A data.frame, matrix, or array containing the posterior draws.

See Also

`as_draws`, `subset_draws`

`as.mcmc.brmsfit`

(*Deprecated*) Extract posterior samples for use with the `coda` package

Description

The `as.mcmc` method is deprecated. We recommend using the more modern and consistent `as_draws_*` extractor functions of the `posterior` package instead.

Usage

```
## S3 method for class 'brmsfit'
as.mcmc(
  x,
  pars = NA,
  fixed = FALSE,
  combine_chains = FALSE,
  inc_warmup = FALSE,
  ...
)
```

Arguments

<code>x</code>	An R object typically of class <code>brmsfit</code>
<code>pars</code>	Names of parameters for which posterior samples should be returned, as given by a character vector or regular expressions. By default, all posterior samples of all parameters are extracted.
<code>fixed</code>	Indicates whether parameter names should be matched exactly (TRUE) or treated as regular expressions (FALSE). Default is FALSE.
<code>combine_chains</code>	Indicates whether chains should be combined.
<code>inc_warmup</code>	Indicates if the warmup samples should be included. Default is FALSE. Warmup samples are used to tune the parameters of the sampling algorithm and should not be analyzed.
...	currently unused

Value

If `combine_chains` = TRUE an `mcmc` object is returned. If `combine_chains` = FALSE an `mcmc.list` object is returned.

AsymLaplace

*The Asymmetric Laplace Distribution***Description**

Density, distribution function, quantile function and random generation for the asymmetric Laplace distribution with location `mu`, scale `sigma` and asymmetry parameter `quantile`.

Usage

```
dasym_laplace(x, mu = 0, sigma = 1, quantile = 0.5, log = FALSE)

pasym_laplace(
  q,
  mu = 0,
  sigma = 1,
  quantile = 0.5,
  lower.tail = TRUE,
  log.p = FALSE
)

qasym_laplace(
  p,
  mu = 0,
  sigma = 1,
  quantile = 0.5,
  lower.tail = TRUE,
  log.p = FALSE
)

rasym_laplace(n, mu = 0, sigma = 1, quantile = 0.5)
```

Arguments

<code>x, q</code>	Vector of quantiles.
<code>mu</code>	Vector of locations.
<code>sigma</code>	Vector of scales.
<code>quantile</code>	Asymmetry parameter corresponding to quantiles in quantile regression (hence the name).
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.

log.p	Logical; If TRUE, values are returned on the log scale.
p	Vector of probabilities.
n	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

Description

Specify autocorrelation terms in **brms** models. Currently supported terms are `arma`, `ar`, `ma`, `cosy`, `unstr`, `sar`, `car`, and `fcor`. Terms can be directly specified within the formula, or passed to the `autocor` argument of `brmsformula` in the form of a one-sided formula. For deprecated ways of specifying autocorrelation terms, see `cor_brms`.

Details

The autocor term functions are almost solely useful when called in formulas passed to the **brms** package. They do not evaluate its arguments – but exist purely to help set up a model with autocorrelation terms.

See Also

`brmsformula`, `acformula`, `arma`, `ar`, `ma`, `cosy`, `unstr`, `sar`, `car`, `fcor`

Examples

```
# specify autocor terms within the formula
y ~ x + arma(p = 1, q = 1) + car(M)

# specify autocor terms in the 'autocor' argument
bf(y ~ x, autocor = ~ arma(p = 1, q = 1) + car(M))

# specify autocor terms via 'acformula'
bf(y ~ x) + acformula(~ arma(p = 1, q = 1) + car(M))
```

`autocor.brmsfit` (*Deprecated*) Extract Autocorrelation Objects

Description

(Deprecated) Extract Autocorrelation Objects

Usage

```
## S3 method for class 'brmsfit'
autocor(object, resp = NULL, ...)

autocor(object, ...)
```

Arguments

- | | |
|---------------------|--|
| <code>object</code> | An object of class <code>brmsfit</code> . |
| <code>resp</code> | Optional names of response variables. If specified, predictions are performed only for the specified response variables. |
| <code>...</code> | Currently unused. |

Value

A `cor_brms` object or a list of such objects for multivariate models. Not supported for models fitted with `brms` 2.11.1 or higher.

`bayes_factor.brmsfit` *Bayes Factors from Marginal Likelihoods*

Description

Compute Bayes factors from marginal likelihoods.

Usage

```
## S3 method for class 'brmsfit'
bayes_factor(x1, x2, log = FALSE, ...)
```

Arguments

- | | |
|------------------|--|
| <code>x1</code> | A <code>brmsfit</code> object |
| <code>x2</code> | Another <code>brmsfit</code> object based on the same responses. |
| <code>log</code> | Report Bayes factors on the log-scale? |
| <code>...</code> | Additional arguments passed to <code>bridge_sampler</code> . |

Details

Computing the marginal likelihood requires samples of all variables defined in Stan's parameters block to be saved. Otherwise `bayes_factor` cannot be computed. Thus, please set `save_all_pars = TRUE` in the call to `brm`, if you are planning to apply `bayes_factor` to your models.

The computation of Bayes factors based on bridge sampling requires a lot more posterior samples than usual. A good conservative rule of thumb is perhaps 10-fold more samples (read: the default of 4000 samples may not be enough in many cases). If not enough posterior samples are provided, the bridge sampling algorithm tends to be unstable, leading to considerably different results each time it is run. We thus recommend running `bayes_factor` multiple times to check the stability of the results.

More details are provided under [bridgesampling::bayes_factor](#).

See Also

[bridge_sampler](#), [post_prob](#)

Examples

```
## Not run:
# model with the treatment effect
fit1 <- brm(
  count ~ zAge + zBase + Trt,
  data = epilepsy, family = negbinomial(),
  prior = prior(normal(0, 1), class = b),
  save_all_pars = TRUE
)
summary(fit1)

# model without the treatment effect
fit2 <- brm(
  count ~ zAge + zBase,
  data = epilepsy, family = negbinomial(),
  prior = prior(normal(0, 1), class = b),
  save_all_pars = TRUE
)
summary(fit2)

# compute the bayes factor
bayes_factor(fit1, fit2)

## End(Not run)
```

Description

Compute a Bayesian version of R-squared for regression models

Usage

```
## S3 method for class 'brmsfit'
bayes_R2(
  object,
  resp = NULL,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)
```

Arguments

object	An object of class <code>brmsfit</code> .
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
summary	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .
robust	If <code>FALSE</code> (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If <code>TRUE</code> , the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is <code>TRUE</code> .
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is <code>TRUE</code> .
...	Further arguments passed to <code>posterior_epred</code> , which is used in the computation of the R-squared values.

Details

For an introduction to the approach, see Gelman et al. (2019) and https://github.com/jgabry/bayes_R2/.

Value

If `summary = TRUE`, an $M \times C$ matrix is returned ($M =$ number of response variables and $C = \text{length(probs)} + 2$) containing summary statistics of the Bayesian R-squared values. If `summary = FALSE`, the posterior draws of the Bayesian R-squared values are returned in an $S \times M$ matrix (S is the number of draws).

References

Andrew Gelman, Ben Goodrich, Jonah Gabry & Aki Vehtari. (2019). R-squared for Bayesian regression models, *The American Statistician*, 73(3):307-309. 10.1080/00031305.2018.1549100 (Preprint available at https://acris.aalto.fi/ws/portalfiles/portal/34206843/bayes_R2_v3.pdf)

Examples

```
## Not run:
fit <- brm(mpg ~ wt + cyl, data = mtcars)
summary(fit)
bayes_R2(fit)

# compute R2 with new data
nd <- data.frame(mpg = c(10, 20, 30), wt = c(4, 3, 2), cyl = c(8, 6, 4))
bayes_R2(fit, newdata = nd)

## End(Not run)
```

BetaBinomial

The Beta-binomial Distribution

Description

Cumulative density & mass functions, and random number generation for the Beta-binomial distribution using the following re-parameterisation of the [Stan Beta-binomial definition](#):

- `mu` = $\alpha * \beta$ mean probability of trial success.
- `phi` = $(1 - \mu) * \beta$ precision or over-dispersion, component.

Usage

```
dbeta_binomial(x, size, mu, phi, log = FALSE)

pbeta_binomial(q, size, mu, phi, lower.tail = TRUE, log.p = FALSE)

rbeta_binomial(n, size, mu, phi)
```

Arguments

<code>x, q</code>	Vector of quantiles.
<code>size</code>	Vector of number of trials (zero or more).
<code>mu</code>	Vector of means.
<code>phi</code>	Vector of precisions.
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If TRUE, values are returned on the log scale.
<code>n</code>	Number of draws to sample from the distribution.

bridge_sampler.brmsfit*Log Marginal Likelihood via Bridge Sampling***Description**

Computes log marginal likelihood via bridge sampling, which can be used in the computation of bayes factors and posterior model probabilities. The `brmsfit` method is just a thin wrapper around the corresponding method for `stanfit` objects.

Usage

```
## S3 method for class 'brmsfit'
bridge_sampler(samples, recompile = FALSE, ...)
```

Arguments

- | | |
|------------------------|---|
| <code>samples</code> | A <code>brmsfit</code> object. |
| <code>recompile</code> | Logical, indicating whether the Stan model should be recompiled. This may be necessary if you are running bridge sampling on another machine than the one used to fit the model. No recompilation is done by default. |
| <code>...</code> | Additional arguments passed to bridge_sampler.stanfit . |

Details

Computing the marginal likelihood requires samples of all variables defined in Stan's parameters block to be saved. Otherwise `bridge_sampler` cannot be computed. Thus, please set `save_pars = save_pars(all = TRUE)` in the call to `brm`, if you are planning to apply `bridge_sampler` to your models.

The computation of marginal likelihoods based on bridge sampling requires a lot more posterior draws than usual. A good conservative rule of thumb is perhaps 10-fold more draws (read: the default of 4000 draws may not be enough in many cases). If not enough posterior draws are provided, the bridge sampling algorithm tends to be unstable leading to considerably different results each time it is run. We thus recommend running `bridge_sampler` multiple times to check the stability of the results.

More details are provided under [bridgesampling::bridge_sampler](#).

See Also

[bayes_factor](#), [post_prob](#)

Examples

```
## Not run:
# model with the treatment effect
fit1 <- brm(
```

```

count ~ zAge + zBase + Trt,
data = epilepsy, family = negbinomial(),
prior = prior(normal(0, 1), class = b),
save_pars = save_pars(all = TRUE)
)
summary(fit1)
bridge_sampler(fit1)

# model without the treatment effect
fit2 <- brm(
  count ~ zAge + zBase,
  data = epilepsy, family = negbinomial(),
  prior = prior(normal(0, 1), class = b),
  save_pars = save_pars(all = TRUE)
)
summary(fit2)
bridge_sampler(fit2)

## End(Not run)

```

brm*Fit Bayesian Generalized (Non-)Linear Multivariate Multilevel Models*

Description

Fit Bayesian generalized (non-)linear multivariate multilevel models using Stan for full Bayesian inference. A wide range of distributions and link functions are supported, allowing users to fit – among others – linear, robust linear, count data, survival, response times, ordinal, zero-inflated, hurdle, extended-support beta regression, and even self-defined mixture models all in a multilevel context. Further modeling options include non-linear and smooth terms, auto-correlation structures, censored data, meta-analytic standard errors, and quite a few more. In addition, all parameters of the response distributions can be predicted in order to perform distributional regression. Prior specifications are flexible and explicitly encourage users to apply prior distributions that actually reflect their beliefs. In addition, model fit can easily be assessed and compared with posterior predictive checks and leave-one-out cross-validation.

Usage

```

brm(
  formula,
  data,
  family = gaussian(),
  prior = NULL,
  autocor = NULL,
  data2 = NULL,
  cov_ranef = NULL,
  sample_prior = "no",

```

```

sparse = NULL,
knots = NULL,
drop_unused_levels = TRUE,
stanvars = NULL,
stan_funs = NULL,
fit = NA,
save_pars = getOption("brms.save_pars", NULL),
save_ranef = NULL,
save_mevars = NULL,
save_all_pars = NULL,
init = NULL,
inits = NULL,
chains = 4,
iter = getOption("brms.iter", 2000),
warmup = floor(iter/2),
thin = 1,
cores = getOption("mc.cores", 1),
threads = getOption("brms.threads", NULL),
opencl = getOption("brms.opencl", NULL),
normalize = getOption("brms.normalize", TRUE),
control = NULL,
algorithm = getOption("brms.algorithm", "sampling"),
backend = getOption("brms.backend", "rstan"),
future = getOption("future", FALSE),
silent = 1,
seed = NA,
save_model = NULL,
stan_model_args = list(),
file = NULL,
file_compress = TRUE,
file_refit = getOption("brms.file_refit", "never"),
empty = FALSE,
rename = TRUE,
...
)

```

Arguments

formula	An object of class formula , brmsformula , or mvbrmsformula (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in brmsformula .
data	An object of class data.frame (or one that can be coerced to that class) containing data of all variables used in the model.
family	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see brmsfamily . By

	default, a linear gaussian model is applied. In multivariate models, <code>family</code> might also be a list of families.
<code>prior</code>	One or more <code>brmsprior</code> objects created by <code>set_prior</code> or related functions and combined using the <code>c</code> method or the <code>+</code> operator. See also <code>default_prior</code> for more help.
<code>autocor</code>	(Deprecated) An optional <code>cor_brms</code> object describing the correlation structure within the response variable (i.e., the 'autocorrelation'). See the documentation of <code>cor_brms</code> for a description of the available correlation structures. Defaults to <code>NULL</code> , corresponding to no correlations. In multivariate models, <code>autocor</code> might also be a list of autocorrelation structures. It is now recommended to specify autocorrelation terms directly within <code>formula</code> . See <code>brmsformula</code> for more details.
<code>data2</code>	A named list of objects containing data, which cannot be passed via argument <code>data</code> . Required for some objects used in autocorrelation structures to specify dependency structures as well as for within-group covariance matrices.
<code>cov_ranef</code>	(Deprecated) A list of matrices that are proportional to the (within) covariance structure of the group-level effects. The names of the matrices should correspond to columns in <code>data</code> that are used as grouping factors. All levels of the grouping factor should appear as rownames of the corresponding matrix. This argument can be used, among others to model pedigrees and phylogenetic effects. It is now recommended to specify those matrices in the formula interface using the <code>gr</code> and related functions. See <code>vignette("brms_phylogenetics")</code> for more details.
<code>sample_prior</code>	Indicate if draws from priors should be drawn additionally to the posterior draws. Options are "no" (the default), "yes", and "only". Among others, these draws can be used to calculate Bayes factors for point hypotheses via <code>hypothesis</code> . Please note that improper priors are not sampled, including the default improper priors used by <code>brm</code> . See <code>set_prior</code> on how to set (proper) priors. Please also note that prior draws for the overall intercept are not obtained by default for technical reasons. See <code>brmsformula</code> how to obtain prior draws for the intercept. If <code>sample_prior</code> is set to "only", draws are drawn solely from the priors ignoring the likelihood, which allows among others to generate draws from the prior predictive distribution. In this case, all parameters must have proper priors.
<code>sparse</code>	(Deprecated) Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to FALSE). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased. It is now recommended to use the <code>sparse</code> argument of <code>brmsformula</code> and related functions.
<code>knots</code>	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See <code>gamm</code> for more details.
<code>drop_unused_levels</code>	Should unused factors levels in the data be dropped? Defaults to TRUE.
<code>stanvars</code>	An optional <code>stanvars</code> object generated by function <code>stanvar</code> to define additional variables for use in Stan's program blocks.
<code>stan_funs</code>	(Deprecated) An optional character string containing self-defined Stan functions, which will be included in the functions block of the generated Stan code. It is now recommended to use the <code>stanvars</code> argument for this purpose instead.

<code>fit</code>	An instance of S3 class <code>brmsfit</code> derived from a previous fit; defaults to NA. If <code>fit</code> is of class <code>brmsfit</code> , the compiled model associated with the fitted result is re-used and all arguments modifying the model code or data are ignored. It is not recommended to use this argument directly, but to call the <code>update</code> method, instead.
<code>save_pars</code>	An object generated by <code>save_pars</code> controlling which parameters should be saved in the model. The argument has no impact on the model fitting itself.
<code>save_ranef</code>	(Deprecated) A flag to indicate if group-level effects for each level of the grouping factor(s) should be saved (default is TRUE). Set to FALSE to save memory. The argument has no impact on the model fitting itself.
<code>save_mevars</code>	(Deprecated) A flag to indicate if draws of latent noise-free variables obtained by using <code>me</code> and <code>mi</code> terms should be saved (default is FALSE). Saving these draws allows to better use methods such as <code>predict</code> with the latent variables but leads to very large R objects even for models of moderate size and complexity.
<code>save_all_pars</code>	(Deprecated) A flag to indicate if draws from all variables defined in Stan's <code>parameters</code> block should be saved (default is FALSE). Saving these draws is required in order to apply the methods <code>bridge_sampler</code> , <code>bayes_factor</code> , and <code>post_prob</code> . Can be set globally for the current R session via the " <code>brms.save_pars</code> " option (see options).
<code>init</code>	Initial values for the sampler. If NULL (the default) or "random", Stan will randomly generate initial values for parameters in a reasonable range. If 0, all parameters are initialized to zero on the unconstrained space. This option is sometimes useful for certain families, as it happens that default random initial values cause draws to be essentially constant. Generally, setting <code>init = 0</code> is worth a try, if chains do not initialize or behave well. Alternatively, <code>init</code> can be a list of lists containing the initial values, or a function (or function name) generating initial values. The latter options are mainly implemented for internal testing but are available to users if necessary. If specifying initial values using a list or a function then currently the parameter names must correspond to the names used in the generated Stan code (not the names used in R). For more details on specifying initial values you can consult the documentation of the selected backend.
<code>inits</code>	(Deprecated) Alias of <code>init</code> .
<code>chains</code>	Number of Markov chains (defaults to 4).
<code>iter</code>	Number of total iterations per chain (including warmup; defaults to 2000). Can be set globally for the current R session via the " <code>brms.iter</code> " option (see options).
<code>warmup</code>	A positive integer specifying number of warmup (aka burnin) iterations. This also specifies the number of iterations used for stepsize adaptation, so warmup draws should not be used for inference. The number of warmup should not be larger than <code>iter</code> and the default is <code>iter/2</code> .
<code>thin</code>	Thinning rate. Must be a positive integer. Set <code>thin > 1</code> to save memory and computation time if <code>iter</code> is large.
<code>cores</code>	Number of cores to use when executing the chains in parallel, which defaults to 1 but we recommend setting the <code>mc.cores</code> option to be as many processors as the hardware and RAM allow (up to the number of chains). For non-Windows OS in non-interactive R sessions, forking is used instead of PSOCK clusters.

threads	Number of threads to use in within-chain parallelization. For more control over the threading process, <code>threads</code> may also be a <code>brms\$threads</code> object created by threading . Within-chain parallelization is experimental! We recommend its use only if you are experienced with Stan's <code>reduce_sum</code> function and have a slow running model that cannot be sped up by any other means. Can be set globally for the current R session via the " <code>brms.threads</code> " option (see options).
opencl	The platform and device IDs of the OpenCL device to use for fitting using GPU support. If you don't know the IDs of your OpenCL device, <code>c(0, 0)</code> is most likely what you need. For more details, see opencl . Can be set globally for the current R session via the " <code>brms.opencl</code> " option
normalize	Logical. Indicates whether normalization constants should be included in the Stan code (defaults to TRUE). Setting it to FALSE requires Stan version >= 2.25 to work. If FALSE, sampling efficiency may be increased but some post processing functions such as bridge_sampler will not be available. Can be controlled globally for the current R session via the ' <code>brms.normalize</code> ' option.
control	A named list of parameters to control the sampler's behavior. It defaults to NULL so all the default values are used. The most important control parameters are discussed in the 'Details' section below. For a comprehensive overview see stan .
algorithm	Character string naming the estimation approach to use. Options are "sampling" for MCMC (the default), "meanfield" for variational inference with independent normal distributions, "fullrank" for variational inference with a multivariate normal distribution, "pathfinder" for the pathfinder algorithm, "laplace" for the laplace approximation, or "fixed_param" for sampling from fixed parameter values. Can be set globally for the current R session via the " <code>brms.algorithm</code> " option (see options).
backend	Character string naming the package to use as the backend for fitting the Stan model. Options are "rstan" (the default) or "cmdstanr". Can be set globally for the current R session via the " <code>brms.backend</code> " option (see options). Details on the rstan and cmdstanr packages are available at https://mc-stan.org/rstan/ and https://mc-stan.org/cmdstanr/ , respectively. Additionally a "mock" backend is available to make testing brms and packages that depend on it easier. The "mock" backend does not actually do any fitting, it only checks the generated Stan code for correctness and then returns whatever is passed in an additional <code>mock_fit</code> argument as the result of the fit.
future	Logical; If TRUE, the future package is used for parallel execution of the chains and argument <code>cores</code> will be ignored. Can be set globally for the current R session via the " <code>future</code> " option. The execution type is controlled via plan (see the examples section below).
silent	Verbosity level between 0 and 2. If 1 (the default), most of the informational messages of compiler and sampler are suppressed. If 2, even more messages are suppressed. The actual sampling progress is still printed. Set <code>refresh = 0</code> to turn this off as well. If using <code>backend = "rstan"</code> you can also set <code>open_progress = FALSE</code> to prevent opening additional progress bars.
seed	The seed for random number generation to make results reproducible. If NA (the default), Stan will set the seed randomly.

<code>save_model</code>	Either NULL or a character string. In the latter case, the model's Stan code is saved via <code>cat</code> in a text file named after the string supplied in <code>save_model</code> .
<code>stan_model_args</code>	A list of further arguments passed to <code>rstan::stan_model</code> for <code>backend = "rstan"</code> or to <code>cmdstanr::cmdstan_model</code> for <code>backend = "cmdstanr"</code> , which allows to change how models are compiled.
<code>file</code>	Either NULL or a character string. In the latter case, the fitted model object is saved via <code>saveRDS</code> in a file named after the string supplied in <code>file</code> . The <code>.rds</code> extension is added automatically. If the file already exists, <code>brm</code> will load and return the saved model object instead of refitting the model. Unless you specify the <code>file_refit</code> argument as well, the existing files won't be overwritten, you have to manually remove the file in order to refit and save the model under an existing file name. The file name is stored in the <code>brmsfit</code> object for later usage.
<code>file_compress</code>	Logical or a character string, specifying one of the compression algorithms supported by <code>saveRDS</code> . If the <code>file</code> argument is provided, this compression will be used when saving the fitted model object.
<code>file_refit</code>	Modifies when the fit stored via the <code>file</code> argument is re-used. Can be set globally for the current R session via the " <code>brms.file_refit</code> " option (see options). For "never" (default) the fit is always loaded if it exists and fitting is skipped. For "always" the model is always refitted. If set to "on_change", <code>brms</code> will refit the model if model, data or algorithm as passed to Stan differ from what is stored in the file. This also covers changes in priors, <code>sample_prior</code> , <code>stanvars</code> , covariance structure, etc. If you believe there was a false positive, you can use <code>brmsfit_needs_refit</code> to see why refit is deemed necessary. Refit will not be triggered for changes in additional parameters of the fit (e.g., initial values, number of iterations, control arguments, ...). A known limitation is that a refit will be triggered if within-chain parallelization is switched on/off.
<code>empty</code>	Logical. If TRUE, the Stan model is not created and compiled and the corresponding 'fit' slot of the <code>brmsfit</code> object will be empty. This is useful if you have estimated a <code>brms</code> -created Stan model outside of <code>brms</code> and want to feed it back into the package.
<code>rename</code>	For internal use only.
...	Further arguments passed to Stan. For <code>backend = "rstan"</code> the arguments are passed to <code>sampling</code> or <code>vb</code> . For <code>backend = "cmdstanr"</code> the arguments are passed to the <code>cmdstanr::sample</code> or <code>cmdstanr::variational</code> method.

Details

Fit a generalized (non-)linear multivariate multilevel model via full Bayesian inference using Stan. A general overview is provided in the vignettes `vignette("brms_overview")` and `vignette("brms_multilevel")`. For a full list of available vignettes see `vignette(package = "brms")`.

Formula syntax of `brms` models

Details of the formula syntax applied in `brms` can be found in `brmsformula`.

Families and link functions

Details of families supported by `brms` can be found in `brmsfamily`.

Prior distributions

Priors should be specified using the [set_prior](#) function. Its documentation contains detailed information on how to correctly specify priors. To find out on which parameters or parameter classes priors can be defined, use [default_prior](#). Default priors are chosen to be non or very weakly informative so that their influence on the results will be negligible and you usually don't have to worry about them. However, after getting more familiar with Bayesian statistics, I recommend you to start thinking about reasonable informative priors for your model parameters: Nearly always, there is at least some prior information available that can be used to improve your inference.

Adjusting the sampling behavior of Stan

In addition to choosing the number of iterations, warmup draws, and chains, users can control the behavior of the NUTS sampler, by using the `control` argument. The most important reason to use `control` is to decrease (or eliminate at best) the number of divergent transitions that cause a bias in the obtained posterior draws. Whenever you see the warning "There were x divergent transitions after warmup." you should really think about increasing `adapt_delta`. To do this, write `control = list(adapt_delta = <x>)`, where `<x>` should usually be value between 0.8 (current default) and 1. Increasing `adapt_delta` will slow down the sampler but will decrease the number of divergent transitions threatening the validity of your posterior draws.

Another problem arises when the depth of the tree being evaluated in each iteration is exceeded. This is less common than having divergent transitions, but may also bias the posterior draws. When it happens, **Stan** will throw out a warning suggesting to increase `max_treedepth`, which can be accomplished by writing `control = list(max_treedepth = <x>)` with a positive integer `<x>` that should usually be larger than the current default of 10. For more details on the `control` argument see [stan](#).

Value

An object of class `brmsfit`, which contains the posterior draws along with many other useful information about the model. Use `methods(class = "brmsfit")` for an overview on available methods.

Parallelization with multiple CPU cores

brms can make use of multiple CPU cores in parallel to speed up computations in various ways. For efficient use of the available resources it is recommended to only use parallelism to an extend such that the available physical CPUs are not oversubscribed. For example, when you have 8 CPU cores locally available, then you may consider to run 4 chains with 2 threads per chain for best performance if you happen to just run a single model. In case you run a simulation study which requires to run many times a given model, then neither chain nor within-chain parallelization is advisable as the computational resources are already exhausted by the simulation study and any further parallelization beyond the simulation study itself will in fact slow down the overall runtime. Please be aware that for historical reasons the nomenclature of the arguments is possibly confusing. The `cores` argument refers to running different chains in parallel and the within-chain parallelization will allocate for each chain as many threads as requested. The requested threads therefore increase the use of overall CPUs in a multiplicative way.

For more advanced parallelization (including beyond single model fits), **brms** also integrates with the **future** package. Importantly, this enables seamless integration with the **mirai** parallelization framework through the use of the **future.mirai** adapter. With **mirai** local and remote machines can be used in a fully transparent manner to the user. This includes the possibility to use large number

of remote machines running in the context of a computer cluster, which are managed with queuing systems. Please refer to the section on distributed computing of [mirai::daemons](#).

Author(s)

Paul-Christian Buerkner <paul.buerkner@gmail.com>

References

- Paul-Christian Buerkner (2017). brms: An R Package for Bayesian Multilevel Models Using Stan. *Journal of Statistical Software*, 80(1), 1-28. doi:10.18637/jss.v080.i01
- Paul-Christian Buerkner (2018). Advanced Bayesian Multilevel Modeling with the R Package brms. *The R Journal*. 10(1), 395–411. doi:10.32614/RJ-2018-017

See Also

[brms](#), [brmsformula](#), [brmsfamily](#), [brmsfit](#)

Examples

```
## Not run:
# Poisson regression for the number of seizures in epileptic patients
fit1 <- brm(
  count ~ zBase * Trt + (1|patient),
  data = epilepsy, family = poisson(),
  prior = prior(normal(0, 10), class = b) +
    prior(cauchy(0, 2), class = sd)
)

# generate a summary of the results
summary(fit1)

# plot the MCMC chains as well as the posterior distributions
plot(fit1)

# predict responses based on the fitted model
head(predict(fit1))

# plot conditional effects for each predictor
plot(conditional_effects(fit1), ask = FALSE)

# investigate model fit
loo(fit1)
pp_check(fit1)

# Ordinal regression modeling patient's rating of inhaler instructions
# category specific effects are estimated for variable 'treat'
fit2 <- brm(rating ~ period + carry + cs(treat),
             data = inhaler, family = sratio("logit"),
             prior = set_prior("normal(0,5)", chains = 2)
summary(fit2)
```

```
plot(fit2, ask = FALSE)
WAIC(fit2)

# Survival regression modeling the time between the first
# and second recurrence of an infection in kidney patients.
fit3 <- brm(time | cens(censored) ~ age * sex + disease + (1|patient),
             data = kidney, family = lognormal())
summary(fit3)
plot(fit3, ask = FALSE)
plot(conditional_effects(fit3), ask = FALSE)

# Probit regression using the binomial family
ntrials <- sample(1:10, 100, TRUE)
success <- rbinom(100, size = ntrials, prob = 0.4)
x <- rnorm(100)
data4 <- data.frame(ntrials, success, x)
fit4 <- brm(success | trials(ntrials) ~ x, data = data4,
              family = binomial("probit"))
summary(fit4)

# Non-linear Gaussian model
fit5 <- brm(
  bf(cum ~ ult * (1 - exp(-(dev/theta)^omega)),
     ult ~ 1 + (1|AY), omega ~ 1, theta ~ 1,
     nl = TRUE),
  data = loss, family = gaussian(),
  prior = c(
    prior(normal(5000, 1000), nlpar = "ult"),
    prior(normal(1, 2), nlpar = "omega"),
    prior(normal(45, 10), nlpar = "theta")
  ),
  control = list(adapt_delta = 0.9)
)
summary(fit5)
conditional_effects(fit5)

# Normal model with heterogeneous variances
data_het <- data.frame(
  y = c(rnorm(50), rnorm(50, 1, 2)),
  x = factor(rep(c("a", "b"), each = 50))
)
fit6 <- brm(bf(y ~ x, sigma ~ 0 + x), data = data_het)
summary(fit6)
plot(fit6)
conditional_effects(fit6)

# extract estimated residual SDs of both groups
sigmas <- exp(as.data.frame(fit6, variable = "^b_sigma_", regex = TRUE))
ggplot(stack(sigmas), aes(values)) +
```

```

geom_density(aes(fill = ind))

# Quantile regression predicting the 25%-quantile
fit7 <- brm(bf(y ~ x, quantile = 0.25), data = data_het,
             family = asym_laplace())
summary(fit7)
conditional_effects(fit7)

# use the future package for more flexible parallelization
library(future)
plan(multisession, workers = 4)
fit7a <- update(fit7, future = TRUE)

# we may also use mirai allowing the use of remote machines as
# well, including the use of computer clusters managed by queuing
# systems
mirai::daemons(4) # distributed computing with url and remote argument
plan(future.mirai::mirai_cluster)
fit7b <- update(fit7, future = TRUE)
mirai::daemons(0) # shuts down the 4 workers

# fit a model manually via rstan
scode <- stancode(count ~ Trt, data = epilepsy)
sdata <- standata(count ~ Trt, data = epilepsy)
stanfit <- rstan::stan(model_code = scode, data = sdata)
# feed the Stan model back into brms
fit8 <- brm(count ~ Trt, data = epilepsy, empty = TRUE)
fit8$fit <- stanfit
fit8 <- rename_pars(fit8)
summary(fit8)

## End(Not run)

```

Description

Family objects provide a convenient way to specify the details of the models used by many model fitting functions. The family functions presented here are for use with **brms** only and will ****not**** work with other model fitting functions such as **glm** or **glmer**. However, the standard family functions as described in [family](#) will work with **brms**. You can also specify custom families for use in **brms** with the [custom_family](#) function.

Usage

```
brmsfamily(
  family,
```

```
link = NULL,
link_sigma = "log",
link_shape = "log",
link_nu = "logm1",
link_phi = "log",
link_kappa = "log",
link_beta = "log",
link_zi = "logit",
link_hu = "logit",
link_zoi = "logit",
link_coi = "logit",
link_disc = "log",
link_bs = "log",
link_ndt = "log",
link_bias = "logit",
link_xi = "log1p",
link_alpha = "identity",
link_quantile = "logit",
threshold = "flexible",
refcat = NULL
)

student(link = "identity", link_sigma = "log", link_nu = "logm1")

bernoulli(link = "logit")

beta_binomial(link = "logit", link_phi = "log")

negbinomial(link = "log", link_shape = "log")

geometric(link = "log")

lognormal(link = "identity", link_sigma = "log")

shifted_lognormal(link = "identity", link_sigma = "log", link_ndt = "log")

skew_normal(link = "identity", link_sigma = "log", link_alpha = "identity")

exponential(link = "log")

weibull(link = "log", link_shape = "log")

frechet(link = "log", link_nu = "logm1")

gen_extreme_value(link = "identity", link_sigma = "log", link_xi = "log1p")

exgaussian(link = "identity", link_sigma = "log", link_beta = "log")
```

```

wiener(
  link = "identity",
  link_bs = "log",
  link_ndt = "log",
  link_bias = "logit"
)

Beta(link = "logit", link_phi = "log")

xbeta(link = "logit", link_phi = "log", link_kappa = "log")

dirichlet(link = "logit", link_phi = "log", refcat = NULL)

logistic_normal(link = "identity", link_sigma = "log", refcat = NULL)

von_mises(link = "tan_half", link_kappa = "log")

asym_laplace(link = "identity", link_sigma = "log", link_quantile = "logit")

cox(link = "log")

hurdle_poisson(link = "log", link_hu = "logit")

hurdle_negbinomial(link = "log", link_shape = "log", link_hu = "logit")

hurdle_gamma(link = "log", link_shape = "log", link_hu = "logit")

hurdle_lognormal(link = "identity", link_sigma = "log", link_hu = "logit")

hurdle_cumulative(
  link = "logit",
  link_hu = "logit",
  link_disc = "log",
  threshold = "flexible"
)

zero_inflated_beta(link = "logit", link_phi = "log", link_zi = "logit")

zero_one_inflated_beta(
  link = "logit",
  link_phi = "log",
  link_zoi = "logit",
  link_coi = "logit"
)

zero_inflated_poisson(link = "log", link_zi = "logit")

zero_inflated_negbinomial(link = "log", link_shape = "log", link_zi = "logit")

```

```

zero_inflated_binomial(link = "logit", link_zi = "logit")

zero_inflated_beta_binomial(
  link = "logit",
  link_phi = "log",
  link_zi = "logit"
)

categorical(link = "logit", refcat = NULL)

multinomial(link = "logit", refcat = NULL)

dirichlet_multinomial(link = "logit", link_phi = "log", refcat = NULL)

cumulative(link = "logit", link_disc = "log", threshold = "flexible")

sratio(link = "logit", link_disc = "log", threshold = "flexible")

cratio(link = "logit", link_disc = "log", threshold = "flexible")

acat(link = "logit", link_disc = "log", threshold = "flexible")

```

Arguments

family	A character string naming the distribution family of the response variable to be used in the model. Currently, the following families are supported: gaussian, student, binomial, bernoulli, beta-binomial, poisson, negbinomial, geometric, Gamma, skew_normal, lognormal, shifted_lognormal, exgaussian, wiener, inverse.gaussian, exponential, weibull, frechet, Beta, dirichlet, von_mises, asym_laplace, gen_extreme_value, categorical, multinomial, dirichlet_multinomial, cumulative, cratio, sratio, acat, hurdle_poisson, hurdle_negbinomial, hurdle_gamma, hurdle_lognormal, hurdle_cumulative, zero_inflated_binomial, zero_inflated_beta_binomial, zero_inflated_beta, zero_inflated_negbinomial, zero_inflated_poisson, zero_one_inflated_beta, and xbeta.
link	A specification for the model link function. This can be a name/expression or character string. See the 'Details' section for more information on link functions supported by each family.
link_sigma	Link of auxiliary parameter sigma if being predicted.
link_shape	Link of auxiliary parameter shape if being predicted.
link_nu	Link of auxiliary parameter nu if being predicted.
link_phi	Link of auxiliary parameter phi if being predicted.
link_kappa	Link of auxiliary parameter kappa if being predicted.
link_beta	Link of auxiliary parameter beta if being predicted.
link_zi	Link of auxiliary parameter zi if being predicted.
link_hu	Link of auxiliary parameter hu if being predicted.

link_zoi	Link of auxiliary parameter zoi if being predicted.
link_coi	Link of auxiliary parameter coi if being predicted.
link_disc	Link of auxiliary parameter disc if being predicted.
link_bs	Link of auxiliary parameter bs if being predicted.
link_ndt	Link of auxiliary parameter ndt if being predicted.
link_bias	Link of auxiliary parameter bias if being predicted.
link_xi	Link of auxiliary parameter xi if being predicted.
link_alpha	Link of auxiliary parameter alpha if being predicted.
link_quantile	Link of auxiliary parameter quantile if being predicted.
threshold	A character string indicating the type of thresholds (i.e. intercepts) used in an ordinal model. "flexible" provides the standard unstructured thresholds, "equidistant" restricts the distance between consecutive thresholds to the same value, and "sum_to_zero" ensures the thresholds sum to zero.
refcat	Optional name of the reference response category used in <code>categorical</code> , <code>multinomial</code> , <code>dirichlet</code> , <code>dirichlet_multinomial</code> and <code>logistic_normal</code> models. If <code>NULL</code> (the default), the first category is used as the reference. If <code>NA</code> , all categories will be predicted, which requires strong priors or carefully specified predictor terms in order to lead to an identified model.

Details

Below, we list common use cases for the different families. This list is not meant to be exhaustive.

- Family `gaussian` can be used for linear regression.
- Family `student` can be used for robust linear regression that is less influenced by outliers.
- Family `skew_normal` can handle skewed responses in linear regression.
- Families `poisson`, `negbinomial`, and `geometric` can be used for regression of unbounded count data.
- Families `bernoulli`, `binomial`, and `beta_binomial` can be used for binary regression (i.e., most commonly logistic regression).
- Families `categorical`, `multinomial` and `dirichlet_multinomial` can be used for multi-logistic regression when there are more than two possible outcomes.
- Families `cumulative`, `cratio` ('continuation ratio'), `sratio` ('stopping ratio'), and `acat` ('adjacent category') leads to ordinal regression.
- Families `Gamma`, `weibull`, `exponential`, `lognormal`, `frechet`, `inverse.gaussian`, and `cox` (Cox proportional hazards model) can be used (among others) for time-to-event regression also known as survival regression.
- Families `weibull`, `frechet`, and `gen_extreme_value` ('generalized extreme value') allow for modeling extremes.
- Families `beta`, `dirichlet`, and `logistic_normal` can be used to model responses representing rates or probabilities.

- Family `xbeta` extends the beta family to support $[0, 1]$ responses with exact 0s and / or 1s, when each response takes values 0, 1, and $(0, 1)$ according to a single process. If there is merit in assuming that 0 and 1 values arise from different processes than $(0, 1)$ values, then the `zero_inflated_beta`, `zero_one_inflated_beta` families provide more flexibility. For details see Kosmidis & Zeileis (2024).
- Family `asym_laplace` allows for quantile regression when fixing the auxiliary quantile parameter to the quantile of interest.
- Family `exgaussian` ('exponentially modified Gaussian') and `shifted_lognormal` are especially suited to model reaction times.
- Family `wiener` provides an implementation of the Wiener diffusion model. For this family, the main formula predicts the drift parameter 'delta' and all other parameters are modeled as auxiliary parameters (see `brmsformula` for details).
- Families `hurdle_poisson`, `hurdle_negbinomial`, `hurdle_gamma`, `hurdle_lognormal`, `zero_inflated_poisson`, `zero_inflated_negbinomial`, `zero_inflated_binomial`, `zero_inflated_beta_binomial`, `zero_inflated_beta`, `zero_one_inflated_beta`, and `hurdle_cumulative` allow to estimate zero-inflated and hurdle models. These models can be very helpful when there are many zeros in the data (or ones in case of one-inflated models) that cannot be explained by the primary distribution of the response.

Below, we list all possible links for each family. The first link mentioned for each family is the default.

- Families `gaussian`, `student`, `skew_normal`, `exgaussian`, `asym_laplace`, and `gen_extreme_value` support the links (as names) `identity`, `log`, `inverse`, and `softplus`.
- Families `poisson`, `negbinomial`, `geometric`, `zero_inflated_poisson`, `zero_inflated_negbinomial`, `hurdle_poisson`, and `hurdle_negbinomial` support `log`, `identity`, `sqrt`, and `softplus`.
- Families `binomial`, `bernoulli`, `beta_binomial`, `zero_inflated_binomial`, `zero_inflated_beta_binomial`, `Beta`, `zero_inflated_beta`, `zero_one_inflated_beta`, and `xbeta` support `logit`, `probit`, `probit_approx`, `cloglog`, `cauchit`, `identity`, and `log`.
- Families `cumulative`, `cratio`, `sratio`, `acat`, and `hurdle_cumulative` support `logit`, `probit`, `probit_approx`, `cloglog`, and `cauchit`.
- Families `categorical`, `multinomial`, `dirichlet_multinomial` and `dirichlet` support `logit`.
- Families `Gamma`, `weibull`, `exponential`, `frechet`, and `hurdle_gamma` support `log`, `identity`, `inverse`, and `softplus`.
- Families `lognormal` and `hurdle_lognormal` support `identity` and `inverse`.
- Family `logistic_normal` supports `identity`.
- Family `inverse.gaussian` supports $1/\mu^2$, `inverse`, `identity`, `log`, and `softplus`.
- Family `von_mises` supports `tan_half` and `identity`.
- Family `cox` supports `log`, `identity`, and `softplus` for the proportional hazards parameter.
- Family `wiener` supports `identity`, `log`, and `softplus` for the main parameter which represents the drift rate.

Please note that when calling the `Gamma` family function of the `stats` package, the default link will be `inverse` instead of `log` although the latter is the default in `brms`. Also, when using the family functions `gaussian`, `binomial`, `poisson`, and `Gamma` of the `stats` package (see `family`), special link functions such as `softplus` or `cauchit` won't work. In this case, you have to use `brmsfamily` to specify the family with corresponding link function.

References

Kosmidis I, Zeileis A (2024). Extended-Support Beta Regression for [0, 1] Responses. *arXiv Preprint*. doi:10.48550/arXiv.2409.07233

See Also

[brm](#), [family](#), [customfamily](#)

Examples

```
# create a family object
(fam1 <- student("log"))
# alternatively use the brmsfamily function
(fam2 <- brmsfamily("student", "log"))
# both leads to the same object
identical(fam1, fam2)
```

brmsfit-class

*Class brmsfit of models fitted with the **brms** package*

Description

Models fitted with the [brms](#) package are represented as a **brmsfit** object, which contains the posterior draws (samples), model formula, Stan code, relevant data, and other information.

Details

See `methods(class = "brmsfit")` for an overview of available methods.

Slots

- `formula` A [brmsformula](#) object.
- `data` A `data.frame` containing all variables used in the model.
- `data2` A list of data objects which cannot be passed via `data`.
- `prior` A [brmsprior](#) object containing information on the priors used in the model.
- `stanvars` A [stanvars](#) object.
- `model` The model code in **Stan** language.
- `exclude` The names of the parameters for which draws are not saved.
- `algorithm` The name of the algorithm used to fit the model.
- `backend` The name of the backend used to fit the model.
- `threads` An object of class ‘`brmstthreads`‘ created by [threading](#).
- `opencl` An object of class ‘`brmsopencl`‘ created by [opencl](#).

`stan_args` Named list of additional control arguments that were passed to the Stan backend directly.

`fit` An object of class `stanfit` among others containing the posterior draws.

`basis` An object that contains a small subset of the Stan data created at fitting time, which is needed to process new data correctly.

`criteria` An empty list for adding model fit criteria after estimation of the model.

`file` Optional name of a file in which the model object was stored in or loaded from.

`version` The versions of `brms` and `rstan` with which the model was fitted.

`family` (Deprecated) A `brmsfamily` object.

`autocor` (Deprecated) An `cor_brms` object containing the autocorrelation structure if specified.

`ranef` (Deprecated) A `data.frame` containing the group-level structure.

`cov_ranef` (Deprecated) A list of customized group-level covariance matrices.

`stan_funs` (Deprecated) A character string of length one or `NULL`.

`data.name` (Deprecated) The name of data as specified by the user.

See Also

`brms`, `brm`, `brmsformula`, `brmsfamily`

brmsformula

Set up a model formula for use in brms

Description

Set up a model formula for use in the `brms` package allowing to define (potentially non-linear) additive multilevel models for all parameters of the assumed response distribution.

Usage

```
brmsformula(  
  formula,  
  ...,  
  flist = NULL,  
  family = NULL,  
  autocor = NULL,  
  nl = NULL,  
  loop = NULL,  
  center = NULL,  
  cmc = NULL,  
  sparse = NULL,  
  decomp = NULL,  
  unused = NULL  
)
```

Arguments

<code>formula</code>	An object of class <code>formula</code> (or one that can be coerced to that class): a symbolic description of the model to be fitted. The details of model specification are given in 'Details'.
<code>...</code>	Additional <code>formula</code> objects to specify predictors of non-linear and distributional parameters. Formulas can either be named directly or contain names on their left-hand side. Alternatively, it is possible to fix parameters to certain values by passing numbers or character strings in which case arguments have to be named to provide the parameter names. See 'Details' for more information.
<code>flist</code>	Optional list of formulas, which are treated in the same way as formulas passed via the <code>...</code> argument.
<code>family</code>	Same argument as in <code>brm</code> . If <code>family</code> is specified in <code>brmsformula</code> , it will overwrite the value specified in other functions.
<code>autocor</code>	An optional formula which contains autocorrelation terms as described in autocor-terms or alternatively a <code>cor_brms</code> object (deprecated). If <code>autocor</code> is specified in <code>brmsformula</code> , it will overwrite the value specified in other functions.
<code>n1</code>	Logical; Indicates whether <code>formula</code> should be treated as specifying a non-linear model. By default, <code>formula</code> is treated as an ordinary linear model formula.
<code>loop</code>	Logical; Only used in non-linear models. Indicates if the computation of the non-linear formula should be done inside (TRUE) or outside (FALSE) a loop over observations. Defaults to TRUE.
<code>center</code>	Logical; Indicates if the population-level design matrix should be centered, which usually increases sampling efficiency. See the 'Details' section for more information. Defaults to TRUE for distributional parameters and to FALSE for non-linear parameters.
<code>cmc</code>	Logical; Indicates whether automatic cell-mean coding should be enabled when removing the intercept by adding <code>0</code> to the right-hand of model formulas. Defaults to TRUE to mirror the behavior of standard R formula parsing.
<code>sparse</code>	Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to FALSE). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased.
<code>decomp</code>	Optional name of the decomposition used for the population-level design matrix. Defaults to NULL that is no decomposition. Other options currently available are "QR" for the QR decomposition that helps in fitting models with highly correlated predictors.
<code>unused</code>	An optional <code>formula</code> which contains variables that are unused in the model but should still be stored in the model's data frame. This can be useful, for example, if those variables are required for post-processing the model.

Details

General formula structure

The `formula` argument accepts formulas of the following syntax:

```
response | aterms ~ pterms + (gterms | group)
```

The `pterms` part contains effects that are assumed to be the same across observations. We call them 'population-level' or 'overall' effects, or (adopting frequentist vocabulary) 'fixed' effects. The optional `gterms` part may contain effects that are assumed to vary across grouping variables specified in `group`. We call them 'group-level' or 'varying' effects, or (adopting frequentist vocabulary) 'random' effects, although the latter name is misleading in a Bayesian context. For more details type `vignette("brms_overview")` and `vignette("brms_multilevel")`.

Group-level terms

Multiple grouping factors each with multiple group-level effects are possible. (Of course we can also run models without any group-level effects.) Instead of `|` you may use `||` in grouping terms to prevent correlations from being modeled. Equivalently, the `cor` argument of the `gr` function can be used for this purpose, for example, `(1 + x || g)` is equivalent to `(1 + x | gr(g, cor = FALSE))`.

It is also possible to model different group-level terms of the same grouping factor as correlated (even across different formulas, e.g., in non-linear models) by using `|<ID>|` instead of `|`. All group-level terms sharing the same ID will be modeled as correlated. If, for instance, one specifies the terms `(1+x|i|g)` and `(1+z|i|g)` somewhere in the formulas passed to `brmsformula`, correlations between the corresponding group-level effects will be estimated. In the above example, `i` is not a variable in the data but just a symbol to indicate correlations between multiple group-level terms. Equivalently, the `id` argument of the `gr` function can be used as well, for example, `(1 + x | gr(g, id = "i"))`.

If levels of the grouping factor belong to different sub-populations, it may be reasonable to assume a different covariance matrix for each of the sub-populations. For instance, the variation within the treatment group and within the control group in a randomized control trial might differ. Suppose that `y` is the outcome, and `x` is the factor indicating the treatment and control group. Then, we could estimate different hyper-parameters of the varying effects (in this case a varying intercept) for treatment and control group via `y ~ x + (1 | gr(subject, by = x))`.

You can specify multi-membership terms using the `mm` function. For instance, a multi-membership term with two members could be `(1 | mm(g1, g2))`, where `g1` and `g2` specify the first and second member, respectively. Moreover, if a covariate `x` varies across the levels of the grouping-factors `g1` and `g2`, we can save the respective covariate values in the variables `x1` and `x2` and then model the varying effect as `(1 + mmc(x1, x2) | mm(g1, g2))`.

Special predictor terms

Flexible non-linear smooth terms can be modeled using the `s` and `t2` functions in the `pterms` part of the model formula. This allows to fit generalized additive mixed models (GAMMs) with `brms`. The implementation is similar to that used in the `gamm4` package. For more details on this model class see `gam` and `gamm`.

Gaussian process terms can be fitted using the `gp` function in the `pterms` part of the model formula. Similar to smooth terms, Gaussian processes can be used to model complex non-linear relationships, for instance temporal or spatial autocorrelation. However, they are computationally demanding and are thus not recommended for very large datasets or approximations need to be used.

The `pterms` and `gterms` parts may contain four non-standard effect types namely monotonic, measurement error, missing value, and category specific effects, which can be specified using terms of the form `mo(predictor)`, `me(predictor, sd_predictor)`, `mi(predictor)`, and `cs(<predictors>)`, respectively. Category specific effects can only be estimated in ordinal models and are explained in more detail in the package's main vignette (type `vignette("brms_overview")`). The other three effect types are explained in the following.

A monotonic predictor must either be integer valued or an ordered factor, which is the first difference to an ordinary continuous predictor. More importantly, predictor categories (or integers) are not assumed to be equidistant with respect to their effect on the response variable. Instead, the distance between adjacent predictor categories (or integers) is estimated from the data and may vary across categories. This is realized by parameterizing as follows: One parameter takes care of the direction and size of the effect similar to an ordinary regression parameter, while an additional parameter vector estimates the normalized distances between consecutive predictor categories. A main application of monotonic effects are ordinal predictors that can this way be modeled without (falsely) treating them as continuous or as unordered categorical predictors. For more details and examples see `vignette("brms_monotonic")`.

Quite often, predictors are measured and as such naturally contain measurement error. Although most researchers are well aware of this problem, measurement error in predictors is ignored in most regression analyses, possibly because only few packages allow for modeling it. Notably, measurement error can be handled in structural equation models, but many more general regression models (such as those featured by **brms**) cannot be transferred to the SEM framework. In **brms**, effects of noise-free predictors can be modeled using the `me` (for 'measurement error') function. If, say, `y` is the response variable and `x` is a measured predictor with known measurement error `sdx`, we can simply include it on the right-hand side of the model formula via `y ~ me(x, sdx)`. This can easily be extended to more general formulas. If `x2` is another measured predictor with corresponding error `sdx2` and `z` is a predictor without error (e.g., an experimental setting), we can model all main effects and interactions of the three predictors in the well known manner: `y ~ me(x, sdx) * me(x2, sdx2) * z`. The `me` function is soft deprecated in favor of the more flexible and consistent `mi` function (see below).

When a variable contains missing values, the corresponding rows will be excluded from the data by default (row-wise exclusion). However, quite often we want to keep these rows and instead estimate the missing values. There are two approaches for this: (a) Impute missing values before the model fitting for instance via multiple imputation (see `brm_multiple` for a way to handle multiple imputed datasets). (b) Impute missing values on the fly during model fitting. The latter approach is explained in the following. Using a variable with missing values as predictors requires two things, First, we need to specify that the predictor contains missings that should be imputed. If, say, `y` is the primary response, `x` is a predictor with missings and `z` is a predictor without missings, we go for `y ~ mi(x) + z`. Second, we need to model `x` as an additional response with corresponding predictors and the addition term `mi()`. In our example, we could write `x | mi() ~ z`. Measurement error may be included via the `sdy` argument, say, `x | mi(sdy = se) ~ z`. See `mi` for examples with real data.

Autocorrelation terms

Autocorrelation terms can be directly specified inside the `pterms` part as well. Details can be found in [autocor-terms](#).

Additional response information

Another special of the **brms** formula syntax is the optional `aterms` part, which may contain multiple terms of the form `fun(<variable>)` separated by `+` each providing special information on the response variable. `fun` can be replaced with either `se`, `weights`, `subset`, `cens`, `trunc`, `trials`, `cat`, `dec`, `rate`, `vreal`, or `vint`. Their meanings are explained below (see also [addition-terms](#)).

For families `gaussian`, `student` and `skew_normal`, it is possible to specify standard errors of the observations, thus allowing to perform meta-analysis. Suppose that the variable `yi` contains the effect sizes from the studies and `sei` the corresponding standard errors. Then, fixed and random effects meta-analyses can be conducted using the formulas `yi | se(sei) ~ 1` and `yi | se(sei) ~ 1`

$+ (1|study)$, respectively, where `study` is a variable uniquely identifying every study. If desired, meta-regression can be performed via $y_i | se(sei) \sim 1 + mod1 + mod2 + (1|study)$ or $y_i | se(sei) \sim 1 + mod1 + mod2 + (1 + mod1 + mod2|study)$, where `mod1` and `mod2` represent moderator variables. By default, the standard errors replace the parameter `sigma`. To model `sigma` in addition to the known standard errors, set argument `sigma` in function `se` to `TRUE`, for instance, $y_i | se(sei, sigma = TRUE) \sim 1$.

For all families, weighted regression may be performed using `weights` in the `aterms` part. Internally, this is implemented by multiplying the log-posterior values of each observation by their corresponding weights. Suppose that variable `wei` contains the weights and that `yi` is the response variable. Then, formula $y_i | weights(wei) \sim predictors$ implements a weighted regression.

For multivariate models, `subset` may be used in the `aterms` part, to use different subsets of the data in different univariate models. For instance, if `sub` is a logical variable and `y` is the response of one of the univariate models, we may write $y | subset(sub) \sim predictors$ so that `y` is predicted only for those observations for which `sub` evaluates to `TRUE`.

For log-linear models such as Poisson models, `rate` may be used in the `aterms` part to specify the denominator of a response that is expressed as a rate. The numerator is given by the actual response variable and has a distribution according to the family as usual. In Poisson models, using `rate(denom)` is equivalent to adding `offset(log(denom))` to the linear predictor of the main parameter. For negative-binomial models, this equivalence no longer holds and `rate(denom)` remains the statistically correct approach.

With the exception of categorical and ordinal families, left, right, and interval censoring can be modeled through $y | cens(censored) \sim predictors$. The censoring variable (named `censored` in this example) should contain the values 'left', 'none', 'right', and 'interval' (or equivalently -1, 0, 1, and 2) to indicate that the corresponding observation is left censored, not censored, right censored, or interval censored. For interval censored data, a second variable (let's call it `y2`) has to be passed to `cens`. In this case, the formula has the structure $y | cens(censored, y2) \sim predictors$. While the lower bounds are given in `y`, the upper bounds are given in `y2` for interval censored data. Intervals are assumed to be open on the left and closed on the right: $(y, y2]$.

With the exception of categorical and ordinal families, the response distribution can be truncated using the `trunc` function in the addition part. If the response variable is truncated between, say, 0 and 100, we can specify this via $y_i | trunc(lb = 0, ub = 100) \sim predictors$. Instead of numbers, variables in the data set can also be passed allowing for varying truncation points across observations. Defining only one of the two arguments in `trunc` leads to one-sided truncation.

For all continuous families, missing values in the responses can be imputed within Stan by using the addition term `mi`. This is mostly useful in combination with `mi` predictor terms as explained above under 'Special predictor terms'.

For families `binomial` and `zero_inflated_binomial`, `addition` should contain a variable indicating the number of trials underlying each observation. In `lme4` syntax, we may write for instance `cbind(success, n - success)`, which is equivalent to `success | trials(n)` in `brms` syntax. If the number of trials is constant across all observations, say 10, we may also write `success | trials(10)`. **Please note that the `cbind()` syntax will not work in `brms` in the expected way because this syntax is reserved for other purposes.**

For all ordinal families, `aterms` may contain a term `thres(number)` to specify the number thresholds (e.g, `thres(6)`), which should be equal to the total number of response categories - 1. If not given, the number of thresholds is calculated from the data. If different threshold vectors should be used for different subsets of the data, the `gr` argument can be used to provide the grouping variable

(e.g., `thres(6, gr = item)`, if `item` is the grouping variable). In this case, the number of thresholds can also be a variable in the data with different values per group.

A deprecated quasi alias of `thres()` is `cat()` with which the total number of response categories (i.e., number of thresholds + 1) can be specified.

In Wiener diffusion models (family `wiener`) the addition term `dec` is mandatory to specify the (vector of) binary decisions corresponding to the reaction times. Non-zero values will be treated as a response on the upper boundary of the diffusion process and zeros will be treated as a response on the lower boundary. Alternatively, the variable passed to `dec` might also be a character vector consisting of 'lower' and 'upper'.

All families support the `index` addition term to uniquely identify each observation of the corresponding response variable. Currently, `index` is primarily useful in combination with the subset addition and `mi` terms.

For custom families, it is possible to pass an arbitrary number of real and integer vectors via the addition terms `vreal` and `vint`, respectively. An example is provided in `vignette('brms_customfamilies')`. To pass multiple vectors of the same data type, provide them separated by commas inside a single `vreal` or `vint` statement.

Multiple addition terms of different types may be specified at the same time using the `+` operator. For example, the formula `formula = yi | se(sei) + cens(censored) ~ 1` implies a censored meta-analytic model.

The addition argument `disp` (short for dispersion) has been removed in version 2.0. You may instead use the distributional regression approach by specifying `sigma ~ 1 + offset(log(xdisp))` or `shape ~ 1 + offset(log(xdisp))`, where `xdisp` is the variable being previously passed to `disp`.

Parameterization of the population-level intercept

By default, the population-level intercept (if incorporated) is estimated separately and not as part of population-level parameter vector `b`. As a result, priors on the intercept also have to be specified separately. Furthermore, to increase sampling efficiency, the population-level design matrix `X` is centered around its column means `X_means` if the intercept is incorporated. This leads to a temporary bias in the intercept equal to $\langle X_{\text{means}}, b \rangle$, where \langle , \rangle is the scalar product. The bias is corrected after fitting the model, but be aware that you are effectively defining a prior on the intercept of the centered design matrix not on the real intercept. You can turn off this special handling of the intercept by setting argument `center` to `FALSE`. For more details on setting priors on population-level intercepts, see [set_prior](#).

This behavior can be avoided by using the reserved (and internally generated) variable `Intercept`. Instead of `y ~ x`, you may write `y ~ 0 + Intercept + x`. This way, priors can be defined on the real intercept, directly. In addition, the intercept is just treated as an ordinary population-level effect and thus priors defined on `b` will also apply to it. Note that this parameterization may be less efficient than the default parameterization discussed above.

Formula syntax for non-linear models

In `brms`, it is possible to specify non-linear models of arbitrary complexity. The non-linear model can just be specified within the `formula` argument. Suppose, that we want to predict the response `y` through the predictor `x`, where `x` is linked to `y` through `y = alpha - beta * lambda^x`, with parameters `alpha`, `beta`, and `lambda`. This is certainly a non-linear model being defined via `formula = y ~ alpha - beta * lambda^x` (addition arguments can be added in the same way as for ordinary formulas). To tell `brms` that this is a non-linear model, we set argument `n1` to `TRUE`. Now we have to specify a model for each of the non-linear parameters. Let's say we just want to estimate those

three parameters with no further covariates or random effects. Then we can pass $\text{alpha} + \text{beta} + \text{lambda} \sim 1$ or equivalently (and more flexible) $\text{alpha} \sim 1$, $\text{beta} \sim 1$, $\text{lambda} \sim 1$ to the \dots argument. This can, of course, be extended. If we have another predictor z and observations nested within the grouping factor g , we may write for instance $\text{alpha} \sim 1$, $\text{beta} \sim 1 + z + (1|g)$, $\text{lambda} \sim 1$. The formula syntax described above applies here as well. In this example, we are using z and g only for the prediction of beta , but we might also use them for the other non-linear parameters (provided that the resulting model is still scientifically reasonable).

By default, non-linear covariates are treated as real vectors in Stan. However, if the data of the covariates is of type ‘integer’ in R (which can be enforced by the ‘`as.integer`’ function), the Stan type will be changed to an integer array. That way, covariates can also be used for indexing purposes in Stan.

Non-linear models may not be uniquely identified and / or show bad convergence. For this reason it is mandatory to specify priors on the non-linear parameters. For instructions on how to do that, see `set_prior`. For some examples of non-linear models, see `vignette("brms_nonlinear")`.

Formula syntax for predicting distributional parameters

It is also possible to predict parameters of the response distribution such as the residual standard deviation `sigma` in gaussian models or the hurdle probability `hu` in hurdle models. The syntax closely resembles that of a non-linear parameter, for instance $\text{sigma} \sim x + s(z) + (1+x|g)$. For some examples of distributional models, see `vignette("brms_distreg")`.

Parameter `mu` exists for every family and can be used as an alternative to specifying terms in `formula`. If both `mu` and `formula` are given, the right-hand side of `formula` is ignored. Accordingly, specifying terms on the right-hand side of both `formula` and `mu` at the same time is deprecated. In future versions, `formula` might be updated by `mu`.

The following are distributional parameters of specific families (all other parameters are treated as non-linear parameters): `sigma` (residual standard deviation or scale of the gaussian, student, skew_normal, lognormal exgaussian, and asym_laplace families); `shape` (shape parameter of the Gamma, weibull, negbinomial, and related zero-inflated / hurdle families); `nu` (degrees of freedom parameter of the student and frechet families); `phi` (precision parameter of the beta, zero_inflated_beta, and xbeta families); `kappa` (precision parameter of the von_mises family); `beta` (mean parameter of the exponential component of the exgaussian family); `quantile` (quantile parameter of the asym_laplace family); `zi` (zero-inflation probability); `hu` (hurdle probability); `zoi` (zero-one-inflation probability); `coi` (conditional one-inflation probability); `disc` (discrimination) for ordinal models; `bs`, `ndt`, and `bias` (boundary separation, non-decision time, and initial bias of the wiener diffusion model). By default, distributional parameters are modeled on the log scale if they can be positive only or on the logit scale if they can only be within the unit interval.

Alternatively, one may fix distributional parameters to certain values. However, this is mainly useful when models become too complicated and otherwise have convergence issues. We thus suggest to be generally careful when making use of this option. The `quantile` parameter of the `asym_laplace` distribution is a good example where it is useful. By fixing `quantile`, one can perform quantile regression for the specified quantile. For instance, `quantile = 0.25` allows predicting the 25%-quantile. Furthermore, the `bias` parameter in drift-diffusion models, is assumed to be 0.5 (i.e. no bias) in many applications. To achieve this, simply write `bias = 0.5`. Other possible applications are the Cauchy distribution as a special case of the Student-t distribution with `nu = 1`, or the geometric distribution as a special case of the negative binomial distribution with `shape = 1`. Furthermore, the parameter `disc` ('discrimination') in ordinal models is fixed to 1 by default and not estimated, but may be modeled as any other distributional parameter if desired (see examples). For reasons of identification, 'disc' can only be positive, which is achieved by applying the log-link.

In categorical models, distributional parameters do not have fixed names. Instead, they are named after the response categories (excluding the first one, which serves as the reference category), with the prefix 'mu'. If, for instance, categories are named cat1, cat2, and cat3, the distributional parameters will be named mucat2 and mucat3.

Some distributional parameters currently supported by *brmsformula* have to be positive (a negative standard deviation or precision parameter does not make any sense) or are bounded between 0 and 1 (for zero-inflated / hurdle probabilities, quantiles, or the initial bias parameter of drift-diffusion models). However, linear predictors can be positive or negative, and thus the log link (for positive parameters) or logit link (for probability parameters) are used by default to ensure that distributional parameters are within their valid intervals. This implies that, by default, effects for such distributional parameters are estimated on the log / logit scale and one has to apply the inverse link function to get to the effects on the original scale. Alternatively, it is possible to use the identity link to predict parameters on their original scale, directly. However, this is much more likely to lead to problems in the model fitting, if the parameter actually has a restricted range.

See also [brmsfamily](#) for an overview of valid link functions.

Formula syntax for mixture models

The specification of mixture models closely resembles that of non-mixture models. If not specified otherwise (see below), all mean parameters of the mixture components are predicted using the right-hand side of *formula*. All types of predictor terms allowed in non-mixture models are allowed in mixture models as well.

Distributional parameters of mixture distributions have the same name as those of the corresponding ordinary distributions, but with a number at the end to indicate the mixture component. For instance, if you use family *mixture(gaussian, gaussian)*, the distributional parameters are *sigma1* and *sigma2*. Distributional parameters of the same class can be fixed to the same value. For the above example, we could write *sigma2 = "sigma1"* to make sure that both components have the same residual standard deviation, which is in turn estimated from the data.

In addition, there are two types of special distributional parameters. The first are named *mu<ID>*, that allow for modeling different predictors for the mean parameters of different mixture components. For instance, if you want to predict the mean of the first component using predictor *x* and the mean of the second component using predictor *z*, you can write *mu1 ~ x* as well as *mu2 ~ z*. The second are named *theta<ID>*, which constitute the mixing proportions. If the mixing proportions are fixed to certain values, they are internally normalized to form a probability vector. If one seeks to predict the mixing proportions, all but one of the them has to be predicted, while the remaining one is used as the reference category to identify the model. The so-called 'softmax' transformation is applied on the linear predictor terms to form a probability vector.

For more information on mixture models, see the documentation of [mixture](#).

Formula syntax for multivariate models

Multivariate models may be specified using *mvbind* notation or with help of the [mvbf](#) function. Suppose that *y1* and *y2* are response variables and *x* is a predictor. Then *mvbind(y1, y2) ~ x* specifies a multivariate model. The effects of all terms specified at the RHS of the formula are assumed to vary across response variables. For instance, two parameters will be estimated for *x*, one for the effect on *y1* and another for the effect on *y2*. This is also true for group-level effects. When writing, for instance, *mvbind(y1, y2) ~ x + (1+x|g)*, group-level effects will be estimated separately for each response. To model these effects as correlated across responses, use the ID syntax (see above). For the present example, this would look as follows: *mvbind(y1, y2) ~ x + (1+x|2|g)*. Of course, you could also use any value other than 2 as ID.

It is also possible to specify different formulas for different responses. If, for instance, y_1 should be predicted by x and y_2 should be predicted by z , we could write `mvbf(y1 ~ x, y2 ~ z)`. Alternatively, multiple `brmsformula` objects can be added to specify a joint multivariate model (see 'Examples').

Value

An object of class `brmsformula`, which is essentially a list containing all model formulas as well as some additional information.

See Also

[mvbrmsformula](#), [brmsformula-helpers](#)

Examples

```
# multilevel model with smoothing terms
brmsformula(y ~ x1*x2 + s(z) + (1+x1|1) + (1|g2))

# additionally predict 'sigma'
brmsformula(y ~ x1*x2 + s(z) + (1+x1|1) + (1|g2),
            sigma ~ x1 + (1|g2))

# use the shorter alias 'bf'
(formula1 <- brmsformula(y ~ x + (x|g)))
(formula2 <- bf(y ~ x + (x|g)))
# will be TRUE
identical(formula1, formula2)

# incorporate censoring
bf(y | cens(censor_variable) ~ predictors)

# define a simple non-linear model
bf(y ~ a1 - a2^x, a1 + a2 ~ 1, nl = TRUE)

# predict a1 and a2 differently
bf(y ~ a1 - a2^x, a1 ~ 1, a2 ~ x + (x|g), nl = TRUE)

# correlated group-level effects across parameters
bf(y ~ a1 - a2^x, a1 ~ 1 + (1 |2| g), a2 ~ x + (x |2| g), nl = TRUE)
# alternative but equivalent way to specify the above model
bf(y ~ a1 - a2^x, a1 ~ 1 + (1 | gr(g, id = 2)),
   a2 ~ x + (x | gr(g, id = 2)), nl = TRUE)

# define a multivariate model
bf(mvbind(y1, y2) ~ x * z + (1|g))

# define a zero-inflated model
# also predicting the zero-inflation part
bf(y ~ x * z + (1+x|ID1|g), zi ~ x + (1|ID1|g))

# specify a predictor as monotonic
bf(y ~ mo(x) + more_predictors)
```

```

# for ordinal models only
# specify a predictor as category specific
bf(y ~ cs(x) + more_predictors)
# add a category specific group-level intercept
bf(y ~ cs(x) + (cs(1)|g))
# specify parameter 'disc'
bf(y ~ person + item, disc ~ item)

# specify variables containing measurement error
bf(y ~ me(x, sdx))

# specify predictors on all parameters of the wiener diffusion model
# the main formula models the drift rate 'delta'
bf(rt | dec(decision) ~ x, bs ~ x, ndt ~ x, bias ~ x)

# fix the bias parameter to 0.5
bf(rt | dec(decision) ~ x, bias = 0.5)

# specify different predictors for different mixture components
mix <- mixture(gaussian, gaussian)
bf(y ~ 1, mu1 ~ x, mu2 ~ z, family = mix)

# fix both residual standard deviations to the same value
bf(y ~ x, sigma2 = "sigma1", family = mix)

# use the '+' operator to specify models
bf(y ~ 1) +
  nlf(sigma ~ a * exp(b * x), a ~ x) +
  lf(b ~ z + (1|g), dpar = "sigma") +
  gaussian()

# specify a multivariate model using the '+' operator
bf(y1 ~ x + (1|g)) +
  gaussian() + cor_ar(~1|g) +
  bf(y2 ~ z) + poisson()

# specify correlated residuals of a gaussian and a poisson model
form1 <- bf(y1 ~ 1 + x + (1|c|obs), sigma = 1) + gaussian()
form2 <- bf(y2 ~ 1 + x + (1|c|obs)) + poisson()

# model missing values in predictors
bf(bmi ~ age * mi(chl)) +
  bf(chl | mi() ~ age) +
  set_rescor(FALSE)

# model sigma as a function of the mean
bf(y ~ eta, nl = TRUE) +
  lf(eta ~ 1 + x) +
  nlf(sigma ~ tau * sqrt(eta)) +
  lf(tau ~ 1)

```

Description

Helper functions to specify linear and non-linear formulas for use with `brmsformula`.

Usage

```
nlf(formula, ..., flist = NULL, dpar = NULL, resp = NULL, loop = NULL)

lf(
  ...,
  flist = NULL,
  dpar = NULL,
  resp = NULL,
  center = NULL,
  cmc = NULL,
  sparse = NULL,
  decomp = NULL
)

acformula(autocor, resp = NULL)

set_nl(nl = TRUE, dpar = NULL, resp = NULL)

set_rescor(rescor = TRUE)

set_mecor(mecor = TRUE)
```

Arguments

<code>formula</code>	Non-linear formula for a distributional parameter. The name of the distributional parameter can either be specified on the left-hand side of <code>formula</code> or via argument <code>dpar</code> .
<code>...</code>	Additional <code>formula</code> objects to specify predictors of non-linear and distributional parameters. Formulas can either be named directly or contain names on their left-hand side. Alternatively, it is possible to fix parameters to certain values by passing numbers or character strings in which case arguments have to be named to provide the parameter names. See 'Details' for more information.
<code>flist</code>	Optional list of formulas, which are treated in the same way as formulas passed via the <code>...</code> argument.
<code>dpar</code>	Optional character string specifying the distributional parameter to which the formulas passed via <code>...</code> and <code>flist</code> belong.
<code>resp</code>	Optional character string specifying the response variable to which the formulas passed via <code>...</code> and <code>flist</code> belong. Only relevant in multivariate models.

loop	Logical; Only used in non-linear models. Indicates if the computation of the non-linear formula should be done inside (TRUE) or outside (FALSE) a loop over observations. Defaults to TRUE.
center	Logical; Indicates if the population-level design matrix should be centered, which usually increases sampling efficiency. See the 'Details' section for more information. Defaults to TRUE for distributional parameters and to FALSE for non-linear parameters.
cmc	Logical; Indicates whether automatic cell-mean coding should be enabled when removing the intercept by adding 0 to the right-hand of model formulas. Defaults to TRUE to mirror the behavior of standard R formula parsing.
sparse	Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to FALSE). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased.
decomp	Optional name of the decomposition used for the population-level design matrix. Defaults to NULL that is no decomposition. Other options currently available are "QR" for the QR decomposition that helps in fitting models with highly correlated predictors.
autocor	A one sided formula containing autocorrelation terms. All none autocorrelation terms in autocor will be silently ignored.
nl	Logical; Indicates whether formula should be treated as specifying a non-linear model. By default, formula is treated as an ordinary linear model formula.
rescor	Logical; Indicates if residual correlation between the response variables should be modeled. Currently this is only possible in multivariate gaussian and student models. Only relevant in multivariate models.
mecor	Logical; Indicates if correlations between latent variables defined by me terms should be modeled. Defaults to TRUE.

Value

For lf and nlf a list that can be passed to [brmsformula](#) or added to an existing brmsformula or mvbrmsformula object. For set_nl and set_rescor a logical value that can be added to an existing brmsformula or mvbrmsformula object.

See Also

[brmsformula](#), [mvbrmsformula](#)

Examples

```
# add more formulas to the model
bf(y ~ 1) +
  nlf(sigma ~ a * exp(b * x)) +
  lf(a ~ x, b ~ z + (1|g)) +
  gaussian()

# specify 'nl' later on
```

```

bf(y ~ a * inv_logit(x * b)) +
  lf(a + b ~ z) +
  set_nl(TRUE)

# specify a multivariate model
bf(y1 ~ x + (1|g)) +
  bf(y2 ~ z) +
  set_rescor(TRUE)

# add autocorrelation terms
bf(y ~ x) + acformula(~ arma(p = 1, q = 1) + car(W))

```

brmshypothesis	<i>Descriptions of brmshypothesis Objects</i>
----------------	---

Description

A `brmshypothesis` object contains posterior draws as well as summary statistics of non-linear hypotheses as returned by [hypothesis](#).

Usage

```

## S3 method for class 'brmshypothesis'
print(
  x,
  digits = 2,
  chars = 20,
  short = getOption("brms.short_summary", FALSE),
  ...
)

## S3 method for class 'brmshypothesis'
plot(
  x,
  nvariables = 5,
  N = NULL,
  ignore_prior = FALSE,
  chars = 40,
  colors = NULL,
  theme = NULL,
  ask = TRUE,
  plot = TRUE,
  ...
)

```

Arguments

x	An object of class <code>brmsfit</code> .
digits	Minimal number of significant digits, see print.default .
chars	Maximum number of characters of each hypothesis to print or plot. If NULL, print the full hypotheses. Defaults to 20.
short	A flag indicating whether to provide a shorter summary with less informational text. Defaults to FALSE. Can be set globally for the current session via the <code>brms.short_summary</code> option.
...	Currently ignored.
nvariables	The number of variables (parameters) plotted per page.
N	Deprecated alias of <code>nvariables</code> .
ignore_prior	A flag indicating if prior distributions should also be plotted. Only used if priors were specified on the relevant parameters.
colors	Two values specifying the colors of the posterior and prior density respectively. If NULL (the default) colors are taken from the current color scheme of the <code>bayesplot</code> package.
theme	A <code>theme</code> object modifying the appearance of the plots. For some basic themes see <code>ggtheme</code> and <code>theme_default</code> .
ask	Logical; indicates if the user is prompted before a new page is plotted. Only used if <code>plot</code> is TRUE.
plot	Logical; indicates if plots should be plotted directly in the active graphic device. Defaults to TRUE.

Details

The two most important elements of a `brmshypothesis` object are `hypothesis`, which is a `data.frame` containing the summary estimates of the hypotheses, and `samples`, which is a `data.frame` containing the corresponding posterior draws.

See Also

[hypothesis](#)

Description

Parse formulas objects for use in **brms**.

Usage

```
brmsterms(formula, ...)

## Default S3 method:
brmsterms(formula, ...)

## S3 method for class 'brmsformula'
brmsterms(formula, check_response = TRUE, resp_rhs_all = TRUE, ...)

## S3 method for class 'mvbrmsformula'
brmsterms(formula, ...)
```

Arguments

formula	An object of class formula , brmsformula , or mvbrmsformula (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in brmsformula .
...	Further arguments passed to or from other methods.
check_response	Logical; Indicates whether the left-hand side of formula (i.e. response variables and addition arguments) should be parsed. If FALSE, formula may also be one-sided.
resp_rhs_all	Logical; Indicates whether to also include response variables on the right-hand side of formula .\$allvars, where . represents the output of brmsterms.

Details

This is the main formula parsing function of [brms](#). It should usually not be called directly, but is exported to allow package developers making use of the formula syntax implemented in [brms](#). As long as no other packages depend on this functions, it may be changed without deprecation warnings, when new features make this necessary.

Value

An object of class [brmsterms](#) or [mvbrmsterms](#) (for multivariate models), which is a [list](#) containing all required information initially stored in formula in an easier to use format, basically a list of formulas (not an abstract syntax tree).

See Also

[brm](#), [brmsformula](#), [mvbrmsformula](#)

brm_multiple	<i>Run the same brms model on multiple datasets</i>
--------------	---

Description

Run the same **brms** model on multiple datasets and then combine the results into one fitted model object. This is useful in particular for multiple missing value imputation, where the same model is fitted on multiple imputed data sets. Models can be run in parallel using the **future** package.

Usage

```
brm_multiple(
  formula,
  data,
  family = gaussian(),
  prior = NULL,
  data2 = NULL,
  autocor = NULL,
  cov_ranef = NULL,
  sample_prior = c("no", "yes", "only"),
  sparse = NULL,
  knots = NULL,
  stanvars = NULL,
  stan_funs = NULL,
  silent = 1,
  recompile = FALSE,
  combine = TRUE,
  fit = NA,
  algorithm = getOption("brms.algorithm", "sampling"),
  seed = NA,
  file = NULL,
  file_compress = TRUE,
  file_refit = getOption("brms.file_refit", "never"),
  ...
)
```

Arguments

formula	An object of class formula , brmsformula , or mvbrmsformula (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in brmsformula .
data	A <i>list</i> of <code>data.frames</code> each of which will be used to fit a separate model. Alternatively, a <code>mids</code> object from the mice package.
family	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to

	specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see brmsfamily . By default, a linear gaussian model is applied. In multivariate models, family might also be a list of families.
prior	One or more <code>brmsprior</code> objects created by set_prior or related functions and combined using the <code>c</code> method or the <code>+</code> operator. See also default_prior for more help.
data2	A <i>list</i> of named lists each of which will be used to fit a separate model. Each of the named lists contains objects representing data which cannot be passed via argument <code>data</code> (see <code>brm</code> for examples). The length of the outer list should match the length of the list passed to the <code>data</code> argument.
autocor	(Deprecated) An optional <code>cor_brms</code> object describing the correlation structure within the response variable (i.e., the 'autocorrelation'). See the documentation of <code>cor_brms</code> for a description of the available correlation structures. Defaults to <code>NULL</code> , corresponding to no correlations. In multivariate models, <code>autocor</code> might also be a list of autocorrelation structures. It is now recommend to specify autocorrelation terms directly within <code>formula</code> . See brmsformula for more details.
cov_ranef	(Deprecated) A list of matrices that are proportional to the (within) covariance structure of the group-level effects. The names of the matrices should correspond to columns in <code>data</code> that are used as grouping factors. All levels of the grouping factor should appear as rownames of the corresponding matrix. This argument can be used, among others to model pedigrees and phylogenetic effects. It is now recommended to specify those matrices in the formula interface using the <code>gr</code> and related functions. See <code>vignette("brms_phylogenetics")</code> for more details.
sample_prior	Indicate if draws from priors should be drawn additionally to the posterior draws. Options are "no" (the default), "yes", and "only". Among others, these draws can be used to calculate Bayes factors for point hypotheses via hypothesis . Please note that improper priors are not sampled, including the default improper priors used by <code>brm</code> . See set_prior on how to set (proper) priors. Please also note that prior draws for the overall intercept are not obtained by default for technical reasons. See brmsformula how to obtain prior draws for the intercept. If <code>sample_prior</code> is set to "only", draws are drawn solely from the priors ignoring the likelihood, which allows among others to generate draws from the prior predictive distribution. In this case, all parameters must have proper priors.
sparse	(Deprecated) Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to <code>FALSE</code>). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased. It is now recommended to use the <code>sparse</code> argument of brmsformula and related functions.
knots	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See gamm for more details.
stanvars	An optional <code>stanvars</code> object generated by function stanvar to define additional variables for use in Stan 's program blocks.
stan_funs	(Deprecated) An optional character string containing self-defined Stan functions, which will be included in the functions block of the generated Stan code. It is now recommended to use the <code>stanvars</code> argument for this purpose instead.

<code>silent</code>	Verbosity level between 0 and 2. If 1 (the default), most of the informational messages of compiler and sampler are suppressed. If 2, even more messages are suppressed. The actual sampling progress is still printed. Set <code>refresh = 0</code> to turn this off as well. If using <code>backend = "rstan"</code> you can also set <code>open_progress = FALSE</code> to prevent opening additional progress bars.
<code>recompile</code>	Logical, indicating whether the Stan model should be recompiled for every imputed data set. Defaults to FALSE. If NULL, <code>brm_multiple</code> tries to figure out internally, if recompilation is necessary, for example because data-dependent priors have changed. Using the default of no recompilation should be fine in most cases.
<code>combine</code>	Logical; Indicates if the fitted models should be combined into a single fitted model object via combine_models . Defaults to TRUE.
<code>fit</code>	An instance of S3 class <code>brmsfit_multiple</code> derived from a previous fit; defaults to NA. If <code>fit</code> is of class <code>brmsfit_multiple</code> , the compiled model associated with the fitted result is re-used and all arguments modifying the model code or data are ignored. It is not recommended to use this argument directly, but to call the update method, instead.
<code>algorithm</code>	Character string naming the estimation approach to use. Options are "sampling" for MCMC (the default), "meanfield" for variational inference with independent normal distributions, "fullrank" for variational inference with a multivariate normal distribution, "pathfinder" for the pathfinder algorithm, "laplace" for the laplace approximation, or "fixed_param" for sampling from fixed parameter values. Can be set globally for the current R session via the "brms.algorithm" option (see options).
<code>seed</code>	The seed for random number generation to make results reproducible. If NA (the default), Stan will set the seed randomly.
<code>file</code>	Either NULL or a character string. In the latter case, the fitted model object is saved via saveRDS in a file named after the string supplied in <code>file</code> . The .rds extension is added automatically. If the file already exists, <code>brm</code> will load and return the saved model object instead of refitting the model. Unless you specify the <code>file_refit</code> argument as well, the existing files won't be overwritten, you have to manually remove the file in order to refit and save the model under an existing file name. The file name is stored in the <code>brmsfit</code> object for later usage.
<code>file_compress</code>	Logical or a character string, specifying one of the compression algorithms supported by saveRDS . If the <code>file</code> argument is provided, this compression will be used when saving the fitted model object.
<code>file_refit</code>	Modifies when the fit stored via the <code>file</code> argument is re-used. Can be set globally for the current R session via the "brms.file_refit" option (see options). For "never" (default) the fit is always loaded if it exists and fitting is skipped. For "always" the model is always refitted. If set to "on_change", brms will refit the model if model, data or algorithm as passed to Stan differ from what is stored in the file. This also covers changes in priors, <code>sample_prior</code> , <code>stanvars</code> , covariance structure, etc. If you believe there was a false positive, you can use <code>brmsfit_needs_refit</code> to see why refit is deemed necessary. Refit will not be triggered for changes in additional parameters of the fit (e.g., initial values, number of iterations, control arguments, ...). A known limitation is that a refit will be triggered if within-chain parallelization is switched on/off.

... Further arguments passed to [brm](#).

Details

The combined model may issue false positive convergence warnings, as the MCMC chains corresponding to different datasets may not necessarily overlap, even if each of the original models did converge. To find out whether each of the original models converged, subset the draws belonging to the individual models and then run convergence diagnostics. See Examples below for details.

Value

If `combine = TRUE` a `brmsfit_multiple` object, which inherits from class `brmsfit` and behaves essentially the same. If `combine = FALSE` a list of `brmsfit` objects.

Parallelization with multiple CPU cores

`brms` can make use of multiple CPU cores in parallel to speed up computations in various ways. For efficient use of the available resources it is recommended to only use parallelism to an extend such that the available physical CPUs are not oversubscribed. For example, when you have 8 CPU cores locally available, then you may consider to run 4 chains with 2 threads per chain for best performance if you happen to just run a single model. In case you run a simulation study which requires to run many times a given model, then neither chain nor within-chain parallelization is advisable as the computational resources are already exhausted by the simulation study and any further parallelization beyond the simulation study itself will in fact slow down the overall runtime. Please be aware that for historical reasons the nomenclature of the arguments is possibly confusing. The `cores` argument refers to running different chains in parallel and the within-chain parallelization will allocate for each chain as many threads as requested. The requested threads therefore increase the use of overall CPUs in a multiplicative way.

For more advanced parallelization (including beyond single model fits), `brms` also integrates with the `future` package. Importantly, this enables seamless integration with the `mirai` parallelization framework through the use of the `future.mirai` adapter. With `mirai` local and remote machines can be used in a fully transparent manner to the user. This includes the possibility to use large number of remote machines running in the context of a computer cluster, which are managed with queuing systems. Please refer to the section on distributed computing of [mirai::daemons](#).

Examples

```
## Not run:
library(mice)
m <- 5
imp <- mice(nhanes2, m = m)

# fit the model using mice and lm
fit_imp1 <- with(lm(bmi ~ age + hyp + chl), data = imp)
summary(pool(fit_imp1))

# fit the model using brms
fit_imp2 <- brm_multiple(bmi ~ age + hyp + chl, data = imp, chains = 1)
summary(fit_imp2)
plot(fit_imp2, variable = "b_",
     regex = TRUE)
```

```
# investigate convergence of the original models
library(posterior)
draws <- as_draws_array(fit_imp2)
# every dataset has just one chain here
draws_per_dat <- lapply(1:m, \((i) subset_draws(draws, chain = i)))
lapply(draws_per_dat, summarise_draws, default_convergence_measures())

# use the future package for parallelization
library(future)
plan(multisession, workers = 4)
fit_imp3 <- brm_multiple(bmi ~ age + hyp + chl, data = imp, chains = 1)
summary(fit_imp3)

## End(Not run)
```

car*Spatial conditional autoregressive (CAR) structures***Description**

Set up an spatial conditional autoregressive (CAR) term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with CAR terms.

Usage

```
car(M, gr = NA, type = "escar")
```

Arguments

M	Adjacency matrix of locations. All non-zero entries are treated as if the two locations are adjacent. If gr is specified, the row names of M have to match the levels of the grouping factor.
gr	An optional grouping factor mapping observations to spatial locations. If not specified, each observation is treated as a separate location. It is recommended to always specify a grouping factor to allow for handling of new data in post-processing methods.
type	Type of the CAR structure. Currently implemented are "escar" (exact sparse CAR), "esicar" (exact sparse intrinsic CAR), "icar" (intrinsic CAR), and "bym2". More information is provided in the 'Details' section.

Details

The escar and esicar types are implemented based on the case study of Max Joseph (<https://github.com/mbjoseph/CARstan>). The icar and bym2 type is implemented based on the case study of Mitzi Morris (https://mc-stan.org/users/documentation/case-studies/icar_stan.html).

Value

An object of class 'car_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#)

Examples

```
## Not run:
# generate some spatial data
east <- north <- 1:10
Grid <- expand.grid(east, north)
K <- nrow(Grid)

# set up distance and neighbourhood matrices
distance <- as.matrix(dist(Grid))
W <- array(0, c(K, K))
W[distance == 1] <- 1
rownames(W) <- 1:nrow(W)

# generate the covariates and response data
x1 <- rnorm(K)
x2 <- rnorm(K)
theta <- rnorm(K, sd = 0.05)
phi <- rmulti_normal(
  1, mu = rep(0, K), Sigma = 0.4 * exp(-0.1 * distance)
)
eta <- x1 + x2 + phi
prob <- exp(eta) / (1 + exp(eta))
size <- rep(50, K)
y <- rbinom(n = K, size = size, prob = prob)
g <- 1:length(y)
dat <- data.frame(y, size, x1, x2, g)

# fit a CAR model
fit <- brm(y | trials(size) ~ x1 + x2 + car(W, gr = g),
            data = dat, data2 = list(W = W),
            family = binomial())
summary(fit)

## End(Not run)
```

Description

Extract model coefficients, which are the sum of population-level effects and corresponding group-level effects

Usage

```
## S3 method for class 'brmsfit'
coef(object, summary = TRUE, robust = FALSE, probs = c(0.025, 0.975), ...)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>summary</code>	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .
<code>robust</code>	If <code>FALSE</code> (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If <code>TRUE</code> , the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is <code>TRUE</code> .
<code>probs</code>	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is <code>TRUE</code> .
<code>...</code>	Further arguments passed to <code>fixef.brmsfit</code> and <code>ranef.brmsfit</code> .

Value

A list of 3D arrays (one per grouping factor). If `summary` is `TRUE`, the 1st dimension contains the factor levels, the 2nd dimension contains the summary statistics (see `posterior_summary`), and the 3rd dimension contains the group-level effects. If `summary` is `FALSE`, the 1st dimension contains the posterior draws, the 2nd dimension contains the factor levels, and the 3rd dimension contains the group-level effects.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt + (1+Trt|visit),
            data = epilepsy, family = gaussian(), chains = 2)
## extract population and group-level coefficients separately
fixef(fit)
ranef(fit)
## extract combined coefficients
coef(fit)

## End(Not run)
```

combine_models*Combine Models fitted with brms*

Description

Combine multiple `brmsfit` objects, which fitted the same model. This is usefully for instance when having manually run models in parallel.

Usage

```
combine_models(..., mlist = NULL, check_data = TRUE)
```

Arguments

...	One or more <code>brmsfit</code> objects.
<code>mlist</code>	Optional list of one or more <code>brmsfit</code> objects.
<code>check_data</code>	Logical; indicates if the data should be checked for being the same across models (defaults to <code>TRUE</code>). Setting it to <code>FALSE</code> may be useful for instance when combining models fitted on multiple imputed data sets.

Details

This function just takes the first model and replaces its `stanfit` object (slot `fit`) by the combined `stanfit` objects of all models.

Value

A `brmsfit` object.

compare_ic*Compare Information Criteria of Different Models*

Description

Compare information criteria of different models fitted with `waic` or `loo`. Deprecated and will be removed in the future. Please use `loo_compare` instead.

Usage

```
compare_ic(..., x = NULL, ic = c("loo", "waic", "kfold"))
```

Arguments

- ... At least two objects returned by [waic](#) or [loo](#). Alternatively, [brmsfit](#) objects with information criteria precomputed via [add_ic](#) may be passed, as well.
- x A list containing the same types of objects as can be passed via ...
- ic The name of the information criterion to be extracted from [brmsfit](#) objects. Ignored if information criterion objects are only passed directly.

Details

See [loo_compare](#) for the recommended way of comparing models with the [loo](#) package.

Value

An object of class [iclist](#).

See Also

[loo](#), [loo_compare](#) [add_criterion](#)

Examples

```
## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler)
waic1 <- waic(fit1)

# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler)
waic2 <- waic(fit2)

# compare both models
compare_ic(waic1, waic2)

## End(Not run)
```

Description

Display conditional effects of one or more numeric and/or categorical predictors including two-way interaction effects.

Usage

```
## S3 method for class 'brmsfit'
conditional_effects(
  x,
  effects = NULL,
  conditions = NULL,
  int_conditions = NULL,
  re_formula = NA,
  prob = 0.95,
  robust = TRUE,
  method = "posterior_epred",
  spaghetti = FALSE,
  surface = FALSE,
  categorical = FALSE,
  ordinal = FALSE,
  transform = NULL,
  resolution = 100,
  select_points = 0,
  too_far = 0,
  probs = NULL,
  ...
)

conditional_effects(x, ...)

## S3 method for class 'brms_conditional_effects'
plot(
  x,
  ncol = NULL,
  points = getOption("brms.plot_points", FALSE),
  rug = getOption("brms.plot_rug", FALSE),
  mean = TRUE,
  jitter_width = 0,
  stype = c("contour", "raster"),
  line_args = list(),
  cat_args = list(),
  errorbar_args = list(),
  surface_args = list(),
  spaghetti_args = list(),
  point_args = list(),
  rug_args = list(),
  facet_args = list(),
  theme = NULL,
  ask = TRUE,
  plot = TRUE,
  ...
)
```

Arguments

x	An object of class <code>brmsfit</code> .
effects	An optional character vector naming effects (main effects or interactions) for which to compute conditional plots. Interactions are specified by a <code>:</code> between variable names. If <code>NULL</code> (the default), plots are generated for all main effects and two-way interactions estimated in the model. When specifying <code>effects</code> manually, <i>all</i> two-way interactions (including grouping variables) may be plotted even if not originally modeled.
conditions	An optional <code>data.frame</code> containing variable values to condition on. Each effect defined in <code>effects</code> will be plotted separately for each row of <code>conditions</code> . Values in the <code>cond_</code> column will be used as titles of the subplots. If <code>cond_</code> is not given, the row names will be used for this purpose instead. It is recommended to only define a few rows in order to keep the plots clear. See make_conditions for an easy way to define conditions. If <code>NULL</code> (the default), numeric variables will be conditionalized by using their means and factors will get their first level assigned. NA values within factors are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding.
int_conditions	An optional named list whose elements are vectors of values of the variables specified in <code>effects</code> . At these values, predictions are evaluated. The names of <code>int_conditions</code> have to match the variable names exactly. Additionally, the elements of the vectors may be named themselves, in which case their names appear as labels for the conditions in the plots. Instead of vectors, functions returning vectors may be passed and are applied on the original values of the corresponding variable. If <code>NULL</code> (the default), predictions are evaluated at the <code>mean</code> and at <code>mean + / - sd</code> for numeric predictors and at all categories for factor-like predictors.
re_formula	A formula containing group-level effects to be considered in the conditional predictions. If <code>NULL</code> , include all group-level effects; if <code>NA</code> (default), include no group-level effects.
prob	A value between 0 and 1 indicating the desired probability to be covered by the uncertainty intervals. The default is 0.95.
robust	If <code>TRUE</code> (the default) the median is used as the measure of central tendency. If <code>FALSE</code> the mean is used instead.
method	Method used to obtain predictions. Can be set to <code>"posterior_epred"</code> (the default), <code>"posterior_predict"</code> , or <code>"posterior_linpred"</code> . For more details, see the respective function documentations.
spaghetti	Logical. Indicates if predictions should be visualized via spaghetti plots. Only applied for numeric predictors. If <code>TRUE</code> , it is recommended to set argument <code>ndraws</code> to a relatively small value (e.g., 100) in order to reduce computation time.
surface	Logical. Indicates if interactions or two-dimensional smooths should be visualized as a surface. Defaults to <code>FALSE</code> . The surface type can be controlled via argument <code>stype</code> of the related plotting method.
categorical	Logical. Indicates if effects of categorical or ordinal models should be shown in terms of probabilities of response categories. Defaults to <code>FALSE</code> .

<code>ordinal</code>	(Deprecated) Please use argument <code>categorical</code> . Logical. Indicates if effects in ordinal models should be visualized as a raster with the response categories on the y-axis. Defaults to FALSE.
<code>transform</code>	A function or a character string naming a function to be applied on the predicted responses before summary statistics are computed. Only allowed if <code>method = "posterior_predict"</code> .
<code>resolution</code>	Number of support points used to generate the plots. Higher resolution leads to smoother plots. Defaults to 100. If <code>surface</code> is TRUE, this implies 10000 support points for interaction terms, so it might be necessary to reduce <code>resolution</code> when only few RAM is available.
<code>select_points</code>	Positive number. Only relevant if <code>points</code> or <code>rug</code> are set to TRUE: Actual data points of numeric variables that are too far away from the values specified in <code>conditions</code> can be excluded from the plot. Values are scaled into the unit interval and then points more than <code>select_points</code> from the values in <code>conditions</code> are excluded. By default, all points are used.
<code>too_far</code>	Positive number. For surface plots only: Grid points that are too far away from the actual data points can be excluded from the plot. <code>too_far</code> determines what is too far. The grid is scaled into the unit square and then grid points more than <code>too_far</code> from the predictor variables are excluded. By default, all grid points are used. Ignored for non-surface plots.
<code>probs</code>	(Deprecated) The quantiles to be used in the computation of uncertainty intervals. Please use argument <code>prob</code> instead.
<code>...</code>	Further arguments such as <code>draw_ids</code> or <code>ndraws</code> passed to <code>posterior_predict</code> or <code>posterior_epred</code> .
<code>ncol</code>	Number of plots to display per column for each effect. If <code>NULL</code> (default), <code>ncol</code> is computed internally based on the number of rows of <code>conditions</code> .
<code>points</code>	Logical. Indicates if the original data points should be added via <code>geom_jitter</code> . Default is FALSE. Can be controlled globally via the <code>brms.plot_points</code> option. Note that only those data points will be added that match the specified conditions defined in <code>conditions</code> . For categorical predictors, the conditions have to match exactly. For numeric predictors, argument <code>select_points</code> is used to determine, which points do match a condition.
<code>rug</code>	Logical. Indicates if a rug representation of predictor values should be added via <code>geom_rug</code> . Default is FALSE. Depends on <code>select_points</code> in the same way as <code>points</code> does. Can be controlled globally via the <code>brms.plot_rug</code> option.
<code>mean</code>	Logical. Only relevant for spaghetti plots. If TRUE (the default), display the mean regression line on top of the regression lines for each sample.
<code>jitter_width</code>	Only used if <code>points = TRUE</code> : Amount of horizontal jittering of the data points. Mainly useful for ordinal models. Defaults to 0 that is no jittering.
<code>stype</code>	Indicates how surface plots should be displayed. Either "contour" or "raster".
<code>line_args</code>	Only used in plots of continuous predictors: A named list of arguments passed to <code>geom_smooth</code> .
<code>cat_args</code>	Only used in plots of categorical predictors: A named list of arguments passed to <code>geom_point</code> .

errorbar_args	Only used in plots of categorical predictors: A named list of arguments passed to <code>geom_errorbar</code> .
surface_args	Only used in surface plots: A named list of arguments passed to <code>geom_contour</code> or <code>geom_raster</code> (depending on argument <code>stype</code>).
spaghetti_args	Only used in spaghetti plots: A named list of arguments passed to <code>geom_smooth</code> .
point_args	Only used if <code>points</code> = TRUE: A named list of arguments passed to <code>geom_jitter</code> .
rug_args	Only used if <code>rug</code> = TRUE: A named list of arguments passed to <code>geom_rug</code> .
facet_args	Only used if if multiple conditions are provided: A named list of arguments passed to <code>facet_wrap</code> .
theme	A <code>theme</code> object modifying the appearance of the plots. For some basic themes see <code>ggtheme</code> and <code>theme_default</code> .
ask	Logical; indicates if the user is prompted before a new page is plotted. Only used if <code>plot</code> is TRUE.
plot	Logical; indicates if plots should be plotted directly in the active graphic device. Defaults to TRUE.

Details

When creating `conditional_effects` for a particular predictor (or interaction of two predictors), one has to choose the values of all other predictors to condition on. By default, the mean is used for continuous variables and the reference category is used for factors, but you may change these values via argument `conditions`. This also has an implication for the `points` argument: In the created plots, only those points will be shown that correspond to the factor levels actually used in the conditioning, in order not to create the false impression of bad model fit, where it is just due to conditioning on certain factor levels.

To fully change colors of the created plots, one has to amend both `scale_colour` and `scale_fill`. See `scale_colour_grey` or `scale_colour_gradient` for more details.

Value

An object of class 'brms_conditional_effects' which is a named list with one `data.frame` per effect containing all information required to generate conditional effects plots. Among others, these `data.frames` contain some special variables, namely `estimate_` (predicted values of the response), `se_` (standard error of the predicted response), `lower_` and `upper_` (lower and upper bounds of the uncertainty interval of the response), as well as `cond_` (used in faceting when `conditions` contains multiple rows).

The corresponding `plot` method returns a named list of `ggplot` objects, which can be further customized using the `ggplot2` package.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt + (1 | patient),
            data = epilepsy, family = poisson())

## plot all conditional effects
```

```

plot(conditional_effects(fit), ask = FALSE)

## change colours to grey scale
library(ggplot2)
ce <- conditional_effects(fit, "zBase:Trt")
plot(ce, plot = FALSE)[[1]] +
  scale_color_grey() +
  scale_fill_grey()

## only plot the conditional interaction effect of 'zBase:Trt'
## for different values for 'zAge'
conditions <- data.frame(zAge = c(-1, 0, 1))
plot(conditional_effects(fit, effects = "zBase:Trt",
                        conditions = conditions))

## also incorporate group-level effects variance over patients
## also add data points and a rug representation of predictor values
plot(conditional_effects(fit, effects = "zBase:Trt",
                        conditions = conditions, re_formula = NULL),
     points = TRUE, rug = TRUE)

## change handling of two-way interactions
int_conditions <- list(
  zBase = setNames(c(-2, 1, 0), c("b", "c", "a"))
)
conditional_effects(fit, effects = "Trt:zBase",
                    int_conditions = int_conditions)
conditional_effects(fit, effects = "Trt:zBase",
                    int_conditions = list(zBase = quantile))

## fit a model to illustrate how to plot 3-way interactions
fit3way <- brm(count ~ zAge * zBase * Trt, data = epilepsy)
conditions <- make_conditions(fit3way, "zAge")
conditional_effects(fit3way, "zBase:Trt", conditions = conditions)
## only include points close to the specified values of zAge
ce <- conditional_effects(
  fit3way, "zBase:Trt", conditions = conditions,
  select_points = 0.1
)
plot(ce, points = TRUE)

## End(Not run)

```

conditional_smooths.brmsfit

*Display Smooth Terms***Description**

Display smooth s and t2 terms of models fitted with **brms**.

Usage

```
## S3 method for class 'brmsfit'
conditional_smooths(
  x,
  smooths = NULL,
  int_conditions = NULL,
  prob = 0.95,
  spaghetti = FALSE,
  surface = TRUE,
  resolution = 100,
  too_far = 0,
  ndraws = NULL,
  draw_ids = NULL,
  nsamples = NULL,
  subset = NULL,
  probs = NULL,
  ...
)
conditional_smooths(x, ...)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code> .
<code>smooths</code>	Optional character vector of smooth terms to display. If <code>NULL</code> (the default) all smooth terms are shown.
<code>int_conditions</code>	An optional named <code>list</code> whose elements are vectors of values of the variables specified in <code>effects</code> . At these values, predictions are evaluated. The names of <code>int_conditions</code> have to match the variable names exactly. Additionally, the elements of the vectors may be named themselves, in which case their names appear as labels for the conditions in the plots. Instead of vectors, functions returning vectors may be passed and are applied on the original values of the corresponding variable. If <code>NULL</code> (the default), predictions are evaluated at the <code>mean</code> and at <code>mean + / - sd</code> for numeric predictors and at all categories for factor-like predictors.
<code>prob</code>	A value between 0 and 1 indicating the desired probability to be covered by the uncertainty intervals. The default is 0.95.
<code>spaghetti</code>	Logical. Indicates if predictions should be visualized via spaghetti plots. Only applied for numeric predictors. If <code>TRUE</code> , it is recommended to set argument <code>ndraws</code> to a relatively small value (e.g., 100) in order to reduce computation time.
<code>surface</code>	Logical. Indicates if interactions or two-dimensional smooths should be visualized as a surface. Defaults to <code>TRUE</code> . The surface type can be controlled via argument <code>stype</code> of the related plotting method.
<code>resolution</code>	Number of support points used to generate the plots. Higher resolution leads to smoother plots. Defaults to 100. If <code>surface</code> is <code>TRUE</code> , this implies 10000 support

	points for interaction terms, so it might be necessary to reduce resolution when only few RAM is available.
too_far	Positive number. For surface plots only: Grid points that are too far away from the actual data points can be excluded from the plot. <code>too_far</code> determines what is too far. The grid is scaled into the unit square and then grid points more than <code>too_far</code> from the predictor variables are excluded. By default, all grid points are used. Ignored for non-surface plots.
ndraws	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
draw_ids	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
nsamples	Deprecated alias of <code>ndraws</code> .
subset	Deprecated alias of <code>draw_ids</code> .
probs	(Deprecated) The quantiles to be used in the computation of uncertainty intervals. Please use argument <code>prob</code> instead.
...	Currently ignored.

Details

Two-dimensional smooth terms will be visualized using either contour or raster plots.

Value

For the `brmsfit` method, an object of class `brms_conditional_effects`. See [conditional_effects](#) for more details and documentation of the related plotting function.

Examples

```
## Not run:
set.seed(0)
dat <- mgcv:::gamSim(1, n = 200, scale = 2)
fit <- brm(y ~ s(x0) + s(x1) + s(x2) + s(x3), data = dat)
# show all smooth terms
plot(conditional_smooths(fit), rug = TRUE, ask = FALSE)
# show only the smooth term s(x2)
plot(conditional_smooths(fit, smooths = "s(x2)'), ask = FALSE)

# fit and plot a two-dimensional smooth term
fit2 <- brm(y ~ t2(x0, x2), data = dat)
ms <- conditional_smooths(fit2)
plot(ms, stype = "contour")
plot(ms, stype = "raster")

## End(Not run)
```

constant*Constant priors in brms*

Description

Function used to set up constant priors in **brms**. The function does not evaluate its arguments – it exists purely to help set up the model.

Usage

```
constant(const, broadcast = TRUE)
```

Arguments

<code>const</code>	Numeric value, vector, matrix of values to which the parameters should be fixed to. Can also be a valid Stan variable in the model.
<code>broadcast</code>	Should <code>const</code> be automatically broadcasted to the correct size of the parameter? Defaults to <code>TRUE</code> . If you supply vectors or matrices in <code>const</code> or vector/matrix valued Stan variables, you need to set <code>broadcast</code> to <code>TRUE</code> (see Examples).

Value

A named list with elements `const` and `broadcast`.

See Also

[set_prior](#)

Examples

```
stancode(count ~ Base + Age, data = epilepsy,
         prior = prior(constant(1), class = "b"))

# will fail parsing because brms will try to broadcast a vector into a vector
stancode(count ~ Base + Age, data = epilepsy,
         prior = prior(constant(alpha), class = "b"),
         stanvars = stanvar(c(1, 0), name = "alpha"))

stancode(count ~ Base + Age, data = epilepsy,
         prior = prior(constant(alpha, broadcast = FALSE), class = "b"),
         stanvars = stanvar(c(1, 0), name = "alpha"))
```

<code>control_params</code>	<i>Extract Control Parameters of the NUTS Sampler</i>
-----------------------------	---

Description

Extract control parameters of the NUTS sampler such as `adapt_delta` or `max_treedepth`.

Usage

```
control_params(x, ...)
## S3 method for class 'brmsfit'
control_params(x, pars = NULL, ...)
```

Arguments

<code>x</code>	An R object
<code>...</code>	Currently ignored.
<code>pars</code>	Optional names of the control parameters to be returned. If <code>NULL</code> (the default) all control parameters are returned. See stan for more details.

Value

A named list with control parameter values.

<code>cor_ar</code>	<i>(Deprecated) AR(p) correlation structure</i>
---------------------	---

Description

This function is deprecated. Please see [ar](#) for the new syntax. This function is a constructor for the `cor_arma` class, allowing for autoregression terms only.

Usage

```
cor_ar(formula = ~1, p = 1, cov = FALSE)
```

Arguments

<code>formula</code>	A one sided formula of the form <code>~ t</code> , or <code>~ t g</code> , specifying a time covariate <code>t</code> and, optionally, a grouping factor <code>g</code> . A covariate for this correlation structure must be integer valued. When a grouping factor is present in <code>formula</code> , the correlation structure is assumed to apply only to observations within the same grouping level; observations with different grouping levels are assumed to be uncorrelated. Defaults to <code>~ 1</code> , which corresponds to using the order of the observations in the data as a covariate, and no groups.
----------------------	--

p	A non-negative integer specifying the autoregressive (AR) order of the ARMA structure. Default is 1.
cov	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If FALSE (the default) a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Details

AR refers to autoregressive effects of residuals, which is what is typically understood as autoregressive effects. However, one may also model autoregressive effects of the response variable, which is called ARR in **brms**.

Value

An object of class `cor_arma` containing solely autoregression terms.

See Also

[cor_arma](#)

Examples

```
cor_ar(~visit|patient, p = 2)
```

`cor_arma`

(Deprecated) ARMA(p,q) correlation structure

Description

This function is deprecated. Please see [arma](#) for the new syntax. This functions is a constructor for the `cor_arma` class, representing an autoregression-moving average correlation structure of order (p, q).

Usage

```
cor_arma(formula = ~1, p = 0, q = 0, r = 0, cov = FALSE)
```

Arguments

<code>formula</code>	A one sided formula of the form $\sim t$, or $\sim t \mid g$, specifying a time covariate t and, optionally, a grouping factor g . A covariate for this correlation structure must be integer valued. When a grouping factor is present in <code>formula</code> , the correlation structure is assumed to apply only to observations within the same grouping level; observations with different grouping levels are assumed to be uncorrelated. Defaults to ~ 1 , which corresponds to using the order of the observations in the data as a covariate, and no groups.
<code>p</code>	A non-negative integer specifying the autoregressive (AR) order of the ARMA structure. Default is 0.
<code>q</code>	A non-negative integer specifying the moving average (MA) order of the ARMA structure. Default is 0.
<code>r</code>	No longer supported.
<code>cov</code>	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If <code>FALSE</code> (the default) a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Value

An object of class `cor_arma`, representing an autoregression-moving-average correlation structure.

See Also

[cor_ar](#), [cor_ma](#)

Examples

```
cor_arma(~ visit | patient, p = 2, q = 2)
```

Description

Classes of correlation structures available in the **brms** package. `cor_brms` is not a correlation structure itself, but the class common to all correlation structures implemented in **brms**.

Available correlation structures

- cor_arma** autoregressive-moving average (ARMA) structure, with arbitrary orders for the autoregressive and moving average components
- cor_ar** autoregressive (AR) structure of arbitrary order
- cor_ma** moving average (MA) structure of arbitrary order
- cor_car** Spatial conditional autoregressive (CAR) structure
- cor_sar** Spatial simultaneous autoregressive (SAR) structure
- cor_fixed** fixed user-defined covariance structure

See Also

[cor_arma](#), [cor_ar](#), [cor_ma](#), [cor_car](#), [cor_sar](#), [cor_fixed](#)

cor_car

(Deprecated) Spatial conditional autoregressive (CAR) structures

Description

These function are deprecated. Please see [car](#) for the new syntax. These functions are constructors for the `cor_car` class implementing spatial conditional autoregressive structures.

Usage

```
cor_car(W, formula = ~1, type = "escar")
cor_icar(W, formula = ~1)
```

Arguments

- | | |
|---------|---|
| W | Adjacency matrix of locations. All non-zero entries are treated as if the two locations are adjacent. If <code>formula</code> contains a grouping factor, the row names of <code>W</code> have to match the levels of the grouping factor. |
| formula | An optional one-sided formula of the form <code>~ 1 g</code> , where <code>g</code> is a grouping factor mapping observations to spatial locations. If not specified, each observation is treated as a separate location. It is recommended to always specify a grouping factor to allow for handling of new data in post-processing methods. |
| type | Type of the CAR structure. Currently implemented are "escar" (exact sparse CAR), "esicar" (exact sparse intrinsic CAR), "icar" (intrinsic CAR), and "bym2". More information is provided in the 'Details' section. |

Details

The escar and esicar types are implemented based on the case study of Max Joseph (<https://github.com/mbjoseph/CARstan>). The icar and bym2 type is implemented based on the case study of Mitzi Morris (https://mc-stan.org/users/documentation/case-studies/icar_stan.html).

Examples

```

## Not run:
# generate some spatial data
east <- north <- 1:10
Grid <- expand.grid(east, north)
K <- nrow(Grid)

# set up distance and neighbourhood matrices
distance <- as.matrix(dist(Grid))
W <- array(0, c(K, K))
W[distance == 1] <- 1

# generate the covariates and response data
x1 <- rnorm(K)
x2 <- rnorm(K)
theta <- rnorm(K, sd = 0.05)
phi <- rmulti_normal(
  1, mu = rep(0, K), Sigma = 0.4 * exp(-0.1 * distance)
)
eta <- x1 + x2 + phi
prob <- exp(eta) / (1 + exp(eta))
size <- rep(50, K)
y <- rbinom(n = K, size = size, prob = prob)
dat <- data.frame(y, size, x1, x2)

# fit a CAR model
fit <- brm(y | trials(size) ~ x1 + x2, data = dat,
            family = binomial(), autocor = cor_car(W))
summary(fit)

## End(Not run)

```

Description

This function is deprecated. Please see [cosy](#) for the new syntax. This functions is a constructor for the `cor_cosy` class, representing a compound symmetry structure corresponding to uniform correlation.

Usage

```
cor_cosy(formula = ~1)
```

Arguments

<code>formula</code>	A one sided formula of the form $\sim t$, or $\sim t g$, specifying a time covariate t and, optionally, a grouping factor g . A covariate for this correlation structure must be integer valued. When a grouping factor is present in <code>formula</code> , the correlation structure is assumed to apply only to observations within the same grouping level; observations with different grouping levels are assumed to be uncorrelated. Defaults to ~ 1 , which corresponds to using the order of the observations in the data as a covariate, and no groups.
----------------------	--

Value

An object of class `cor_cosy`, representing a compound symmetry correlation structure.

Examples

```
cor_cosy(~ visit | patient)
```

cor_fixed

(*Deprecated*) Fixed user-defined covariance matrices

Description

This function is deprecated. Please see `fcor` for the new syntax. Define a fixed covariance matrix of the response variable for instance to model multivariate effect sizes in meta-analysis.

Usage

```
cor_fixed(V)
```

Arguments

<code>V</code>	Known covariance matrix of the response variable. If a vector is passed, it will be used as diagonal entries (variances) and covariances will be set to zero.
----------------	---

Value

An object of class `cor_fixed`.

Examples

```
## Not run:
dat <- data.frame(y = rnorm(3))
V <- cbind(c(0.5, 0.3, 0.2), c(0.3, 1, 0.1), c(0.2, 0.1, 0.2))
fit <- brm(y~1, data = dat, autocor = cor_fixed(V))

## End(Not run)
```

cor_ma*(Deprecated) MA(q) correlation structure*

Description

This function is deprecated. Please see [ma](#) for the new syntax. This function is a constructor for the `cor_arma` class, allowing for moving average terms only.

Usage

```
cor_ma(formula = ~1, q = 1, cov = FALSE)
```

Arguments

<code>formula</code>	A one sided formula of the form $\sim t$, or $\sim t g$, specifying a time covariate t and, optionally, a grouping factor g . A covariate for this correlation structure must be integer valued. When a grouping factor is present in <code>formula</code> , the correlation structure is assumed to apply only to observations within the same grouping level; observations with different grouping levels are assumed to be uncorrelated. Defaults to ~ 1 , which corresponds to using the order of the observations in the data as a covariate, and no groups.
<code>q</code>	A non-negative integer specifying the moving average (MA) order of the ARMA structure. Default is 1.
<code>cov</code>	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If FALSE (the default) a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Value

An object of class `cor_arma` containing solely moving average terms.

See Also

[cor_arma](#)

Examples

```
cor_ma(~visit|patient, q = 2)
```

cor_sar

(Deprecated) Spatial simultaneous autoregressive (SAR) structures

Description

These functions are deprecated. Please see [sar](#) for the new syntax. These functions are constructors for the `cor_sar` class implementing spatial simultaneous autoregressive structures. The `lagsar` structure implements SAR of the response values:

$$y = \rho W y + \eta + e$$

The `errorsar` structure implements SAR of the residuals:

$$y = \eta + u, u = \rho W u + e$$

In the above equations, η is the predictor term and e are independent normally or t-distributed residuals.

Usage

```
cor_sar(W, type = c("lag", "error"))

cor_lagsar(W)

cor_errorsar(W)
```

Arguments

- | | |
|-------------------|--|
| <code>W</code> | An object specifying the spatial weighting matrix. Can be either the spatial weight matrix itself or an object of class <code>listw</code> or <code>nb</code> , from which the spatial weighting matrix can be computed. |
| <code>type</code> | Type of the SAR structure. Either "lag" (for SAR of the response values) or "error" (for SAR of the residuals). |

Details

Currently, only families `gaussian` and `student` support SAR structures.

Value

An object of class `cor_sar` to be used in calls to [brm](#).

Examples

```
## Not run:
data(oldcol, package = "spdep")
fit1 <- brm(CRIME ~ INC + HOVAL, data = COL.OLD,
             autocor = cor_lagsar(COL.nb),
             chains = 2, cores = 2)
```

```
summary(fit1)
plot(fit1)

fit2 <- brm(CRIME ~ INC + HOVAL, data = COL.OLD,
             autocor = cor_errorsar(COL.nb),
             chains = 2, cores = 2)
summary(fit2)
plot(fit2)

## End(Not run)
```

cosy

Set up COSY correlation structures

Description

Set up a compounds symmetry (COSY) term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with COSY terms.

Usage

```
cosy(time = NA, gr = NA)
```

Arguments

- | | |
|------|---|
| time | An optional time variable specifying the time ordering of the observations. By default, the existing order of the observations in the data is used. |
| gr | An optional grouping variable. If specified, the correlation structure is assumed to apply only to observations within the same grouping level. |

Value

An object of class 'cosy_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#)

Examples

```
## Not run:
data("lh")
lh <- as.data.frame(lh)
fit <- brm(x ~ cosy(), data = lh)
summary(fit)

## End(Not run)
```

`create_priorsense_data.brmsfit`
Prior sensitivity: Create priorsense data

Description

The `create_priorsense_data.brmsfit` method can be used to create the data structure needed by the **priorsense** package for performing power-scaling sensitivity analysis. This method is called automatically when performing powerscaling via `powerscale` or other related functions, so you will rarely need to call it manually yourself.

Usage

```
create_priorsense_data.brmsfit(x, ...)
```

Arguments

- | | |
|----------------|--|
| <code>x</code> | A <code>brmsfit</code> object. |
| ... | Additional arguments passed to <code>log_lik</code> , for example <code>newdata</code> . |

Value

A `priorsense_data` object to be used in conjunction with the **priorsense** package.

Examples

```
## Not run:
# fit a model with non-uniform priors
fit <- brm(rating ~ treat + period + carry,
            data = inhaler, family = sratio(),
            prior = set_prior("normal(0, 0.5)"))
summary(fit)

# The following code requires the 'priorsense' package to be installed:
library(priorsense)

# perform power-scaling of the prior
powerscale(fit, alpha = 1.5, component = "prior")

# perform power-scaling sensitivity checks
powerscale_sensitivity(fit)

# create power-scaling sensitivity plots (for one variable)
powerscale_plot_dens(fit, variable = "b_treat")

## End(Not run)
```

cs*Category Specific Predictors in brms Models*

Description

Category Specific Predictors in **brms** Models

Usage

```
cs(expr)
```

Arguments

`expr` Expression containing predictors, for which category specific effects should be estimated. For evaluation, R formula syntax is applied.

Details

For detailed documentation see `help(brmsformula)` as well as `vignette("brms_overview")`.

This function is almost solely useful when called in formulas passed to the **brms** package.

See Also

[brmsformula](#)

Examples

```
## Not run:  
fit <- brm(rating ~ period + carry + cs(treat),  
           data = inhaler, family = sratio("cloglog"),  
           prior = set_prior("normal(0,5)"), chains = 2)  
summary(fit)  
plot(fit, ask = FALSE)  
  
## End(Not run)
```

custom_family*Custom Families in brms Models*

Description

Define custom families (i.e. response distribution) for use in **brms** models. It allows users to benefit from the modeling flexibility of **brms**, while applying their self-defined likelihood functions. All of the post-processing methods for `brmsfit` objects can be made compatible with custom families. See `vignette("brms_customfamilies")` for more details. For a list of built-in families see [brmsfamily](#).

Usage

```
custom_family(
  name,
  dpars = "mu",
  links = "identity",
  type = c("real", "int"),
  lb = NA,
  ub = NA,
  vars = NULL,
  loop = TRUE,
  specials = NULL,
  threshold = "flexible",
  log_lik = NULL,
  posterior_predict = NULL,
  posterior_epred = NULL,
  predict = NULL,
  fitted = NULL,
  env = parent.frame()
)
```

Arguments

<code>name</code>	Name of the custom family.
<code>dpars</code>	Names of the distributional parameters of the family. One parameter must be named "mu" and the main formula of the model will correspond to that parameter.
<code>links</code>	Names of the link functions of the distributional parameters.
<code>type</code>	Indicates if the response distribution is continuous ("real") or discrete ("int"). This controls if the corresponding density function will be named with <code><name>_lpdf</code> or <code><name>_lpmf</code> .
<code>lb</code>	Vector of lower bounds of the distributional parameters. Defaults to NA that is no lower bound.
<code>ub</code>	Vector of upper bounds of the distributional parameters. Defaults to NA that is no upper bound.
<code>vars</code>	Names of variables that are part of the likelihood function without being distributional parameters. That is, <code>vars</code> can be used to pass data to the likelihood. Such arguments will be added to the list of function arguments at the end, after the distributional parameters. See stanvar for details about adding self-defined data to the generated Stan model. Additional arguments <code>vreal</code> and <code>vint</code> may be used for this purpose as well (see Examples below). See also brmsformula and addition-terms for more details.
<code>loop</code>	Logical; Should the likelihood be evaluated via a loop (TRUE; the default) over observations in Stan? If FALSE, the Stan code will be written in a vectorized manner over observations if possible.
<code>specials</code>	A character vector of special options to enable for this custom family. Currently for internal use only.

threshold	Optional threshold type for custom ordinal families. Ignored for non-ordinal families.
log_lik	Optional function to compute log-likelihood values of the model in R. This is only relevant if one wants to ensure compatibility with method log_lik .
posterior_predict	Optional function to compute posterior prediction of the model in R. This is only relevant if one wants to ensure compatibility with method posterior_predict .
posterior_epred	Optional function to compute expected values of the posterior predictive distribution of the model in R. This is only relevant if one wants to ensure compatibility with method posterior_epred .
predict	Deprecated alias of ‘posterior_predict’.
fitted	Deprecated alias of ‘posterior_epred’.
env	An environment in which certain post-processing functions related to the custom family can be found, if there were not directly passed to <code>custom_family</code> . This is only relevant if one wants to ensure compatibility with the methods log_lik , posterior_predict , or posterior_epred . By default, <code>env</code> is the environment from which <code>custom_family</code> is called.

Details

The corresponding probability density or mass Stan functions need to have the same name as the custom family. That is if a family is called `myfamily`, then the **Stan** functions should be called `myfamily_lpdf` or `myfamily_lpmf` depending on whether it defines a continuous or discrete distribution.

Value

An object of class `customfamily` inheriting from class `brmsfamily`.

See Also

`brmsfamily`, `brmsformula`, `stanvar`

Examples

```
## Not run:
## demonstrate how to fit a beta-binomial model
## generate some fake data
phi <- 0.7
n <- 300
z <- rnorm(n, sd = 0.2)
ntrials <- sample(1:10, n, replace = TRUE)
eta <- 1 + z
mu <- exp(eta) / (1 + exp(eta))
a <- mu * phi
b <- (1 - mu) * phi
p <- rbeta(n, a, b)
y <- rbinom(n, ntrials, p)
```

```

dat <- data.frame(y, z, ntrials)

# define a custom family
beta_binomial2 <- custom_family(
  "beta_binomial2", dpars = c("mu", "phi"),
  links = c("logit", "log"), lb = c(NA, 0),
  type = "int", vars = "vint1[n]"
)

# define the corresponding Stan density function
stan_density <- "
real beta_binomial2_lpmf(int y, real mu, real phi, int N) {
  return beta_binomial_lpmf(y | N, mu * phi, (1 - mu) * phi);
}
"
stanvars <- stanvar(scode = stan_density, block = "functions")

# fit the model
fit <- brm(y | vint(ntrials) ~ z, data = dat,
            family = beta_binomial2, stanvars = stanvars)
summary(fit)

# define a *vectorized* custom family (no loop over observations)
# notice also that 'vint' no longer has an observation index
beta_binomial2_vec <- custom_family(
  "beta_binomial2", dpars = c("mu", "phi"),
  links = c("logit", "log"), lb = c(NA, 0),
  type = "int", vars = "vint1", loop = FALSE
)

# define the corresponding Stan density function
stan_density_vec <- "
real beta_binomial2_lpmf(array[] int y, vector mu, real phi, array[] int N) {
  return beta_binomial_lpmf(y | N, mu * phi, (1 - mu) * phi);
}
"
stanvars_vec <- stanvar(scode = stan_density_vec, block = "functions")

# fit the model
fit_vec <- brm(y | vint(ntrials) ~ z, data = dat,
                 family = beta_binomial2_vec,
                 stanvars = stanvars_vec)
summary(fit_vec)

## End(Not run)

```

Description

`default_prior` is a generic function that can be used to get default priors for Bayesian models. Its original use is within the **brms** package, but new methods for use with objects from other packages can be registered to the same generic.

Usage

```
default_prior(object, ...)
get_prior(formula, ...)
```

Arguments

<code>object</code>	An object whose class will determine which method will be used. A symbolic description of the model to be fitted.
<code>...</code>	Further arguments passed to the specific method.
<code>formula</code>	Synonym of <code>object</code> for use in <code>get_prior</code> .

Details

See [default_prior.default](#) for the default method applied for **brms** models. You can view the available methods by typing `methods(default_prior)`.

Value

Usually, a `brmsprior` object. See [default_prior.default](#) for more details.

See Also

[set_prior](#), [default_prior.default](#)

Examples

```
## get all parameters and parameters classes to define priors on
(prior <- default_prior(count ~ zAge + zBase * Trt + (1|patient) + (1|obs),
                           data = epilepsy, family = poisson()))
```

`default_prior.default` *Default Priors for brms Models*

Description

Get information on all parameters (and parameter classes) for which priors may be specified including default priors.

Usage

```
## Default S3 method:
default_prior(
  object,
  data,
  family = gaussian(),
  autocor = NULL,
  data2 = NULL,
  knots = NULL,
  drop_unused_levels = TRUE,
  sparse = NULL,
  ...
)
```

Arguments

<code>object</code>	An object of class formula , brmsformula , or mvbrmsformula (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in brmsformula .
<code>data</code>	An object of class <code>data.frame</code> (or one that can be coerced to that class) containing data of all variables used in the model.
<code>family</code>	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see brmsfamily . By default, a linear <code>gaussian</code> model is applied. In multivariate models, <code>family</code> might also be a list of families.
<code>autocor</code>	(Deprecated) An optional cor_brms object describing the correlation structure within the response variable (i.e., the 'autocorrelation'). See the documentation of cor_brms for a description of the available correlation structures. Defaults to <code>NULL</code> , corresponding to no correlations. In multivariate models, <code>autocor</code> might also be a list of autocorrelation structures. It is now recommended to specify autocorrelation terms directly within <code>formula</code> . See brmsformula for more details.
<code>data2</code>	A named list of objects containing data, which cannot be passed via argument <code>data</code> . Required for some objects used in autocorrelation structures to specify dependency structures as well as for within-group covariance matrices.
<code>knots</code>	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See gamm for more details.
<code>drop_unused_levels</code>	Should unused factors levels in the data be dropped? Defaults to <code>TRUE</code> .
<code>sparse</code>	(Deprecated) Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to <code>FALSE</code>). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased. It is now recommended to use the <code>sparse</code> argument of brmsformula and related functions.
<code>...</code>	Other arguments for internal usage only.

Value

A `brmsprior` object. That is, a `data.frame` with specific columns including `prior`, `class`, `coef`, and `group` and several rows, each providing information on a parameter (or parameter class) on which priors can be specified. The `prior` column is empty except for internal default priors.

See Also

[default_prior](#), [set_prior](#)

Examples

```
# get all parameters and parameters classes to define priors on
(prior <- default_prior(count ~ zAge + zBase * Trt + (1|patient) + (1|obs),
                           data = epilepsy, family = poisson()))

# define a prior on all population-level effects a once
prior$prior[1] <- "normal(0,10)"

# define a specific prior on the population-level effect of Trt
prior$prior[5] <- "student_t(10, 0, 5)"

# verify that the priors indeed found their way into Stan's model code
stancode(count ~ zAge + zBase * Trt + (1|patient) + (1|obs),
          data = epilepsy, family = poisson(),
          prior = prior)
```

density_ratio

Compute Density Ratios

Description

Compute the ratio of two densities at given points based on draws of the corresponding distributions.

Usage

```
density_ratio(x, y = NULL, point = 0, n = 4096, ...)
```

Arguments

- x Vector of draws from the first distribution, usually the posterior distribution of the quantity of interest.
- y Optional vector of draws from the second distribution, usually the prior distribution of the quantity of interest. If `NULL` (the default), only the density of x will be evaluated.
- point Numeric values at which to evaluate and compare the densities. Defaults to `0`.
- n Single numeric value. Influences the accuracy of the density estimation. See [density](#) for details.
- ... Further arguments passed to [density](#).

Details

In order to achieve sufficient accuracy in the density estimation, more draws than usual are required. That is you may need an effective sample size of 10,000 or more to reliably estimate the densities.

Value

A vector of length equal to `length(point)`. If `y` is provided, the density ratio of `x` against `y` is returned. Else, only the density of `x` is returned.

Examples

```
x <- rnorm(10000)
y <- rnorm(10000, mean = 1)
density_ratio(x, y, point = c(0, 1))
```

diagnostic-quantities Extract Diagnostic Quantities of brms Models

Description

Extract quantities that can be used to diagnose sampling behavior of the algorithms applied by **Stan** at the back-end of **brms**.

Usage

```
## S3 method for class 'brmsfit'
log_posterior(object, ...)

## S3 method for class 'brmsfit'
nuts_params(object, pars = NULL, ...)

## S3 method for class 'brmsfit'
rhat(x, pars = NULL, ...)

## S3 method for class 'brmsfit'
neff_ratio(object, pars = NULL, ...)
```

Arguments

<code>object, x</code>	A <code>brmsfit</code> object.
<code>...</code>	Arguments passed to individual methods.
<code>pars</code>	An optional character vector of parameter names. For <code>nuts_params</code> these will be NUTS sampler parameter names rather than model parameters. If <code>pars</code> is omitted all parameters are included.

Details

For more details see [bayesplot-extractors](#).

Value

The exact form of the output depends on the method.

Examples

```
## Not run:
fit <- brm(time ~ age * sex, data = kidney)

lp <- log_posterior(fit)
head(lp)

np <- nuts_params(fit)
str(np)
# extract the number of divergence transitions
sum(subset(np, Parameter == "divergent_")$Value)

head(rhat(fit))
head(neff_ratio(fit))

## End(Not run)
```

Description

Density function and random number generation for the dirichlet distribution with shape parameter vector alpha.

Usage

```
ddirichlet(x, alpha, log = FALSE)

rdirichlet(n, alpha)
```

Arguments

x	Matrix of quantiles. Each row corresponds to one probability vector.
alpha	Matrix of positive shape parameters. Each row corresponds to one probability vector.
log	Logical; If TRUE, values are returned on the log scale.
n	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

`draws-brms`

Transform brmsfit to draws objects

Description

Transform a `brmsfit` object to a format supported by the **posterior** package.

Usage

```
## S3 method for class 'brmsfit'
as_draws(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)

## S3 method for class 'brmsfit'
as_draws_matrix(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)

## S3 method for class 'brmsfit'
as_draws_array(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)

## S3 method for class 'brmsfit'
as_draws_df(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)

## S3 method for class 'brmsfit'
as_draws_list(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)

## S3 method for class 'brmsfit'
as_draws_rvars(x, variable = NULL, regex = FALSE, inc_warmup = FALSE, ...)
```

Arguments

<code>x</code>	A <code>brmsfit</code> object or another R object for which the methods are defined.
<code>variable</code>	A character vector providing the variables to extract. By default, all variables are extracted.
<code>regex</code>	Logical; Should variable should be treated as a (vector of) regular expressions? Any variable in <code>x</code> matching at least one of the regular expressions will be selected. Defaults to FALSE.
<code>inc_warmup</code>	Should warmup draws be included? Defaults to FALSE.
<code>...</code>	Arguments passed to individual methods (if applicable).

Details

To subset iterations, chains, or draws, use the `subset_draws` method after transforming the `brmsfit` to a `draws` object.

See Also[draws](#) [subset_draws](#)**Examples**

```
## Not run:  
fit <- brm(count ~ zAge + zBase * Trt + (1|patient),  
           data = epilepsy, family = poisson())  
  
# extract posterior draws in an array format  
(draws_fit <- as_draws_array(fit))  
posterior::summarize_draws(draws_fit)  
  
# extract only certain variables  
as_draws_array(fit, variable = "r_patient")  
as_draws_array(fit, variable = "^b_", regex = TRUE)  
  
# extract posterior draws in a random variables format  
as_draws_rvars(fit)  
  
## End(Not run)
```

draws-index-brms *Index brmsfit objects*

Description

Index `brmsfit` objects

Usage

```
## S3 method for class 'brmsfit'  
variables(x, ...)  
  
## S3 method for class 'brmsfit'  
nvariables(x, ...)  
  
## S3 method for class 'brmsfit'  
niterations(x)  
  
## S3 method for class 'brmsfit'  
nchains(x)  
  
## S3 method for class 'brmsfit'  
ndraws(x)
```

Arguments

- x A `brmsfit` object or another R object for which the methods are defined.
- ... Arguments passed to individual methods (if applicable).

`emmeans-brms-helpers` *Support Functions for emmeans*

Description

Functions required for compatibility of `brms` with `emmeans`. Users are not required to call these functions themselves. Instead, they will be called automatically by the `emmeans` function of the `emmeans` package.

Usage

```
recover_data.brmsfit(
  object,
  data,
  resp = NULL,
  dpar = NULL,
  nlpar = NULL,
  re_formula = NA,
  epred = FALSE,
  ...
)

emm_basis.brmsfit(
  object,
  trms,
  xlev,
  grid,
  vcov.,
  resp = NULL,
  dpar = NULL,
  nlpar = NULL,
  re_formula = NA,
  epred = FALSE,
  ...
)
```

Arguments

- `object` An object of class `brmsfit`.
- `data, trms, xlev, grid, vcov.` Arguments required by `emmeans`.

<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>dpar</code>	Optional name of a predicted distributional parameter. If specified, expected predictions of this parameters are returned.
<code>nlpar</code>	Optional name of a predicted non-linear parameter. If specified, expected predictions of this parameters are returned.
<code>re_formula</code>	Optional formula containing group-level effects to be considered in the prediction. If <code>NULL</code> , include all group-level effects; if <code>NA</code> (default), include no group-level effects.
<code>epred</code>	Logical. If <code>TRUE</code> compute predictions of the posterior predictive distribution's mean (see posterior_epred.brmsfit) while ignoring arguments <code>dpar</code> and <code>nlpar</code> . Defaults to <code>FALSE</code> . If you have specified a response transformation within the formula, you need to set <code>epred</code> to <code>TRUE</code> for emmeans to detect this transformation.
<code>...</code>	Additional arguments passed to emmeans .

Details

In order to ensure compatibility of most **brms** models with **emmeans**, predictions are not generated ‘manually’ via a design matrix and coefficient vector, but rather via [posterior_linpred.brmsfit](#). This appears to generally work well, but note that it produces an ‘`@linfct`’ slot that contains the computed predictions as columns instead of the coefficients.

Examples

```
## Not run:
fit1 <- brm(time | cens(censored) ~ age * sex + disease + (1|patient),
              data = kidney, family = lognormal())
summary(fit1)

# summarize via 'emmeans'
library(emmeans)
rg <- ref_grid(fit1)
em <- emmeans(rg, "disease")
summary(em, point.est = mean)

# obtain estimates for the posterior predictive distribution's mean
epred <- emmeans(fit1, "disease", epred = TRUE)
summary(epred, point.est = mean)

# model with transformed response variable
fit2 <- brm(log(mpg) ~ factor(cyl), data = mtcars)
summary(fit2)

# results will be on the log scale by default
emmeans(fit2, ~ cyl)
# log transform is detected and can be adjusted automatically
emmeans(fit2, ~ cyl, epred = TRUE, type = "response")
```

```
## End(Not run)
```

epilepsy	<i>Epileptic seizure counts</i>
----------	---------------------------------

Description

Breslow and Clayton (1993) analyze data initially provided by Thall and Vail (1990) concerning seizure counts in a randomized trial of anti-convulsant therapy in epilepsy. Covariates are treatment, 8-week baseline seizure counts, and age of the patients in years.

Usage

```
epilepsy
```

Format

A data frame of 236 observations containing information on the following 9 variables.

Age The age of the patients in years

Base The seizure count at 8-weeks baseline

Trt Either 0 or 1 indicating if the patient received anti-convulsant therapy

patient The patient number

visit The session number from 1 (first visit) to 4 (last visit)

count The seizure count between two visits

obs The observation number, that is a unique identifier for each observation

zAge Standardized Age

zBase Standardized Base

Source

Thall, P. F., & Vail, S. C. (1990). Some covariance models for longitudinal count data with overdispersion. *Biometrics*, 46(2), 657-671.

Breslow, N. E., & Clayton, D. G. (1993). Approximate inference in generalized linear mixed models. *Journal of the American Statistical Association*, 88(421), 9-25.

Examples

```
## Not run:
## poisson regression without random effects.
fit1 <- brm(count ~ zAge + zBase * Trt,
             data = epilepsy, family = poisson())
summary(fit1)
plot(fit1)
```

```

## poisson regression with varying intercepts of patients
## as well as normal priors for overall effects parameters.
fit2 <- brm(count ~ zAge + zBase * Trt + (1|patient),
             data = epilepsy, family = poisson(),
             prior = set_prior("normal(0,5)"))
summary(fit2)
plot(fit2)

## End(Not run)

```

Description

Density, distribution function, and random generation for the exponentially modified Gaussian distribution with mean *mu* and standard deviation *sigma* of the gaussian component, as well as scale *beta* of the exponential component.

Usage

```

dexgaussian(x, mu, sigma, beta, log = FALSE)

pexgaussian(q, mu, sigma, beta, lower.tail = TRUE, log.p = FALSE)

rexgaussian(n, mu, sigma, beta)

```

Arguments

<i>x, q</i>	Vector of quantiles.
<i>mu</i>	Vector of means of the combined distribution.
<i>sigma</i>	Vector of standard deviations of the gaussian component.
<i>beta</i>	Vector of scales of the exponential component.
<i>log</i>	Logical; If TRUE, values are returned on the log scale.
<i>lower.tail</i>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<i>log.p</i>	Logical; If TRUE, values are returned on the log scale.
<i>n</i>	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

`expose_functions.brmsfit`

Expose user-defined Stan functions

Description

Export user-defined **Stan** function and optionally vectorize them. For more details see [expose_stan_functions](#).

Usage

```
## S3 method for class 'brmsfit'
expose_FUNCTIONS(x, vectorize = FALSE, env = globalenv(), ...)

expose_FUNCTIONS(x, ...)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code> .
<code>vectorize</code>	Logical; Indicates if the exposed functions should be vectorized via Vectorize . Defaults to FALSE.
<code>env</code>	Environment where the functions should be made available. Defaults to the global environment.
<code>...</code>	Further arguments passed to expose_stan_functions .

`exp1`

Exponential function plus one.

Description

Computes $\exp(x) + 1$.

Usage

```
exp1(x)
```

Arguments

<code>x</code>	A numeric or complex vector.
----------------	------------------------------

family.brmsfit *Extract Model Family Objects*

Description

Extract Model Family Objects

Usage

```
## S3 method for class 'brmsfit'  
family(object, resp = NULL, ...)
```

Arguments

- | | |
|--------|--|
| object | An object of class brmsfit . |
| resp | Optional names of response variables. If specified, predictions are performed only for the specified response variables. |
| ... | Currently unused. |

Value

A **brmsfamily** object or a list of such objects for multivariate models.

fcor *Fixed residual correlation (FCOR) structures*

Description

Set up a fixed residual correlation (FCOR) term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with FCOR terms.

Usage

```
fcor(M)
```

Arguments

- | | |
|---|--|
| M | Known correlation/covariance matrix of the response variable. If a vector is passed, it will be used as diagonal entries (variances) and correlations/covariances will be set to zero. The actual covariance matrix used in the likelihood is obtained by multiplying M by the square of the residual standard deviation parameter sigma estimated as part of the model. |
|---|--|

Value

An object of class 'fcor_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#)

Examples

```
## Not run:
dat <- data.frame(y = rnorm(3))
V <- cbind(c(0.5, 0.3, 0.2), c(0.3, 1, 0.1), c(0.2, 0.1, 0.2))
fit <- brm(y ~ 1 + fcor(V), data = dat, data2 = list(V = V))

## End(Not run)
```

fitted.brmsfit

Expected Values of the Posterior Predictive Distribution

Description

This method is an alias of [posterior_epred.brmsfit](#) with additional arguments for obtaining summaries of the computed draws.

Usage

```
## S3 method for class 'brmsfit'
fitted(
  object,
  newdata = NULL,
  re_formula = NULL,
  scale = c("response", "linear"),
  resp = NULL,
  dpar = NULL,
  nlpar = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)
```

Arguments

object	An object of class <code>brmsfit</code> .
newdata	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
re_formula	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
scale	Either "response" or "linear". If "response", results are returned on the scale of the response variable. If "linear", results are returned on the scale of the linear predictor term, that is without applying the inverse link function or other transformations.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
dpar	Optional name of a predicted distributional parameter. If specified, expected predictions of this parameters are returned.
nlpar	Optional name of a predicted non-linear parameter. If specified, expected predictions of this parameters are returned.
ndraws	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
draw_ids	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
sort	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (<code>FALSE</code> ; default) or in the order of the time series (<code>TRUE</code>).
summary	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .
robust	If <code>FALSE</code> (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If <code>TRUE</code> , the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is <code>TRUE</code> .
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is <code>TRUE</code> .
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Value

An array of predicted *mean* response values. If `summary = FALSE` the output resembles those of `posterior_epred.brmsfit`.

If `summary = TRUE` the output depends on the family: For categorical and ordinal families, the output is an $N \times E \times C$ array, where N is the number of observations, E is the number of summary statistics,

and C is the number of categories. For all other families, the output is an N x E matrix. The number of summary statistics E is equal to 2 + length(probs): The Estimate column contains point estimates (either mean or median depending on argument robust), while the Est.Error column contains uncertainty estimates (either standard deviation or median absolute deviation depending on argument robust). The remaining columns starting with Q contain quantile estimates as specified via argument probs.

In multivariate models, an additional dimension is added to the output which indexes along the different response variables.

See Also

[posterior_epred.brmsfit](#)

Examples

```
## Not run:
## fit a model
fit <- brm(rating ~ treat + period + carry + (1|subject),
           data = inhaler)

## compute expected predictions
fitted_values <- fitted(fit)
head(fitted_values)

## plot expected predictions against actual response
dat <- as.data.frame(cbind(Y = standata(fit)$Y, fitted_values))
ggplot(dat) + geom_point(aes(x = Estimate, y = Y))

## End(Not run)
```

fixef.brmsfit

Extract Population-Level Estimates

Description

Extract the population-level ('fixed') effects from a *brmsfit* object.

Usage

```
## S3 method for class 'brmsfit'
fixef(
  object,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  pars = NULL,
  ...
)
```

Arguments

object	An object of class <code>brmsfit</code> .
summary	Should summary statistics be returned instead of the raw values? Default is TRUE.
robust	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
pars	Optional names of coefficients to extract. By default, all coefficients are extracted.
...	Currently ignored.

Value

If `summary` is TRUE, a matrix returned by `posterior_summary` for the population-level effects. If `summary` is FALSE, a matrix with one row per posterior draw and one column per population-level effect.

Examples

```
## Not run:
fit <- brm(time | cens(censored) ~ age + sex + disease,
            data = kidney, family = "exponential")
fixef(fit)
# extract only some coefficients
fixef(fit, pars = c("age", "sex"))

## End(Not run)
```

Description

Density, distribution function, quantile function and random generation for the Frechet distribution with location `loc`, scale `scale`, and shape `shape`.

Usage

```
dfrechet(x, loc = 0, scale = 1, shape = 1, log = FALSE)

pfrechet(q, loc = 0, scale = 1, shape = 1, lower.tail = TRUE, log.p = FALSE)
```

```
qfrechet(p, loc = 0, scale = 1, shape = 1, lower.tail = TRUE, log.p = FALSE)
rfrechet(n, loc = 0, scale = 1, shape = 1)
```

Arguments

x, q	Vector of quantiles.
loc	Vector of locations.
scale	Vector of scales.
shape	Vector of shapes.
log	Logical; If TRUE, values are returned on the log scale.
lower.tail	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
log.p	Logical; If TRUE, values are returned on the log scale.
p	Vector of probabilities.
n	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

Description

Density, distribution function, and random generation for the generalized extreme value distribution with location `mu`, scale `sigma` and shape `xi`.

Usage

```
dgen_extreme_value(x, mu = 0, sigma = 1, xi = 0, log = FALSE)

pgen_extreme_value(
  q,
  mu = 0,
  sigma = 1,
  xi = 0,
  lower.tail = TRUE,
  log.p = FALSE
)

qgen_extreme_value(
  p,
  mu = 0,
  sigma = 1,
```

```

  xi = 0,
  lower.tail = TRUE,
  log.p = FALSE
)

rgenExtremeValue(n, mu = 0, sigma = 1, xi = 0)

```

Arguments

x, q	Vector of quantiles.
mu	Vector of locations.
sigma	Vector of scales.
xi	Vector of shapes.
log	Logical; If TRUE, values are returned on the log scale.
lower.tail	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
log.p	Logical; If TRUE, values are returned on the log scale.
p	Vector of probabilities.
n	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

get_dpar	<i>Draws of a Distributional Parameter</i>
----------	--

Description

Get draws of a distributional parameter from a `brmsprep` or `mvbrmsprep` object. This function is primarily useful when developing custom families or packages depending on **brms**. This function lets callers easily handle both the case when the distributional parameter is predicted directly, via a (non-)linear predictor or fixed to a constant. See the vignette `vignette("brms_customfamilies")` for an example use case.

Usage

```
get_dpar(prep, dpar, i = NULL, inv_link = NULL)
```

Arguments

prep	A 'brmsprep' or 'mvbrmsprep' object created by prepare_predictions .
dpar	Name of the distributional parameter.
i	The observation numbers for which predictions shall be extracted. If NULL (the default), all observation will be extracted. Ignored if dpar is not predicted.
inv_link	Should the inverse link function be applied? If NULL (the default), the value is chosen internally. In particular, <code>inv_link</code> is TRUE by default for custom families.

Value

If the parameter is predicted and *i* is NULL or `length(i) > 1`, an S x N matrix. If the parameter is not predicted or `length(i) == 1`, a vector of length S. Here S is the number of draws and N is the number of observations or length of *i* if specified.

Examples

```
## Not run:
posterior_predict_my_dist <- function(i, prep, ...) {
  mu <- brms:::get_dpar(prep, "mu", i = i)
  mypar <- brms:::get_dpar(prep, "mypar", i = i)
  my_rng(mu, mypar)
}

## End(Not run)
```

`get_refmodel.brmsfit` *Projection Predictive Variable Selection: Get Reference Model*

Description

The `get_refmodel.brmsfit` method can be used to create the reference model structure which is needed by the **projpred** package for performing a projection predictive variable selection. This method is called automatically when performing variable selection via `varsel` or `cv_varsel`, so you will rarely need to call it manually yourself.

Usage

```
get_refmodel.brmsfit(
  object,
  newdata = NULL,
  resp = NULL,
  cvfun = NULL,
  dis = NULL,
  latent = FALSE,
  brms_seed = NULL,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.

resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
cvfun	Optional cross-validation function (see <code>get_refmodel</code> for details). If NULL (the default), cvfun is defined internally based on <code>kfold.brmsfit</code> .
dis	Passed to argument dis of <code>init_refmodel</code> , but leave this at NULL unless projpred complains about it.
latent	See argument latent of <code>extend_family</code> . Setting this to TRUE requires a projpred version >= 2.4.0.
brms_seed	A seed used to infer seeds for <code>kfold.brmsfit</code> and for sampling group-level effects for new levels (in multilevel models). If NULL, then <code>set.seed</code> is not called at all. If not NULL, then the pseudorandom number generator (PRNG) state is reset (to the state before calling this function) upon exiting this function.
...	Further arguments passed to <code>init_refmodel</code> .

Details

The extract_model_data function used internally by `get_refmodel.brmsfit` ignores arguments wrhs and orhs (a warning is thrown if these are non-NULL). For example, arguments weightsnew and offsetnew of `proj_linpred`, `proj_predict`, and `predict.refmodel` are passed to wrhs and orhs, respectively.

Value

A refmodel object to be used in conjunction with the **projpred** package.

Examples

```
## Not run:
# fit a simple model
fit <- brm(count ~ zAge + zBase * Trt,
            data = epilepsy, family = poisson())
summary(fit)

# The following code requires the 'projpred' package to be installed:
library(projpred)

# perform variable selection without cross-validation
vs <- varsel(fit)
summary(vs)
plot(vs)

# perform variable selection with cross-validation
cv_vs <- cv_varsel(fit)
summary(cv_vs)
plot(cv_vs)

## End(Not run)
```

gp*Set up Gaussian process terms in brms*

Description

Set up a Gaussian process (GP) term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with GP terms.

Usage

```
gp(
  ...,
  by = NA,
  k = NA,
  cov = "exp_quad",
  iso = TRUE,
  gr = TRUE,
  cmc = TRUE,
  scale = TRUE,
  c = 5/4
)
```

Arguments

...	One or more predictors for the GP.
by	A numeric or factor variable of the same length as each predictor. In the numeric vector case, the elements multiply the values returned by the GP. In the factor variable case, a separate GP is fitted for each factor level.
k	Optional number of basis functions for computing Hilbert-space approximate GPs. If NA (the default), exact GPs are computed.
cov	Name of the covariance kernel. Currently supported are "exp_quad" (exponentiated-quadratic kernel; default), "matern32" (Matern 3/2 kernel), "matern52" (Matern 5/2 kernel), and "exponential" (exponential kernel; alias: "matern12").
iso	A flag to indicate whether an isotropic (TRUE; the default) or a non-isotropic GP should be used. In the former case, the same amount of smoothing is applied to all predictors. In the latter case, predictors may have different smoothing. Ignored if only a single predictor is supplied.
gr	Logical; Indicates if auto-grouping should be used (defaults to TRUE). If enabled, observations sharing the same predictor values will be represented by the same latent variable in the GP. This will improve sampling efficiency drastically if the number of unique predictor combinations is small relative to the number of observations.
cmc	Logical; Only relevant if by is a factor. If TRUE (the default), cell-mean coding is used for the by-factor, that is one GP per level is estimated. If FALSE, contrast GPs are estimated according to the contrasts set for the by-factor.

scale	Logical; If TRUE (the default), predictors are scaled so that the maximum Euclidean distance between two points is 1. This often improves sampling speed and convergence. Scaling also affects the estimated length-scale parameters in that they resemble those of scaled predictors (not of the original predictors) if scale is TRUE.
c	Numeric value only used in approximate GPs. Defines the multiplicative constant of the predictors' range over which predictions should be computed. A good default could be $c = 5/4$ but we are still working on providing better recommendations.

Details

A GP is a stochastic process, which describes the relation between one or more predictors $x = (x_1, \dots, x_d)$ and a response $f(x)$, where d is the number of predictors. A GP is the generalization of the multivariate normal distribution to an infinite number of dimensions. Thus, it can be interpreted as a prior over functions. The values of $f()$ at any finite set of locations are jointly multivariate normal, with a covariance matrix defined by the covariance kernel $k_p(x_i, x_j)$, where p is the vector of parameters of the GP:

$$(f(x_1), \dots, f(x_n)) \sim MVN(0, (k_p(x_i, x_j))_{i,j=1}^n).$$

The smoothness and general behavior of the function f depends only on the choice of covariance kernel. For a more detailed introduction to Gaussian processes, see https://en.wikipedia.org/wiki/Gaussian_process.

For mathematical details on the supported kernels, please see the Stan manual: https://mc-stan.org/docs/functions-reference/matrix_operations.html under "Gaussian Process Covariance Functions".

There are several parameters estimated for GPs, the most important of which are as follows:

Parameter	Notation	Support	Meaning
lscale	ℓ	\mathbb{R}^+	length-scale of the GP's covariance kernel
sdgp	σ	\mathbb{R}^+	marginal standard deviation of the GP's covariance kernel
zgp	z	\mathbb{R}	latent variable values of the training data

Assuming an exponentiated quadratic covariance structure, the parameters can be broadly interpreted as follows (see also <https://mc-stan.org/docs/stan-users-guide/gaussian-processes.html#gaussian-process-regression>). Note that in the above documentation the parameter σ is denoted as α instead, and the parameter ℓ as ρ . The length-scale ℓ controls the frequency of the functions represented by the GP prior i.e., values of $\ell \gg 0$ lead to lower-frequency functions, while values of $\ell \approx 0$ lead to higher-frequency functions. In slightly simpler terms, ℓ sets the distance over which observations in the input space are strongly correlated. The marginal standard deviation σ controls the magnitude of the range of the function represented by the GP i.e., it represents how much the values of the function tend to deviate from the mean level. Lastly, the parameter z represents latent variable values per observation (or unique input point).

Value

An object of class 'gp_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[brmsformula](#)

Examples

```

## Not run:
# simulate data using the mgcv package
dat <- mgcv::gamSim(1, n = 30, scale = 2)

# fit a simple GP model
fit1 <- brm(y ~ gp(x2), dat, chains = 2)
summary(fit1)
me1 <- conditional_effects(fit1, ndraws = 200, spaghetti = TRUE)
plot(me1, ask = FALSE, points = TRUE)

# fit a more complicated GP model and use an approximate GP for x2
fit2 <- brm(y ~ gp(x0) + x1 + gp(x2, k = 10) + x3, dat, chains = 2)
summary(fit2)
me2 <- conditional_effects(fit2, ndraws = 200, spaghetti = TRUE)
plot(me2, ask = FALSE, points = TRUE)

# fit a multivariate GP model with Matern 3/2 kernel
fit3 <- brm(y ~ gp(x1, x2, cov = "matern32"), dat, chains = 2)
summary(fit3)
me3 <- conditional_effects(fit3, ndraws = 200, spaghetti = TRUE)
plot(me3, ask = FALSE, points = TRUE)

# compare model fit
loo(fit1, fit2, fit3)

# simulate data with a factor covariate
dat2 <- mgcv::gamSim(4, n = 90, scale = 2)

# fit separate gaussian processes for different levels of 'fac'
fit4 <- brm(y ~ gp(x2, by = fac), dat2, chains = 2)
summary(fit4)
plot(conditional_effects(fit4), points = TRUE)

## End(Not run)

```

Description

Function used to set up a basic grouping term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with grouping terms. **gr** is called implicitly inside the package and there is usually no need to call it directly.

Usage

```
gr(
  ...,
  by = NULL,
  cor = TRUE,
  id = NA,
  pw = NULL,
  cov = NULL,
  dist = "gaussian"
)
```

Arguments

...	One or more terms containing grouping factors.
by	An optional factor variable, specifying sub-populations of the groups. For each level of the by variable, a separate variance-covariance matrix will be fitted. Levels of the grouping factor must be nested in levels of the by variable.
cor	Logical. If TRUE (the default), group-level terms will be modelled as correlated.
id	Optional character string. All group-level terms across the model with the same id will be modeled as correlated (if cor is TRUE). See brmsformula for more details.
pw	Optional numeric variable specifying prior weights. They weight the contribution of each group to the log-prior of the group-level coefficients. Should have one distinct value for each level of the grouping variable.
cov	An optional matrix which is proportional to the within-group covariance matrix of the group-level effects. All levels of the grouping factor should appear as row-names of the corresponding matrix. This argument can be used, among others, to model pedigrees and phylogenetic effects. See vignette("brms_phylogenetics") for more details. By default, levels of the same grouping factor are modeled as independent of each other.
dist	Name of the distribution of the group-level effects. Currently "gaussian" is the only option.

See Also

[brmsformula](#)

Examples

```
## Not run:
# model using basic lme4-style formula
fit1 <- brm(count ~ Trt + (1|patient), data = epilepsy)
summary(fit1)

# equivalent model using 'gr' which is called anyway internally
fit2 <- brm(count ~ Trt + (1|gr(patient)), data = epilepsy)
summary(fit2)
```

```

# include Trt as a by variable
fit3 <- brm(count ~ Trt + (1|gr(patient, by = Trt)), data = epilepsy)
summary(fit3)

# include a group-level weight variable
epilepsy[['patient_samp_wgt']] <- c(1, rep(c(0.9, 1.1), each = 29))
fit4 <- brm(count ~ Trt + (1|gr(patient, pw = patient_samp_wgt)),
            data = epilepsy)
summary(fit4)

## End(Not run)

```

horseshoe*Regularized horseshoe priors in **brms*****Description**

Function used to set up regularized horseshoe priors and related hierarchical shrinkage priors in **brms**. The function does not evaluate its arguments – it exists purely to help set up the model.

Usage

```

horseshoe(
  df = 1,
  scale_global = 1,
  df_global = 1,
  scale_slab = 2,
  df_slab = 4,
  par_ratio = NULL,
  autoscale = TRUE,
  main = FALSE
)

```

Arguments

<code>df</code>	Degrees of freedom of student-t prior of the local shrinkage parameters. Defaults to 1.
<code>scale_global</code>	Scale of the student-t prior of the global shrinkage parameter. Defaults to 1. In linear models, <code>scale_global</code> will internally be multiplied by the residual standard deviation parameter <code>sigma</code> .
<code>df_global</code>	Degrees of freedom of student-t prior of the global shrinkage parameter. Defaults to 1. If <code>df_global</code> is greater 1, the shape of the prior will no longer resemble a horseshoe and it may be more appropriately called an hierarchical shrinkage prior in this case.

<code>scale_slab</code>	Scale of the Student-t slab. Defaults to 2. The original unregularized horseshoe prior is obtained by setting <code>scale_slab</code> to infinite, which we can approximate in practice by setting it to a very large real value.
<code>df_slab</code>	Degrees of freedom of the student-t slab. Defaults to 4.
<code>par_ratio</code>	Ratio of the expected number of non-zero coefficients to the expected number of zero coefficients. If specified, <code>scale_global</code> is ignored and internally computed as <code>par_ratio / sqrt(N)</code> , where N is the total number of observations in the data.
<code>autoscale</code>	Logical; indicating whether the horseshoe prior should be scaled using the residual standard deviation <code>sigma</code> if possible and sensible (defaults to TRUE). Autoscaling is not applied for distributional parameters or when the model does not contain the parameter <code>sigma</code> .
<code>main</code>	Logical (defaults to FALSE); only relevant if the horseshoe prior spans multiple parameter classes. In this case, only arguments given in the single instance where <code>main</code> is TRUE will be used. Arguments given in other instances of the prior will be ignored. See the Examples section below.

Details

The horseshoe prior is a special shrinkage prior initially proposed by Carvalho et al. (2009). It is symmetric around zero with fat tails and an infinitely large spike at zero. This makes it ideal for sparse models that have many regression coefficients, although only a minority of them is non-zero. The horseshoe prior can be applied on all population-level effects at once (excluding the intercept) by using `set_prior("horseshoe(1)")`. The 1 implies that the student-t prior of the local shrinkage parameters has 1 degrees of freedom. This may, however, lead to an increased number of divergent transitions in Stan. Accordingly, increasing the degrees of freedom to slightly higher values (e.g., 3) may often be a better option, although the prior no longer resembles a horseshoe in this case. Further, the scale of the global shrinkage parameter plays an important role in amount of shrinkage applied. It defaults to 1, but this may result in too few shrinkage (Piironen & Vehtari, 2016). It is thus possible to change the scale using argument `scale_global` of the horseshoe prior, for instance `horseshoe(1, scale_global = 0.5)`. In linear models, `scale_global` will internally be multiplied by the residual standard deviation parameter `sigma`. See Piironen and Vehtari (2016) for recommendations how to properly set the global scale. The degrees of freedom of the global shrinkage prior may also be adjusted via argument `df_global`. Piironen and Vehtari (2017) recommend specifying the ratio of the expected number of non-zero coefficients to the expected number of zero coefficients `par_ratio` rather than `scale_global` directly. As proposed by Piironen and Vehtari (2017), an additional regularization is applied that only affects non-zero coefficients. The amount of regularization can be controlled via `scale_slab` and `df_slab`. To make sure that shrinkage can equally affect all coefficients, predictors should be on the same scale. Generally, models with horseshoe priors are more likely than other models to have divergent transitions so that increasing `adapt_delta` from 0.8 to values closer to 1 will often be necessary. See the documentation of `brm` for instructions on how to increase `adapt_delta`.

The prior does not account for scale differences of the terms it is applied on. Accordingly, please make sure that all these terms have a comparable scale to ensure that shrinkage is applied properly.

Currently, the following classes support the horseshoe prior: `b` (overall regression coefficients), `sds` (SDs of smoothing splines), `sdgp` (SDs of Gaussian processes), `ar` (autoregressive coefficients), `ma`

(moving average coefficients), `sderr` (SD of latent residuals), `sdcar` (SD of spatial CAR structures), `sd` (SD of varying coefficients).

Value

A character string obtained by `match.call()` with additional arguments.

References

Carvalho, C. M., Polson, N. G., & Scott, J. G. (2009). Handling sparsity via the horseshoe. Artificial Intelligence and Statistics. <http://proceedings.mlr.press/v5/carvalho09a>

Piironen J. & Vehtari A. (2017). On the Hyperprior Choice for the Global Shrinkage Parameter in the Horseshoe Prior. Artificial Intelligence and Statistics. [https://arxiv.org/pdf/1610.05559v1](https://arxiv.org/pdf/1610.05559v1.pdf)

Piironen, J., and Vehtari, A. (2017). Sparsity information and regularization in the horseshoe and other shrinkage priors. Electronic Journal of Statistics. <https://arxiv.org/abs/1707.01694>

See Also

`set_prior`

Examples

```
set_prior(horseshoe(df = 3, par_ratio = 0.1))

# specify the horseshoe prior across multiple parameter classes
set_prior(horseshoe(df = 3, par_ratio = 0.1, main = TRUE), class = "b") +
  set_prior(horseshoe(), class = "sd")
```

Description

Density and distribution functions for hurdle distributions.

Usage

```
dhurdle_poisson(x, lambda, hu, log = FALSE)

phurdle_poisson(q, lambda, hu, lower.tail = TRUE, log.p = FALSE)

dhurdle_negbinomial(x, mu, shape, hu, log = FALSE)

phurdle_negbinomial(q, mu, shape, hu, lower.tail = TRUE, log.p = FALSE)

dhurdle_gamma(x, shape, scale, hu, log = FALSE)
```

```

phurdle_gamma(q, shape, scale, hu, lower.tail = TRUE, log.p = FALSE)

dhurdle_lognormal(x, mu, sigma, hu, log = FALSE)

phurdle_lognormal(q, mu, sigma, hu, lower.tail = TRUE, log.p = FALSE)

```

Arguments

<code>x</code>	Vector of quantiles.
<code>hu</code>	hurdle probability
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>q</code>	Vector of quantiles.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If TRUE, values are returned on the log scale.
<code>mu, lambda</code>	location parameter
<code>shape</code>	shape parameter
<code>sigma, scale</code>	scale parameter

Details

The density of a hurdle distribution can be specified as follows. If $x = 0$ set $f(x) = \theta$. Else set $f(x) = (1 - \theta) * g(x)/(1 - G(0))$ where $g(x)$ and $G(x)$ are the density and distribution function of the non-hurdle part, respectively.

`hypothesis.brmsfit` *Non-Linear Hypothesis Testing*

Description

Perform non-linear hypothesis testing for all model parameters.

Usage

```

## S3 method for class 'brmsfit'
hypothesis(
  x,
  hypothesis,
  class = "b",
  group = "",
  scope = c("standard", "ranef", "coef"),
  alpha = 0.05,
  robust = FALSE,
  seed = NULL,
  ...
)

```

```
)
hypothesis(x, ...)

## Default S3 method:
hypothesis(x, hypothesis, alpha = 0.05, robust = FALSE, ...)
```

Arguments

<code>x</code>	An R object. If it is no <code>brmsfit</code> object, it must be coercible to a <code>data.frame</code> . In the latter case, the variables used in the <code>hypothesis</code> argument need to correspond to column names of <code>x</code> , while the rows are treated as representing posterior draws of the variables.
<code>hypothesis</code>	A character vector specifying one or more non-linear hypothesis concerning parameters of the model.
<code>class</code>	A string specifying the class of parameters being tested. Default is "b" for population-level effects. Other typical options are "sd" or "cor". If <code>class</code> = <code>NULL</code> , all parameters can be tested against each other, but have to be specified with their full name (see also variables)
<code>group</code>	Name of a grouping factor to evaluate only group-level effects parameters related to this grouping factor.
<code>scope</code>	Indicates where to look for the variables specified in <code>hypothesis</code> . If "standard", use the full parameter names (subject to the restriction given by <code>class</code> and <code>group</code>). If "coef" or "ranef", compute the hypothesis for all levels of the grouping factor given in "group", based on the output of coef.brmsfit and ranef.brmsfit , respectively.
<code>alpha</code>	The alpha-level of the tests (default is 0.05; see 'Details' for more information).
<code>robust</code>	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead.
<code>seed</code>	A single numeric value passed to set.seed to make results reproducible. This is currently only relevant for point hypotheses that scope over at least two parameters (see Details).
...	Currently ignored.

Details

Among others, `hypothesis` computes an evidence ratio (`Evid.Ratio`) for each hypothesis. For a one-sided hypothesis, this is just the posterior probability (`Post.Prob`) under the hypothesis against its alternative. That is, when the hypothesis is of the form $a > b$, the evidence ratio is the ratio of the posterior probability of $a > b$ and the posterior probability of $a < b$. In this example, values greater than one indicate that the evidence in favor of $a > b$ is larger than evidence in favor of $a < b$. For a two-sided (point) hypothesis, the evidence ratio is a Bayes factor between the hypothesis and its alternative computed via the Savage-Dickey density ratio method. That is the posterior density at the point of interest divided by the prior density at that point. Values greater than one indicate that evidence in favor of the point hypothesis has increased after seeing the data. In order to calculate

this Bayes factor, all parameters related to the hypothesis must have proper priors and argument `sample_prior` of function `brm` must be set to "yes". Otherwise `Evid.Ratio` (and `Post.Prob`) will be NA.

Please note that the Savage-Dickey density ratio as implemented here provides only a very basic test of point hypotheses. It is recommended that you use bridge sampling instead (via `bayes_factor` which relies on the `bridgesampling` package). When interpreting Bayes factors for point hypotheses, make sure that your priors are reasonable and carefully chosen, as the result will depend heavily on the priors. In particular, avoid using default priors. Additionally, note that point hypotheses that scope over more than one parameter (e.g., when testing equality between two parameters) involve random sampling of the priors over those parameters (to accommodate the assumption that priors for different parameters are independent of each other). This introduces an element of randomness into such hypothesis tests. Consider repeating the test to assure results are sufficiently stable, and use the argument `seed` for reproducibility. Finally, note that, for technical reasons, we cannot sample from priors of certain parameters classes. Most notably, these include overall intercept parameters (prior class "Intercept") as well as group-level coefficients.

For one-sided hypotheses, the `Evid.Ratio` may sometimes be 0 or Inf implying very small or large evidence, respectively, in favor of the tested hypothesis. For one-sided hypotheses pairs, this basically means that all posterior draws are on the same side of the value dividing the two hypotheses. In that sense, instead of 0 or Inf, you may rather read it as `Evid.Ratio` smaller 1 / S or greater S, respectively, where S denotes the number of posterior draws used in the computations.

The argument `alpha` specifies the size of the credible interval (i.e., Bayesian confidence interval). For instance, if we tested a two-sided hypothesis and set `alpha = 0.05` (5%) an, the credible interval will contain $1 - \alpha/2 = 0.95$ (95%) of the posterior values. Hence, $\alpha * 100\%$ of the posterior values will lie outside of the credible interval. Although this allows testing of hypotheses in a similar manner as in the frequentist null-hypothesis testing framework, we strongly argue against using arbitrary cutoffs (e.g., $p < .05$) to determine the 'existence' of an effect.

Value

A `brmshypothesis` object.

See Also

[brmshypothesis](#)

Examples

```
## Not run:
## define priors
prior <- c(set_prior("normal(0,2)", class = "b"),
           set_prior("student_t(10,0,1)", class = "sigma"),
           set_prior("student_t(10,0,1)", class = "sd"))

## fit a linear mixed effects models
fit <- brm(time ~ age + sex + disease + (1 + age|patient),
            data = kidney, family = lognormal(),
            prior = prior, sample_prior = "yes",
            control = list(adapt_delta = 0.95))
```

```

## perform two-sided hypothesis testing
(hyp1 <- hypothesis(fit, "sexfemale = age + diseasePKD"))
plot(hyp1)
hypothesis(fit, "exp(age) - 3 = 0", alpha = 0.01)

## perform one-sided hypothesis testing
hypothesis(fit, "diseasePKD + diseaseGN - 3 < 0")

hypothesis(fit, "age < Intercept",
           class = "sd", group = "patient")

## test the amount of random intercept variance on all variance
h <- paste("sd_patient_Intercept^2 / (sd_patient_Intercept^2 +",
           "sd_patient_age^2 + sigma^2) = 0")
(hyp2 <- hypothesis(fit, h, class = NULL))
plot(hyp2)

## test more than one hypothesis at once
h <- c("diseaseGN = diseaseAN", "2 * diseaseGN - diseasePKD = 0")
(hyp3 <- hypothesis(fit, h))
plot(hyp3, ignore_prior = TRUE)

## compute hypotheses for all levels of a grouping factor
hypothesis(fit, "age = 0", scope = "coef", group = "patient")

## use the default method
dat <- as.data.frame(fit)
str(dat)
hypothesis(dat, "b_age > 0")

## End(Not run)

```

Description

Ezzet and Whitehead (1991) analyze data from a two-treatment, two-period crossover trial to compare 2 inhalation devices for delivering the drug salbutamol in 286 asthma patients. Patients were asked to rate the clarity of leaflet instructions accompanying each device, using a 4-point ordinal scale.

Usage

inhaler

Format

A data frame of 572 observations containing information on the following 5 variables.

- subject** The subject number
- rating** The rating of the inhaler instructions on a scale ranging from 1 to 4
- treat** A contrast to indicate which of the two inhaler devices was used
- period** A contrast to indicate the time of administration
- carry** A contrast to indicate possible carry over effects

Source

Ezzet, F., & Whitehead, J. (1991). A random effects model for ordinal responses from a crossover trial. *Statistics in Medicine*, 10(6), 901-907.

Examples

```
## Not run:
## ordinal regression with family "sratio"
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler, family = sratio(),
             prior = set_prior("normal(0,5)"))
summary(fit1)
plot(fit1)

## ordinal regression with family "cumulative"
## and random intercept over subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler, family = cumulative(),
             prior = set_prior("normal(0,5)"))
summary(fit2)
plot(fit2)

## End(Not run)
```

inits.brmsfit

Extract Initial Values Used for Each Chain

Description

Extract the initial values used by Stan for each chain. This is currently only available if the model was fitted with the rstan backend.

Usage

```
## S3 method for class 'brmsfit'
inits(x, ...)

inits(x, ...)
```

Arguments

- x A `brmsfit` object.
- ... Currently ignored.

Value

The initial values (either user-specified or generated randomly) for all chains. This is a list with one component per chain. Each component is a named list containing the initial values for each parameter for the corresponding chain.

Description

Density, distribution function, and random generation for the inverse Gaussian distribution with location `mu`, and shape `shape`.

Usage

```
dinv_gaussian(x, mu = 1, shape = 1, log = FALSE)
pinv_gaussian(q, mu = 1, shape = 1, lower.tail = TRUE, log.p = FALSE)
rinv_gaussian(n, mu = 1, shape = 1)
```

Arguments

- x, q Vector of quantiles.
- mu Vector of locations.
- shape Vector of shapes.
- log Logical; If TRUE, values are returned on the log scale.
- lower.tail Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
- log.p Logical; If TRUE, values are returned on the log scale.
- n Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

inv_logit_scaled	<i>Scaled inverse logit-link</i>
------------------	----------------------------------

Description

Computes $\text{inv_logit}(x) * (\text{ub} - \text{lb}) + \text{lb}$

Usage

```
inv_logit_scaled(x, lb = 0, ub = 1)
```

Arguments

- | | |
|----|------------------------------|
| x | A numeric or complex vector. |
| lb | Lower bound defaulting to 0. |
| ub | Upper bound defaulting to 1. |

Value

A numeric or complex vector between lb and ub.

is.brmsfit	<i>Checks if argument is a brmsfit object</i>
------------	---

Description

Checks if argument is a brmsfit object

Usage

```
is.brmsfit(x)
```

Arguments

- | | |
|---|-------------|
| x | An R object |
|---|-------------|

is.brmsfit_multiple *Checks if argument is a brmsfit_multiple object*

Description

Checks if argument is a brmsfit_multiple object

Usage

```
is.brmsfit_multiple(x)
```

Arguments

x An R object

is.brmsformula *Checks if argument is a brmsformula object*

Description

Checks if argument is a brmsformula object

Usage

```
is.brmsformula(x)
```

Arguments

x An R object

is.brmsprior *Checks if argument is a brmsprior object*

Description

Checks if argument is a brmsprior object

Usage

```
is.brmsprior(x)
```

Arguments

x An R object

is.brmsterms	<i>Checks if argument is a brmsterms object</i>
--------------	---

Description

Checks if argument is a `brmsterms` object

Usage

```
is.brmsterms(x)
```

Arguments

x An R object

See Also

`brmsterms`

is.cor_brms	<i>Check if argument is a correlation structure</i>
-------------	---

Description

Check if argument is one of the correlation structures used in `brms`.

Usage

```
is.cor_brms(x)
```

```
is.cor_arma(x)
```

```
is.cor_cosy(x)
```

```
is.cor_sar(x)
```

```
is.cor_car(x)
```

```
is.cor_fixed(x)
```

Arguments

x An R object.

is.mvbrmsformula *Checks if argument is a mvbrmsformula object*

Description

Checks if argument is a `mvbrmsformula` object

Usage

```
is.mvbrmsformula(x)
```

Arguments

x An R object

is.mvbrmsterms *Checks if argument is a mvbrmsterms object*

Description

Checks if argument is a `mvbrmsterms` object

Usage

```
is.mvbrmsterms(x)
```

Arguments

x An R object

See Also

[brmsterms](#)

kfold.brmsfit	<i>K-Fold Cross-Validation</i>
----------------------	--------------------------------

Description

Perform exact K-fold cross-validation by refitting the model K times each leaving out one- K th of the original data. Folds can be run in parallel using the **future** package.

Usage

```
## S3 method for class 'brmsfit'
kfold(
  x,
  ...,
  K = 10,
  Ksub = NULL,
  folds = NULL,
  group = NULL,
  joint = FALSE,
  compare = TRUE,
  resp = NULL,
  model_names = NULL,
  save_fits = FALSE,
  recompile = NULL,
  future_args = list()
)
```

Arguments

x	A <code>brmsfit</code> object.
...	Further arguments passed to <code>brm</code> and <code>log_lik</code> .
K	The number of subsets of equal (if possible) size into which the data will be partitioned for performing K -fold cross-validation. The model is refit K times, each time leaving out one of the K subsets. If K is equal to the total number of observations in the data then K -fold cross-validation is equivalent to exact leave-one-out cross-validation.
Ksub	Optional number of subsets (of those subsets defined by K) to be evaluated. If <code>NULL</code> (the default), K -fold cross-validation will be performed on all subsets. If <code>Ksub</code> is a single integer, <code>Ksub</code> subsets (out of all K) subsets will be randomly chosen. If <code>Ksub</code> consists of multiple integers or a one-dimensional array (created via <code>as.array</code>) potentially of length one, the corresponding subsets will be used. This argument is primarily useful, if evaluation of all subsets is infeasible for some reason.
folds	Determines how the subsets are being constructed. Possible values are <code>NULL</code> (the default), "stratified", "grouped", or "loo". May also be a vector of length equal to the number of observations in the data. Alters the way <code>group</code> is handled. More information is provided in the 'Details' section.

group	Optional name of a grouping variable or factor in the model. What exactly is done with this variable depends on argument folds. More information is provided in the 'Details' section.
joint	Indicates which observations' log likelihoods shall be considered jointly in the ELPD computation. If "obs" or FALSE (the default), each observation is considered separately. This enables comparability of <i>kfold</i> with <i>loo</i> . If "fold" or TRUE, the joint log likelihoods per fold are used. If "group", the joint log likelihoods per group within folds are used (only available if argument group is specified).
compare	A flag indicating if the information criteria of the models should be compared to each other via loo_compare .
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
model_names	If NULL (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.
save_fits	If TRUE, a component fits is added to the returned object to store the cross-validated brmsfit objects and the indices of the omitted observations for each fold. Defaults to FALSE.
recompile	Logical, indicating whether the Stan model should be recompiled. This may be necessary if you are running <i>reloo</i> on another machine than the one used to fit the model.
future_args	A list of further arguments passed to future for additional control over parallel execution if activated.

Details

The *kfold* function performs exact K -fold cross-validation. First the data are partitioned into K folds (i.e. subsets) of equal (or as close to equal as possible) size by default. Then the model is refit K times, each time leaving out one of the K subsets. If K is equal to the total number of observations in the data then K -fold cross-validation is equivalent to exact leave-one-out cross-validation (to which *loo* is an efficient approximation). The *compare_ic* function is also compatible with the objects returned by *kfold*.

The subsets can be constructed in multiple different ways:

- If both *folds* and *group* are NULL, the subsets are randomly chosen so that they have equal (or as close to equal as possible) size.
- If *folds* is NULL but *group* is specified, the data is split up into subsets, each time omitting all observations of one of the factor levels, while ignoring argument *K*.
- If *folds* = "stratified" the subsets are stratified after *group* using [loo::kfold_split_stratified](#).
- If *folds* = "grouped" the subsets are split by *group* using [loo::kfold_split_grouped](#).
- If *folds* = "loo" exact leave-one-out cross-validation will be performed and *K* will be ignored. Further, if *group* is specified, all observations corresponding to the factor level of the currently predicted single value are omitted. Thus, in this case, the predicted values are only a subset of the omitted ones.

- If `folds` is a numeric vector, it must contain one element per observation in the data. Each element of the vector is an integer in `1:K` indicating to which of the `K` folds the corresponding observation belongs. There are some convenience functions available in the `loo` package that create integer vectors to use for this purpose (see the Examples section below and also the [kfold-helpers](#) page).

When running `kfold` on a `brmsfit` created with the `cmdstanr` backend in a different R session, several recompilations will be triggered because by default, `cmdstanr` writes the model executable to a temporary directory. To avoid that, set option `"cmdstanr_write_stan_file_dir"` to a non-temporary path of your choice before creating the original `brmsfit` (see section 'Examples' below). By default, this method uses `sample_new_levels = "gaussian"` to sample parameter values for new grouping-factor levels (see also [prepare_predictions](#)). This default will fail for models with non-Gaussian group-level effects. In this case, we recommend setting `sample_new_levels = "uncertainty"`.

Value

`kfold` returns an object that has a similar structure as the objects returned by the `loo` and `waic` methods and can be used with the same post-processing functions.

Parallelization with multiple CPU cores

`brms` can make use of multiple CPU cores in parallel to speed up computations in various ways. For efficient use of the available resources it is recommended to only use parallelism to an extend such that the available physical CPUs are not oversubscribed. For example, when you have 8 CPU cores locally available, then you may consider to run 4 chains with 2 threads per chain for best performance if you happen to just run a single model. In case you run a simulation study which requires to run many times a given model, then neither chain nor within-chain parallelization is advisable as the computational resources are already exhausted by the simulation study and any further parallelization beyond the simulation study itself will in fact slow down the overall runtime. Please be aware that for historical reasons the nomenclature of the arguments is possibly confusing. The `cores` argument refers to running different chains in parallel and the within-chain parallelization will allocate for each chain as many threads as requested. The requested threads therefore increase the use of overall CPUs in a multiplicative way.

For more advanced parallelization (including beyond single model fits), `brms` also integrates with the `future` package. Importantly, this enables seamless integration with the `mirai` parallelization framework through the use of the `future.mirai` adapter. With `mirai` local and remote machines can be used in a fully transparent manner to the user. This includes the possibility to use large number of remote machines running in the context of a computer cluster, which are managed with queuing systems. Please refer to the section on distributed computing of [mirai::daemons](#).

See Also

[loo](#), [reloo](#)

Examples

```
## Not run:
fit1 <- brm(count ~ zAge + zBase * Trt + (1|patient) + (1|obs),
```

```

    data = epilepsy, family = poisson())
# throws warning about some pareto k estimates being too high
(loo1 <- loo(fit1))
# perform 10-fold cross validation
(kfold1 <- kfold(fit1, chains = 1))

# use joint likelihoods per fold for ELPD evaluation
kfold(fit1, chains = 1, joint = "fold")

# use the future package for parallelization of models
# that is to fit models belonging to different folds in parallel
library(future)
plan(multisession, workers = 4)
kfold(fit1, chains = 1)
plan(sequential)

## to avoid recompilations when running kfold() on a 'cmdstanr'-backend fit
## in a fresh R session, set option 'cmdstanr_write_stan_file_dir' before
## creating the initial 'brmsfit'
## CAUTION: the following code creates some files in the current working
## directory: two 'model_<hash>.stan' files, one 'model_<hash>(.exe)'
## executable, and one 'fit_cmdstanr_<some_number>.rds' file
set.seed(7)
fname <- paste0("fit_cmdstanr_", sample.int(.Machine$integer.max, 1))
options(cmdstanr_write_stan_file_dir = getwd())
fit_cmdstanr <- brm(rate ~ conc + state, data = Puromycin,
                     backend = "cmdstanr", file = fname)

# now restart the R session and run the following (after attaching 'brms')
set.seed(7)
fname <- paste0("fit_cmdstanr_", sample.int(.Machine$integer.max, 1))
fit_cmdstanr <- brm(rate ~ conc + state,
                     data = Puromycin,
                     backend = "cmdstanr",
                     file = fname)
kfold_cmdstanr <- kfold(fit_cmdstanr, K = 2)

## End(Not run)

```

kfold_predict*Predictions from K-Fold Cross-Validation***Description**

Compute and evaluate predictions after performing K-fold cross-validation via [kfold](#).

Usage

```
kfold_predict(x, method = "posterior_predict", resp = NULL, ...)
```

Arguments

- x Object of class 'kfold' computed by [kfold](#). For kfold_predict to work, the fitted model objects need to have been stored via argument save_fits of [kfold](#).
- method Method used to obtain predictions. Can be set to "posterior_predict" (the default), "posterior_epred", or "posterior_linpred". For more details, see the respective function documentations.
- resp Optional names of response variables. If specified, predictions are performed only for the specified response variables.
- ... Further arguments passed to [prepare_predictions](#) that control several aspects of data validation and prediction.

Value

A list with two slots named 'y' and 'yrep'. Slot y contains the vector of observed responses. Slot yrep contains the matrix of predicted responses, with rows being posterior draws and columns being observations.

See Also

[kfold](#)

Examples

```
## Not run:
fit <- brm(count ~ zBase * Trt + (1|patient),
            data = epilepsy, family = poisson())

# perform k-fold cross validation
(kf <- kfold(fit, save_fits = TRUE, chains = 1))

# define a loss function
rmse <- function(y, yrep) {
  yrep_mean <- colMeans(yrep)
  sqrt(mean((yrep_mean - y)^2))
}

# predict responses and evaluate the loss
kfp <- kfold_predict(kf)
rmse(y = kfp$y, yrep = kfp$yrep)

## End(Not run)
```

kidney

Infections in kidney patients

Description

This dataset, originally discussed in McGilchrist and Aisbett (1991), describes the first and second (possibly right censored) recurrence time of infection in kidney patients using portable dialysis equipment. In addition, information on the risk variables age, sex and disease type is provided.

Usage

kidney

Format

A data frame of 76 observations containing information on the following 7 variables.

time The time to first or second recurrence of the infection, or the time of censoring

recur A factor of levels 1 or 2 indicating if the infection recurred for the first or second time for this patient

censored Either 0 or 1, where 0 indicates no censoring of recurrence time and 1 indicates right censoring

patient The patient number

age The age of the patient

sex The sex of the patient

disease A factor of levels other , GN, AN, and PKD specifying the type of disease

Source

McGilchrist, C. A., & Aisbett, C. W. (1991). Regression with frailty in survival analysis. *Biometrics*, 47(2), 461-466.

Examples

```
## Not run:
## performing survival analysis using the "weibull" family
fit1 <- brm(time | cens(censored) ~ age + sex + disease,
             data = kidney, family = weibull, init = "0")
summary(fit1)
plot(fit1)

## adding random intercepts over patients
fit2 <- brm(time | cens(censored) ~ age + sex + disease + (1|patient),
             data = kidney, family = weibull(), init = "0",
             prior = set_prior("cauchy(0,2)", class = "sd"))
summary(fit2)
plot(fit2)
```

```
## End(Not run)
```

lasso

(Defunct) Set up a lasso prior in brms

Description

This functionality is no longer supported as of brms version 2.19.2. Please use the [horseshoe](#) or [R2D2](#) shrinkage priors instead.

Usage

```
lasso(df = 1, scale = 1)
```

Arguments

<code>df</code>	Degrees of freedom of the chi-square prior of the inverse tuning parameter. Defaults to 1.
<code>scale</code>	Scale of the lasso prior. Defaults to 1.

Value

An error indicating that the lasso prior is no longer supported.

References

Park, T., & Casella, G. (2008). The Bayesian Lasso. *Journal of the American Statistical Association*, 103(482), 681-686.

See Also

[set_prior](#), [horseshoe](#), [R2D2](#)

launch_shinystan.brmsfit

Interface to shinystan

Description

Provide an interface to **shinystan** for models fitted with **brms**

Usage

```
launch_shinystan.brmsfit(object, rstudio =getOption("shinystan.rstudio"), ...)
```

Arguments

object	A fitted model object typically of class brmsfit .
rstudio	Only relevant for RStudio users. The default (rstudio=FALSE) is to launch the app in the default web browser rather than RStudio's pop-up Viewer. Users can change the default to TRUE by setting the global option <code>options(shinystan.rstudio = TRUE)</code> .
...	Optional arguments to pass to <code>runApp</code>

Value

An S4 shinystan object

See Also

[launch_shinystan](#)

Examples

```
## Not run:  
fit <- brm(rating ~ treat + period + carry + (1|subject),  
           data = inhaler, family = "gaussian")  
launch_shinystan(fit)  
  
## End(Not run)
```

LogisticNormal*The (Multivariate) Logistic Normal Distribution***Description**

Density function and random generation for the (multivariate) logistic normal distribution with latent mean vector μ and covariance matrix Σ .

Usage

```
dlogistic_normal(x, mu, Sigma, refcat = 1, log = FALSE, check = FALSE)

rlogistic_normal(n, mu, Sigma, refcat = 1, check = FALSE)
```

Arguments

x	Vector or matrix of quantiles. If x is a matrix, each row is taken to be a quantile.
μ	Mean vector with length equal to the number of dimensions.
Σ	Covariance matrix.
refcat	A single integer indicating the reference category. Defaults to 1.
log	Logical; If TRUE, values are returned on the log scale.
check	Logical; Indicates whether several input checks should be performed. Defaults to FALSE to improve efficiency.
n	Number of draws to sample from the distribution.

logit_scaled*Scaled logit-link***Description**

Computes $\text{logit}((x - \text{lb}) / (\text{ub} - \text{lb}))$

Usage

```
logit_scaled(x, lb = 0, ub = 1)
```

Arguments

x	A numeric or complex vector.
lb	Lower bound defaulting to 0.
ub	Upper bound defaulting to 1.

Value

A numeric or complex vector.

logm1*Logarithm with a minus one offset.***Description**

Computes $\log(x - 1)$.

Usage

```
logm1(x, base = exp(1))
```

Arguments

- | | |
|-------------------|---|
| <code>x</code> | A numeric or complex vector. |
| <code>base</code> | A positive or complex number: the base with respect to which logarithms are computed. Defaults to $e = \exp(1)$. |

log_lik.brmsfit*Compute the Pointwise Log-Likelihood***Description**

Compute the Pointwise Log-Likelihood

Usage

```
## S3 method for class 'brmsfit'
log_lik(
  object,
  newdata = NULL,
  re_formula = NULL,
  resp = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  pointwise = FALSE,
  combine = TRUE,
  add_point_estimate = FALSE,
  cores = NULL,
  ...
)
```

Arguments

object	A fitted model object of class <code>brmsfit</code> .
newdata	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
re_formula	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
ndraws	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
draw_ids	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
pointwise	A flag indicating whether to compute the full log-likelihood matrix at once (the default), or just return the likelihood function along with all data and draws required to compute the log-likelihood separately for each observation. The latter option is rarely useful when calling <code>log_liks</code> directly, but rather when computing <code>waic</code> or <code>loo</code> .
combine	Only relevant in multivariate models. Indicates if the log-likelihoods of the submodels should be combined per observation (i.e. added together; the default) or if the log-likelihoods should be returned separately.
add_point_estimate	For internal use only. Ensures compatibility with the <code>loo_subsample</code> method.
cores	Number of cores (defaults to 1). On non-Windows systems, this argument can be set globally via the <code>mc.cores</code> option.
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Details

NA values within factors in `newdata`, are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding.

In multilevel models, it is possible to allow new levels of grouping factors to be used in the predictions. This can be controlled via argument `allow_new_levels`. New levels can be sampled in multiple ways, which can be controlled via argument `sample_new_levels`. Both of these arguments are documented in `prepare_predictions` along with several other useful arguments to control specific aspects of the predictions.

Value

Usually, an $S \times N$ matrix containing the pointwise log-likelihood draws, where S is the number of draws and N is the number of observations in the data. For multivariate models and if `combine` is

FALSE, an S x N x R array is returned, where R is the number of response variables. If `pointwise` = TRUE, the output is a function with a `draws` attribute containing all relevant data and posterior draws.

loo.brmsfit*Efficient approximate leave-one-out cross-validation (LOO)***Description**

Perform approximate leave-one-out cross-validation based on the posterior likelihood using the **loo** package. For more details see [loo](#).

Usage

```
## S3 method for class 'brmsfit'
loo(
  x,
  ...,
  compare = TRUE,
  resp = NULL,
  pointwise = FALSE,
  moment_match = FALSE,
  reloo = FALSE,
  k_threshold = 0.7,
  save_psis = FALSE,
  moment_match_args = list(),
  reloo_args = list(),
  model_names = NULL
)
```

Arguments

- `x` A `brmsfit` object.
- `...` More `brmsfit` objects or further arguments passed to the underlying post-processing functions. In particular, see [prepare_predictions](#) for further supported arguments.
- `compare` A flag indicating if the information criteria of the models should be compared to each other via [loo_compare](#).
- `resp` Optional names of response variables. If specified, predictions are performed only for the specified response variables.
- `pointwise` A flag indicating whether to compute the full log-likelihood matrix at once or separately for each observation. The latter approach is usually considerably slower but requires much less working memory. Accordingly, if one runs into memory issues, `pointwise` = TRUE is the way to go.

moment_match	Logical; Indicate whether <code>loo_moment_match</code> should be applied on problematic observations. Defaults to FALSE. For most models, moment matching will only work if you have set <code>save_pars = save_pars(all = TRUE)</code> when fitting the model with <code>brm</code> . See <code>loo_moment_match.brmsfit</code> for more details.
reloo	Logical; Indicate whether <code>reloo</code> should be applied on problematic observations. Defaults to FALSE.
k_threshold	The Pareto k threshold for which observations <code>loo_moment_match</code> or <code>reloo</code> is applied if argument <code>moment_match</code> or <code>reloo</code> is TRUE. Defaults to 0.7. See <code>pareto_k_ids</code> for more details.
save_psis	Should the "psis" object created internally be saved in the returned object? For more details see <code>loo</code> .
moment_match_args	Optional named list of additional arguments passed to <code>loo_moment_match</code> .
reloo_args	Optional named list of additional arguments passed to <code>reloo</code> . This can be useful, among others, to control how many chains, iterations, etc. to use for the fitted sub-models.
model_names	If NULL (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Details

See `loo_compare` for details on model comparisons. For `brmsfit` objects, LOO is an alias of `loo`. Use method `add_criterion` to store information criteria in the fitted model object for later usage.

Value

If just one object is provided, an object of class `loo`. If multiple objects are provided, an object of class `loolist`.

References

- Vehtari, A., Gelman, A., & Gabry J. (2016). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. In *Statistics and Computing*, doi:10.1007/s11222-016-9696-4. arXiv preprint arXiv:1507.04544.
- Gelman, A., Hwang, J., & Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing*, 24, 997-1016.
- Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *The Journal of Machine Learning Research*, 11, 3571-3594.

Examples

```
## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler)
(loo1 <- loo(fit1))
```

```
# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler)
(loo2 <- loo(fit2))

# compare both models
loo_compare(loo1, loo2)

## End(Not run)
```

loo_compare.brmsfit *Model comparison with the **loo** package*

Description

For more details see [loo_compare](#).

Usage

```
## S3 method for class 'brmsfit'
loo_compare(x, ..., criterion = c("loo", "waic", "kfold"), model_names = NULL)
```

Arguments

<code>x</code>	A <code>brmsfit</code> object.
<code>...</code>	More <code>brmsfit</code> objects.
<code>criterion</code>	The name of the criterion to be extracted from <code>brmsfit</code> objects.
<code>model_names</code>	If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Details

All `brmsfit` objects should contain precomputed criterion objects. See [add_criterion](#) for more help.

Value

An object of class "compare.loo".

Examples

```
## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler)
fit1 <- add_criterion(fit1, "waic")
```

```

# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler)
fit2 <- add_criterion(fit2, "waic")

# compare both models
loo_compare(fit1, fit2, criterion = "waic")

## End(Not run)

```

loo_model_weights.brmsfit*Model averaging via stacking or pseudo-BMA weighting.***Description**

Compute model weights for `brmsfit` objects via stacking or pseudo-BMA weighting. For more details, see [loo::loo_model_weights](#).

Usage

```
## S3 method for class 'brmsfit'
loo_model_weights(x, ..., model_names = NULL)
```

Arguments

- | | |
|--------------------------|--|
| <code>x</code> | A <code>brmsfit</code> object. |
| <code>...</code> | More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments. |
| <code>model_names</code> | If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names. |

Value

A named vector of model weights.

Examples

```

## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler, family = "gaussian")
# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler, family = "gaussian")
loo_model_weights(fit1, fit2)

```

```
## End(Not run)
```

loo_moment_match.brmsfit

Moment matching for efficient approximate leave-one-out cross-validation

Description

Moment matching for efficient approximate leave-one-out cross-validation (LOO-CV). See [loo_moment_match](#) for more details.

Usage

```
## S3 method for class 'brmsfit'
loo_moment_match(
  x,
  loo = NULL,
  k_threshold = 0.7,
  newdata = NULL,
  resp = NULL,
  check = TRUE,
  recompile = FALSE,
  ...
)

## S3 method for class 'loo'
loo_moment_match(x, fit, ...)
```

Arguments

<code>x</code>	An R object of class <code>brmsfit</code> or <code>loo</code> depending on the method.
<code>loo</code>	An R object of class <code>loo</code> . If <code>NULL</code> , <code>brms</code> will try to extract a precomputed <code>loo</code> object from the fitted model, added there via add_criterion .
<code>k_threshold</code>	The Pareto k threshold for which observations moment matching is applied. Defaults to <code>0.7</code> . See pareto_k_ids for more details.
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.

check	Logical; If TRUE (the default), some checks check are performed if the loo object was generated from the brmsfit object passed to argument fit.
recompile	Logical, indicating whether the Stan model should be recompiled. This may be necessary if you are running moment matching on another machine than the one used to fit the model. No recompilation is done by default.
...	Further arguments passed to the underlying methods. Additional arguments initially passed to <code>loo</code> , for example, newdata or resp need to be passed again to <code>loo_moment_match</code> in order for the latter to work correctly.
fit	An R object of class <code>brmsfit</code> .

Details

The moment matching algorithm requires draws of all variables defined in Stan's parameters block to be saved. Otherwise `loo_moment_match` cannot be computed. Thus, please set `save_pars = save_pars(all = TRUE)` in the call to `brm`, if you are planning to apply `loo_moment_match` to your models.

Value

An updated object of class `loo`.

References

Paananen, T., Piironen, J., Buerkner, P.-C., Vehtari, A. (2021). Implicitly Adaptive Importance Sampling. *Statistics and Computing*.

Examples

```
## Not run:
fit1 <- brm(count ~ zAge + zBase * Trt + (1|patient),
             data = epilepsy, family = poisson(),
             save_pars = save_pars(all = TRUE))

# throws warning about some pareto k estimates being too high
(loo1 <- loo(fit1))

# no more warnings after moment matching
(mmlloo1 <- loo_moment_match(fit1, loo = loo1))

## End(Not run)
```

`loo_predict.brmsfit` *Compute Weighted Expectations Using LOO*

Description

These functions are wrappers around the [E_loo](#) function of the **loo** package.

Usage

```
## S3 method for class 'brmsfit'
loo_predict(
  object,
  type = c("mean", "var", "quantile"),
  probs = 0.5,
  psis_object = NULL,
  resp = NULL,
  ...
)

## S3 method for class 'brmsfit'
loo_epred(
  object,
  type = c("mean", "var", "quantile"),
  probs = 0.5,
  psis_object = NULL,
  resp = NULL,
  ...
)

loo_epred(object, ...)

## S3 method for class 'brmsfit'
loo_linpred(
  object,
  type = c("mean", "var", "quantile"),
  probs = 0.5,
  psis_object = NULL,
  resp = NULL,
  ...
)

## S3 method for class 'brmsfit'
loo_predictive_interval(object, prob = 0.9, psis_object = NULL, ...)
```

Arguments

`object` An object of class `brmsfit`.

type	The statistic to be computed on the results. Can by either "mean" (default), "var", or "quantile".
probs	A vector of quantiles to compute. Only used if type = quantile.
psis_object	An optional object returned by <code>psis</code> . If <code>psis_object</code> is missing then <code>psis</code> is executed internally, which may be time consuming for models fit to very large datasets.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
...	Optional arguments passed to the underlying methods that is <code>log_lik</code> , as well as <code>posterior_predict</code> , <code>posterior_epred</code> or <code>posterior_linpred</code> .
prob	For <code>loo_predictive_interval</code> , a scalar in (0, 1) indicating the desired probability mass to include in the intervals. The default is <code>prob = 0.9</code> (90% intervals).

Value

`loo_predict`, `loo_epred`, `loo_linpred`, and `loo_predictive_interval` all return a matrix with one row per observation and one column per summary statistic as specified by arguments `type` and `probs`. In multivariate or categorical models a third dimension is added to represent the response variables or categories, respectively.

`loo_predictive_interval(..., prob = p)` is equivalent to `loo_predict(..., type = "quantile", probs = c(a, 1-a))` with $a = (1 - p)/2$.

Examples

```
## Not run:
## data from help("lm")
ctl <- c(4.17,5.58,5.18,6.11,4.50,4.61,5.17,4.53,5.33,5.14)
trt <- c(4.81,4.17,4.41,3.59,5.87,3.83,6.03,4.89,4.32,4.69)
d <- data.frame(
  weight = c(ctl, trt),
  group = gl(2, 10, 20, labels = c("Ctl", "Trt"))
)
fit <- brm(weight ~ group, data = d)
loo_predictive_interval(fit, prob = 0.8)

## optionally log-weights can be pre-computed and reused
psis <- loo:::psis(-log_lik(fit), cores = 2)
loo_predictive_interval(fit, prob = 0.8, psis_object = psis)
loo_predict(fit, type = "var", psis_object = psis)
loo_epred(fit, type = "var", psis_object = psis)

## End(Not run)
```

loo_R2.brmsfit*Compute a LOO-adjusted R-squared for regression models*

Description

Compute a LOO-adjusted R-squared for regression models

Usage

```
## S3 method for class 'brmsfit'
loo_R2(
  object,
  resp = NULL,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  seed = NULL,
  args_epred = list(),
  args_loglik = list(),
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>summary</code>	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .
<code>robust</code>	If <code>FALSE</code> (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If <code>TRUE</code> , the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is <code>TRUE</code> .
<code>probs</code>	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is <code>TRUE</code> .
<code>seed</code>	Optional integer used to initialize the random number generator.
<code>args_epred</code>	A named list of further arguments passed only to <code>posterior_epred</code> .
<code>args_loglik</code>	A named list of further arguments passed only to <code>log_lik</code> .
<code>...</code>	Further arguments passed to both <code>posterior_epred</code> and <code>log_lik</code> .

Value

If `summary = TRUE`, an $M \times C$ matrix is returned ($M =$ number of response variables and $C = \text{length}(\text{probs}) + 2$) containing Bayesian bootstrap based summary statistics of the LOO-adjusted

R-squared values. If `summary = FALSE`, the Bayesian bootstrap draws of the LOO-adjusted R-squared values are returned in an S x M matrix (S is the number of draws).

@details LOO-R2 uses LOO residuals and is defined as $1 - \text{Var}_{\text{loo-res}}/\text{Var}_y$, with

$$\text{Var}_y = V_{n=1}^N y_n, \text{ and } \text{Var}_{\text{loo-res}} = V_{n=1}^N \hat{e}_{\text{loo},n},$$

where $\hat{e}_{\text{loo},n} = y_n - \hat{y}_{\text{loo},n}$. Bayesian bootstrap is used to draw from the approximated uncertainty distribution as described by Vehtari and Lampinen (2002).

References

Vehtari and Lampinen (2002). Bayesian model assessment and comparison using cross-validation predictive densities. *Neural Computation*, 14(10):2439-2468.

Examples

```
## Not run:
fit <- brm(mpg ~ wt + cyl, data = mtcars)
summary(fit)
loo_R2(fit)

# compute R2 with new data
nd <- data.frame(mpg = c(10, 20, 30), wt = c(4, 3, 2), cyl = c(8, 6, 4))
loo_R2(fit, newdata = nd)

## End(Not run)
```

`loo_subsample.brmsfit` *Efficient approximate leave-one-out cross-validation (LOO) using subsampling*

Description

Efficient approximate leave-one-out cross-validation (LOO) using subsampling

Usage

```
## S3 method for class 'brmsfit'
loo_subsample(x, ..., compare = TRUE, resp = NULL, model_names = NULL)
```

Arguments

- x A `brmsfit` object.
- ... More `brmsfit` objects or further arguments passed to the underlying post-processing functions. In particular, see [prepare_predictions](#) for further supported arguments.

<code>compare</code>	A flag indicating if the information criteria of the models should be compared to each other via loo_compare .
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>model_names</code>	If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Details

More details can be found on [loo_subsample](#).

Examples

```
## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler)
(loo1 <- loo_subsample(fit1))

# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler)
(loo2 <- loo_subsample(fit2))

# compare both models
loo_compare(loo1, loo2)

## End(Not run)
```

`loss`

Cumulative Insurance Loss Payments

Description

This dataset, discussed in Gesmann & Morris (2020), contains cumulative insurance loss payments over the course of ten years.

Usage

`loss`

Format

A data frame of 55 observations containing information on the following 4 variables.

AY Origin year of the insurance (1991 to 2000)

dev Deviation from the origin year in months

cum Cumulative loss payments

premium Achieved premiums for the given origin year

Source

Gesmann M. & Morris J. (2020). Hierarchical Compartmental Reserving Models. *CAS Research Papers*.

Examples

```
## Not run:
# non-linear model to predict cumulative loss payments
fit_loss <- brm(
  bf(cum ~ ult * (1 - exp(-(dev/theta)^omega)),
    ult ~ 1 + (1|AY), omega ~ 1, theta ~ 1,
    nl = TRUE),
  data = loss, family = gaussian(),
  prior = c(
    prior(normal(5000, 1000), nlpar = "ult"),
    prior(normal(1, 2), nlpar = "omega"),
    prior(normal(45, 10), nlpar = "theta")
  ),
  control = list(adapt_delta = 0.9)
)

# basic summaries
summary(fit_loss)
conditional_effects(fit_loss)

# plot predictions per origin year
conditions <- data.frame(AY = unique(loss$AY))
rownames(conditions) <- unique(loss$AY)
me_loss <- conditional_effects(
  fit_loss, conditions = conditions,
  re_formula = NULL, method = "predict"
)
plot(me_loss, ncol = 5, points = TRUE)

## End(Not run)
```

Description

Set up a moving average (MA) term of order q in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with MA terms.

Usage

```
ma(time = NA, gr = NA, q = 1, cov = FALSE)
```

Arguments

<code>time</code>	An optional time variable specifying the time ordering of the observations. By default, the existing order of the observations in the data is used.
<code>gr</code>	An optional grouping variable. If specified, the correlation structure is assumed to apply only to observations within the same grouping level.
<code>q</code>	A non-negative integer specifying the moving average (MA) order of the ARMA structure. Default is 1.
<code>cov</code>	A flag indicating whether ARMA effects should be estimated by means of residual covariance matrices. This is currently only possible for stationary ARMA effects of order 1. If the model family does not have natural residuals, latent residuals are added automatically. If FALSE (the default), a regression formulation is used that is considerably faster and allows for ARMA effects of order higher than 1 but is only available for gaussian models and some of its generalizations.

Value

An object of class 'arma_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#), [arma](#), [ar](#)

Examples

```
## Not run:
data("LakeHuron")
LakeHuron <- as.data.frame(LakeHuron)
fit <- brm(x ~ ma(p = 2), data = LakeHuron)
summary(fit)

## End(Not run)
```

Description

This is a helper function to prepare fully crossed conditions primarily for use with the `conditions` argument of [conditional_effects](#). Automatically creates labels for each row in the `cond_` column.

Usage

`make_conditions(x, vars, ...)`

Arguments

- x An R object from which to extract the variables that should be part of the conditions.
- vars Names of the variables that should be part of the conditions.
- ... Arguments passed to [rows2labels](#).

Details

For factor like variables, all levels are used as conditions. For numeric variables, `mean + (-1:1) * SD` are used as conditions.

Value

A `data.frame` where each row indicates a condition.

See Also

[conditional_effects](#), [rows2labels](#)

Examples

```
df <- data.frame(x = c("a", "b"), y = rnorm(10))
make_conditions(df, vars = c("x", "y"))
```

mcmc_plot.brmsfit *MCMC Plots Implemented in bayesplot*

Description

Convenient way to call MCMC plotting functions implemented in the **bayesplot** package.

Usage

```
## S3 method for class 'brmsfit'
mcmc_plot(
  object,
  pars = NA,
  type = "intervals",
  variable = NULL,
  regex = FALSE,
  fixed = FALSE,
  ...
)
mcmc_plot(object, ...)
```

Arguments

object	An R object typically of class <code>brmsfit</code>
pars	Deprecated alias of <code>variable</code> . Names of the parameters to plot, as given by a character vector or a regular expression.
type	The type of the plot. Supported types are (as names) <code>hist</code> , <code>dens</code> , <code>hist_by_chain</code> , <code>dens_overlay</code> , <code>violin</code> , <code>intervals</code> , <code>areas</code> , <code>acf</code> , <code>acf_bar</code> , <code>trace</code> , <code>trace_highlight</code> , <code>scatter</code> , <code>rhat</code> , <code>rhat_hist</code> , <code>neff</code> , <code>neff_hist</code> , <code>nuts_acceptance</code> , <code>nuts_divergence</code> , <code>nuts_stepsize</code> , <code>nuts_treedepth</code> , and <code>nuts_energy</code> . For an overview on the various plot types see MCMC-overview .
variable	Names of the variables (parameters) to plot, as given by a character vector or a regular expression (if <code>regex = TRUE</code>). By default, a hopefully not too large selection of variables is plotted.
regex	Logical; Indicates whether <code>variable</code> should be treated as regular expressions. Defaults to FALSE.
fixed	(Deprecated) Indicates whether parameter names should be matched exactly (TRUE) or treated as regular expressions (FALSE). Default is FALSE and only works with argument <code>pars</code> .
...	Additional arguments passed to the plotting functions. See MCMC-overview for more details.

Details

Also consider using the **shinystan** package available via method `launch_shinystan` in **brms** for flexible and interactive visual analysis.

Value

A `ggplot` object that can be further customized using the **ggplot2** package.

Examples

```
## Not run:
model <- brm(count ~ zAge + zBase * Trt + (1|patient),
               data = epilepsy, family = "poisson")

# plot posterior intervals
mcmc_plot(model)

# only show population-level effects in the plots
mcmc_plot(model, variable = "^b_", regex = TRUE)

# show histograms of the posterior distributions
mcmc_plot(model, type = "hist")

# plot some diagnostics of the sampler
mcmc_plot(model, type = "neff")
mcmc_plot(model, type = "rhat")
```

```
# plot some diagnostics specific to the NUTS sampler
mcmc_plot(model, type = "nuts_acceptance")
mcmc_plot(model, type = "nuts_divergence")

## End(Not run)
```

me*Predictors with Measurement Error in brms Models*

Description

(Soft deprecated) Specify predictors with measurement error. The function does not evaluate its arguments – it exists purely to help set up a model.

Usage

```
me(x, sdx, gr = NULL)
```

Arguments

- x** The variable measured with error.
- sdx** Known measurement error of x treated as standard deviation.
- gr** Optional grouping factor to specify which values of x correspond to the same value of the latent variable. If `NULL` (the default) each observation will have its own value of the latent variable.

Details

For detailed documentation see `help(brmsformula)`. `me` terms are soft deprecated in favor of the more general and consistent `mi` terms. By default, latent noise-free variables are assumed to be correlated. To change that, add `set_mecor(FALSE)` to your model formula object (see examples).

See Also

[brmsformula](#), [brmsformula-helpers](#)

Examples

```
## Not run:
# sample some data
N <- 100
dat <- data.frame(
  y = rnorm(N), x1 = rnorm(N),
  x2 = rnorm(N), sdx = abs(rnorm(N, 1)))
# fit a simple error-in-variables model
fit1 <- brm(y ~ me(x1, sdx) + me(x2, sdx), data = dat,
             save_pars = save_pars(latent = TRUE))
```

```

summary(fit1)

# turn off modeling of correlations
bform <- bf(y ~ me(x1, sdx) + me(x2, sdx)) + set_mecor(FALSE)
fit2 <- brm(bform, data = dat, save_pars = save_pars(latent = TRUE))
summary(fit2)

## End(Not run)

```

Description

Specify predictor term with missing values in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model. For documentation on how to specify missing values in response variables, see [resp_mi](#).

Usage

```
mi(x, idx = NA)
```

Arguments

- | | |
|------------------|--|
| <code>x</code> | The variable containing missing values. |
| <code>idx</code> | An optional variable containing indices of observations in ‘x‘ that are to be used in the model. This is mostly relevant in partially subsetted models (via <code>resp_subset</code>) but may also have other applications that I haven’t thought of. |

Details

For detailed documentation see `help(brmsformula)`.

See Also

[brmsformula](#)

Examples

```

## Not run:
data("nhanes", package = "mice")
N <- nrow(nhanes)

# simple model with missing data
bform1 <- bf(bmi | mi() ~ age * mi(chl)) +
  bf(chl | mi() ~ age) +
  set_rescor(FALSE)

```

```

fit1 <- brm(bform1, data = nhanes)

summary(fit1)
plot(conditional_effects(fit1, resp = "bmi"), ask = FALSE)
loo(fit1, newdata = na.omit(fit1$data))

# simulate some measurement noise
nhanes$se <- rexp(N, 2)

# measurement noise can be handled within 'mi' terms
# with or without the presence of missing values
bform2 <- bf(bmi | mi() ~ age * mi(chl)) +
  bf(chl | mi(se) ~ age) +
  set_rescor(FALSE)

fit2 <- brm(bform2, data = nhanes)

summary(fit2)
plot(conditional_effects(fit2, resp = "bmi"), ask = FALSE)

# 'mi' terms can also be used when some responses are subsetted
nhanes$sub <- TRUE
nhanes$sub[1:2] <- FALSE
nhanes$id <- 1:N
nhanes$idx <- sample(3:N, N, TRUE)

# this requires the addition term 'index' being specified
# in the subsetted part of the model
bform3 <- bf(bmi | mi() ~ age * mi(chl, idx)) +
  bf(chl | mi(se) + subset(sub) + index(id) ~ age) +
  set_rescor(FALSE)

fit3 <- brm(bform3, data = nhanes)

summary(fit3)
plot(conditional_effects(fit3, resp = "bmi"), ask = FALSE)

## End(Not run)

```

Description

Set up a finite mixture family for use in **brms**.

Usage

```
mixture(..., flist = NULL, nmix = 1, order = NULL)
```

Arguments

...	One or more objects providing a description of the response distributions to be combined in the mixture model. These can be family functions, calls to family functions or character strings naming the families. For details of supported families see brmsfamily .
flist	Optional list of objects, which are treated in the same way as objects passed via the ... argument.
nmix	Optional numeric vector specifying the number of times each family is repeated. If specified, it must have the same length as the number of families passed via ... and flist.
order	Ordering constraint to identify mixture components. If 'mu' or TRUE, population-level intercepts of the mean parameters are ordered in non-ordinal models and fixed to the same value in ordinal models (see details). If 'none' or FALSE, no ordering constraint is applied. If NULL (the default), order is set to 'mu' if all families are the same and 'none' otherwise. Other ordering constraints may be implemented in the future.

Details

Most families supported by **brms** can be used to form mixtures. The response variable has to be valid for all components of the mixture family. Currently, the number of mixture components has to be specified by the user. It is not yet possible to estimate the number of mixture components from the data.

Ordering intercepts in mixtures of ordinal families is not possible as each family has itself a set of vector of intercepts (i.e. ordinal thresholds). Instead, **brms** will fix the vector of intercepts across components in ordinal mixtures, if desired, so that users can try to identify the mixture model via selective inclusion of predictors.

For most mixture models, you may want to specify priors on the population-level intercepts via [set_prior](#) to improve convergence. In addition, it is sometimes necessary to set `init = 0` in the call to [brm](#) to allow chains to initialize properly.

For more details on the specification of mixture models, see [brmsformula](#).

Value

An object of class `mixfamily`.

Examples

```
## Not run:
## simulate some data
set.seed(1234)
dat <- data.frame(
  y = c(rnorm(200), rnorm(100, 6)),
  x = rnorm(300),
  z = sample(0:1, 300, TRUE)
)
## fit a simple normal mixture model
```

```

mix <- mixture(gaussian, gaussian)
prior <- c(
  prior(normal(0, 7), Intercept, dpar = mu1),
  prior(normal(5, 7), Intercept, dpar = mu2)
)
fit1 <- brm(bf(y ~ x + z), dat, family = mix,
            prior = prior, chains = 2)
summary(fit1)
pp_check(fit1)

## use different predictors for the components
fit2 <- brm(bf(y ~ 1, mu1 ~ x, mu2 ~ z), dat, family = mix,
            prior = prior, chains = 2)
summary(fit2)

## fix the mixing proportions
fit3 <- brm(bf(y ~ x + z, theta1 = 1, theta2 = 2),
            dat, family = mix, prior = prior,
            init = 0, chains = 2)
summary(fit3)
pp_check(fit3)

## predict the mixing proportions
fit4 <- brm(bf(y ~ x + z, theta2 ~ x),
            dat, family = mix, prior = prior,
            init = 0, chains = 2)
summary(fit4)
pp_check(fit4)

## compare model fit
loo(fit1, fit2, fit3, fit4)

## End(Not run)

```

mm

*Set up multi-membership grouping terms in **brms***

Description

Function to set up a multi-membership grouping term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with grouping terms.

Usage

```

mm(
  ...,
  weights = NULL,
  scale = TRUE,
  by = NULL,

```

```

  cor = TRUE,
  id = NA,
  pw = NULL,
  cov = NULL,
  dist = "gaussian"
)

```

Arguments

...	One or more terms containing grouping factors.
weights	A matrix specifying the membership weights of each member. It should have as many columns as grouping terms specified in If NULL (the default), equally weights are used.
scale	Logical; if TRUE (the default), membership weights are standardized in order to sum to one per row. If negative weights are specified, scale needs to be set to FALSE.
by	An optional factor matrix, specifying sub-populations of the groups. It should have as many columns as grouping terms specified in For each level of the by variable, a separate variance-covariance matrix will be fitted. Levels of the grouping factor must be nested in levels of the by variable matrix.
cor	Logical. If TRUE (the default), group-level terms will be modelled as correlated.
id	Optional character string. All group-level terms across the model with the same id will be modeled as correlated (if cor is TRUE). See brmsformula for more details.
pw	Optional numeric matrix specifying prior weights. They weight the contribution of each group to the log-prior of the group-level coefficients. Should have as many columns as grouping terms specified in ... and one distinct value for each group level.
cov	An optional matrix which is proportional to the within-group covariance matrix of the group-level effects. All levels of the grouping factor should appear as row-names of the corresponding matrix. This argument can be used, among others, to model pedigrees and phylogenetic effects. See vignette("brms_phylogenetics") for more details. By default, levels of the same grouping factor are modeled as independent of each other.
dist	Name of the distribution of the group-level effects. Currently "gaussian" is the only option.

See Also

[brmsformula](#), [mmc](#)

Examples

```

## Not run:
# simulate some data
dat <- data.frame(
  y = rnorm(100), x1 = rnorm(100), x2 = rnorm(100),

```

```

g1 = sample(1:10, 100, TRUE), g2 = sample(1:10, 100, TRUE)
)

# multi-membership model with two members per group and equal weights
fit1 <- brm(y ~ x1 + (1|mm(g1, g2)), data = dat)
summary(fit1)

# weight the first member two times for than the second member
dat$w1 <- rep(2, 100)
dat$w2 <- rep(1, 100)
fit2 <- brm(y ~ x1 + (1|mm(g1, g2, weights = cbind(w1, w2))), data = dat)
summary(fit2)

# multi-membership model with level specific covariate values
dat$xc <- (dat$x1 + dat$x2) / 2
fit3 <- brm(y ~ xc + (1 + mmc(x1, x2) | mm(g1, g2)), data = dat)
summary(fit3)

## End(Not run)

```

mmc*Multi-Membership Covariates***Description**

Specify covariates that vary over different levels of multi-membership grouping factors thus requiring special treatment. This function is almost solely useful, when called in combination with [mm](#). Outside of multi-membership terms it will behave very much like [cbind](#).

Usage

```
mmc(...)
```

Arguments

- ... One or more terms containing covariates corresponding to the grouping levels specified in [mm](#).

Value

A matrix with covariates as columns.

See Also

[mm](#)

Examples

```
## Not run:
# simulate some data
dat <- data.frame(
  y = rnorm(100), x1 = rnorm(100), x2 = rnorm(100),
  g1 = sample(1:10, 100, TRUE), g2 = sample(1:10, 100, TRUE)
)

# multi-membership model with level specific covariate values
dat$xc <- (dat$x1 + dat$x2) / 2
fit <- brm(y ~ xc + (1 + mmc(x1, x2) | mm(g1, g2)), data = dat)
summary(fit)

## End(Not run)
```

Description

Specify a monotonic predictor term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model.

Usage

```
mo(x, id = NA)
```

Arguments

- | | |
|----|--|
| x | An integer variable or an ordered factor to be modeled as monotonic. |
| id | Optional character string. All monotonic terms with the same id within one formula will be modeled as having the same simplex (shape) parameter vector. If all monotonic terms of the same predictor have the same id, the resulting predictions will be conditionally monotonic for all values of interacting covariates (Bürkner & Charpentier, 2020). |

Details

See Bürkner and Charpentier (2020) for the underlying theory. For detailed documentation of the formula syntax used for monotonic terms, see `help(brmsformula)` as well as `vignette("brms_monotonic")`.

References

- Bürkner P. C. & Charpentier E. (2020). Modeling Monotonic Effects of Ordinal Predictors in Regression Models. *British Journal of Mathematical and Statistical Psychology*. doi:10.1111/bmsp.12195

See Also

[brmsformula](#)

Examples

```
## Not run:
# generate some data
income_options <- c("below_20", "20_to_40", "40_to_100", "greater_100")
income <- factor(sample(income_options, 100, TRUE),
                  levels = income_options, ordered = TRUE)
mean_ls <- c(30, 60, 70, 75)
ls <- mean_ls[income] + rnorm(100, sd = 7)
dat <- data.frame(income, ls)

# fit a simple monotonic model
fit1 <- brm(ls ~ mo(income), data = dat)
summary(fit1)
plot(fit1, N = 6)
plot(conditional_effects(fit1), points = TRUE)

# model interaction with other variables
dat$x <- sample(c("a", "b", "c"), 100, TRUE)
fit2 <- brm(ls ~ mo(income)*x, data = dat)
summary(fit2)
plot(conditional_effects(fit2), points = TRUE)

# ensure conditional monotonicity
fit3 <- brm(ls ~ mo(income, id = "i")*x, data = dat)
summary(fit3)
plot(conditional_effects(fit3), points = TRUE)

## End(Not run)
```

`model_weights.brmsfit` *Model Weighting Methods*

Description

Compute model weights in various ways, for instance, via stacking of posterior predictive distributions, Akaike weights, or marginal likelihoods.

Usage

```
## S3 method for class 'brmsfit'
model_weights(x, ..., weights = "stacking", model_names = NULL)

model_weights(x, ...)
```

Arguments

x	A <code>brmsfit</code> object.
...	More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments.
weights	Name of the criterion to compute weights from. Should be one of "loo", "waic", "kfold", "stacking" (current default), "bma", or "pseudobma". For the former three options, Akaike weights will be computed based on the information criterion values returned by the respective methods. For "stacking" and "pseudobma", method loo_model_weights will be used to obtain weights. For "bma", method post_prob will be used to compute Bayesian model averaging weights based on log marginal likelihood values (make sure to specify reasonable priors in this case). For some methods, weights may also be a numeric vector of pre-specified weights.
model_names	If NULL (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Value

A numeric vector of weights for the models.

Examples

```
## Not run:
# model with 'treat' as predictor
fit1 <- brm(rating ~ treat + period + carry, data = inhaler)
summary(fit1)

# model without 'treat' as predictor
fit2 <- brm(rating ~ period + carry, data = inhaler)
summary(fit2)

# obtain Akaike weights based on the WAIC
model_weights(fit1, fit2, weights = "waic")

## End(Not run)
```

Description

Density function and random generation for the multivariate normal distribution with mean vector μ and covariance matrix Σ .

Usage

```
dmulti_normal(x, mu, Sigma, log = FALSE, check = FALSE)

rmulti_normal(n, mu, Sigma, check = FALSE)
```

Arguments

x	Vector or matrix of quantiles. If x is a matrix, each row is taken to be a quantile.
mu	Mean vector with length equal to the number of dimensions.
Sigma	Covariance matrix.
log	Logical; If TRUE, values are returned on the log scale.
check	Logical; Indicates whether several input checks should be performed. Defaults to FALSE to improve efficiency.
n	Number of draws to sample from the distribution.

Details

See the Stan user's manual <https://mc-stan.org/docs/> for details on the parameterization

Description

Density function and random generation for the multivariate Student-t distribution with location vector mu, covariance matrix Sigma, and degrees of freedom df.

Usage

```
dmulti_student_t(x, df, mu, Sigma, log = FALSE, check = FALSE)

rmulti_student_t(n, df, mu, Sigma, check = FALSE)
```

Arguments

x	Vector or matrix of quantiles. If x is a matrix, each row is taken to be a quantile.
df	Vector of degrees of freedom.
mu	Location vector with length equal to the number of dimensions.
Sigma	Covariance matrix.
log	Logical; If TRUE, values are returned on the log scale.
check	Logical; Indicates whether several input checks should be performed. Defaults to FALSE to improve efficiency.
n	Number of draws to sample from the distribution.

Details

See the Stan user's manual <https://mc-stan.org/docs/> for details on the parameterization

mvbind*Bind response variables in multivariate models***Description**

Can be used to specify a multivariate **brms** model within a single formula. Outside of **brmsformula**, it just behaves like **cbind**.

Usage

```
mvbind(...)
```

Arguments

...	Same as in cbind
-----	-------------------------

See Also

[brmsformula](#), [mvbrmsformula](#)

Examples

```
bf(mvbind(y1, y2) ~ x)
```

mvbrmsformula*Set up a multivariate model formula for use in **brms*****Description**

Set up a multivariate model formula for use in the **brms** package allowing to define (potentially non-linear) additive multilevel models for all parameters of the assumed response distributions.

Usage

```
mvbrmsformula(..., flist = NULL, rescor = NULL)
```

Arguments

...	Objects of class formula or brmsformula , each specifying a univariate model. See brmsformula for details on how to specify univariate models.
flist	Optional list of formulas, which are treated in the same way as formulas passed via the ... argument.
rescor	Logical; Indicates if residual correlation between the response variables should be modeled. Currently, this is only possible in multivariate gaussian and student models. If NULL (the default), rescor is internally set to TRUE when possible.

Details

See vignette("brms_multivariate") for a case study.

Value

An object of class `mvbrmsformula`, which is essentially a list containing all model formulas as well as some additional information for multivariate models.

See Also

[brmsformula](#), [brmsformula-helpers](#)

Examples

```
bf1 <- bf(y1 ~ x + (1|g))
bf2 <- bf(y2 ~ s(z))
mvbf(bf1, bf2)
```

<code>ngrps.brmsfit</code>	<i>Number of Grouping Factor Levels</i>
----------------------------	---

Description

Extract the number of levels of one or more grouping factors.

Usage

```
## S3 method for class 'brmsfit'
ngrps(object, ...)

ngrps(object, ...)
```

Arguments

- | | |
|---------------------|--------------------|
| <code>object</code> | An R object. |
| <code>...</code> | Currently ignored. |

Value

A named list containing the number of levels per grouping factor.

<code>nsamples.brmsfit</code>	<i>(Deprecated) Number of Posterior Samples</i>
-------------------------------	---

Description

Extract the number of posterior samples (draws) stored in a fitted Bayesian model. Method `nsamples` is deprecated. Please use `ndraws` instead.

Usage

```
## S3 method for class 'brmsfit'
nsamples(object, subset = NULL, incl_warmup = FALSE, ...)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>subset</code>	An optional integer vector defining a subset of samples to be considered.
<code>incl_warmup</code>	A flag indicating whether to also count warmup / burn-in samples.
<code>...</code>	Currently ignored.

<code>opencl</code>	<i>GPU support in Stan via OpenCL</i>
---------------------	---------------------------------------

Description

Use OpenCL for GPU support in **Stan** via the **brms** interface. Only some **Stan** functions can be run on a GPU at this point and so a lot of **brms** models won't benefit from OpenCL for now.

Usage

```
opencl(ids = NULL)
```

Arguments

<code>ids</code>	(integer vector of length 2) The platform and device IDs of the OpenCL device to use for fitting. If you don't know the IDs of your OpenCL device, <code>c(0, 0)</code> is most likely what you need.
------------------	---

Details

For more details on OpenCL in **Stan**, check out https://mc-stan.org/docs/2_26/cmdstan-guide/parallelization.html#opencl as well as https://mc-stan.org/docs/2_26/stan-users-guide/opencl.html.

Value

A `brmsopenc1` object which can be passed to the `opencl` argument of `brm` and related functions.

Examples

```
## Not run:
# this model just serves as an illustration
# OpenCL may not actually speed things up here
fit <- brm(count ~ zAge + zBase * Trt + (1|patient),
            data = epilepsy, family = poisson(),
            chains = 2, cores = 2, opencl = opencl(c(0, 0)),
            backend = "cmdstanr")
summary(fit)

## End(Not run)
```

pairs.brmsfit*Create a matrix of output plots from a brmsfit object***Description**

A `pairs` method that is customized for MCMC output.

Usage

```
## S3 method for class 'brmsfit'
pairs(x, pars = NA, variable = NULL, regex = FALSE, fixed = FALSE, ...)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code>
<code>pars</code>	Deprecated alias of <code>variable</code> . Names of the parameters to plot, as given by a character vector or a regular expression.
<code>variable</code>	Names of the variables (parameters) to plot, as given by a character vector or a regular expression (if <code>regex</code> = TRUE). By default, a hopefully not too large selection of variables is plotted.
<code>regex</code>	Logical; Indicates whether <code>variable</code> should be treated as regular expressions. Defaults to FALSE.
<code>fixed</code>	(Deprecated) Indicates whether parameter names should be matched exactly (TRUE) or treated as regular expressions (FALSE). Default is FALSE and only works with argument <code>pars</code> .
<code>...</code>	Further arguments to be passed to <code>mcmc_pairs</code> .

Details

For a detailed description see [mcmc_pairs](#).

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt
           + (1|patient) + (1|visit),
           data = epilepsy, family = "poisson")
pairs(fit, variable = variables(fit)[1:3])
pairs(fit, variable = "^sd_", regex = TRUE)

## End(Not run)
```

parnames

Extract Parameter Names

Description

Extract all parameter names of a given model.

Usage

```
parnames(x, ...)
```

Arguments

- | | |
|-----|--|
| x | An R object |
| ... | Further arguments passed to or from other methods. |

Value

A character vector containing the parameter names of the model.

plot.brmsfit

Trace and Density Plots for MCMC Draws

Description

Trace and Density Plots for MCMC Draws

Usage

```
## S3 method for class 'brmsfit'
plot(
  x,
  pars = NA,
  combo = c("hist", "trace"),
  nvariables = 5,
  N = NULL,
  variable = NULL,
  regex = FALSE,
  fixed = FALSE,
  bins = 30,
  theme = NULL,
  plot = TRUE,
  ask = TRUE,
  newpage = TRUE,
  ...
)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code> .
<code>pars</code>	Deprecated alias of <code>variable</code> . Names of the parameters to plot, as given by a character vector or a regular expression.
<code>combo</code>	A character vector with at least two elements. Each element of <code>combo</code> corresponds to a column in the resulting graphic and should be the name of one of the available MCMC functions (omitting the <code>mcmc_</code> prefix).
<code>nvariables</code>	The number of variables (parameters) plotted per page.
<code>N</code>	Deprecated alias of <code>nvariables</code> .
<code>variable</code>	Names of the variables (parameters) to plot, as given by a character vector or a regular expression (if <code>regex = TRUE</code>). By default, a hopefully not too large selection of variables is plotted.
<code>regex</code>	Logical; Indicates whether <code>variable</code> should be treated as regular expressions. Defaults to FALSE.
<code>fixed</code>	(Deprecated) Indicates whether parameter names should be matched exactly (TRUE) or treated as regular expressions (FALSE). Default is FALSE and only works with argument <code>pars</code> .
<code>bins</code>	Number of bins used for posterior histograms (defaults to 30).
<code>theme</code>	A theme object modifying the appearance of the plots. For some basic themes see ggtheme and theme_default .
<code>plot</code>	Logical; indicates if plots should be plotted directly in the active graphic device. Defaults to TRUE.
<code>ask</code>	Logical; indicates if the user is prompted before a new page is plotted. Only used if <code>plot</code> is TRUE.

`newpage` Logical; indicates if the first set of plots should be plotted to a new page. Only used if `plot` is TRUE.
`...` Further arguments passed to `mcmc_combo`.

Value

An invisible list of `gttable` objects.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt
            + (1|patient) + (1|visit),
            data = epilepsy, family = "poisson")
plot(fit)
## plot population-level effects only
plot(fit, variable = "^b_",
      regex = TRUE)

## End(Not run)
```

posterior_average.brmsfit

Posterior draws of parameters averaged across models

Description

Extract posterior draws of parameters averaged across models. Weighting can be done in various ways, for instance using Akaike weights based on information criteria or marginal likelihoods.

Usage

```
## S3 method for class 'brmsfit'
posterior_average(
  x,
  ...,
  variable = NULL,
  pars = NULL,
  weights = "stacking",
  ndraws = NULL,
  nsamples = NULL,
  missing = NULL,
  model_names = NULL,
  control = list(),
  seed = NULL
)
posterior_average(x, ...)
```

Arguments

x	A <code>brmsfit</code> object.
...	More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments.
variable	Names of variables (parameters) for which to average across models. Only those variables can be averaged that appear in every model. Defaults to all overlapping variables.
pars	Deprecated alias of <code>variable</code> .
weights	Name of the criterion to compute weights from. Should be one of "loo", "waic", "kfold", "stacking" (current default), "bma", or "pseudobma". For the former three options, Akaike weights will be computed based on the information criterion values returned by the respective methods. For "stacking" and "pseudobma", method loo_model_weights will be used to obtain weights. For "bma", method post_prob will be used to compute Bayesian model averaging weights based on log marginal likelihood values (make sure to specify reasonable priors in this case). For some methods, weights may also be a numeric vector of pre-specified weights.
ndraws	Total number of posterior draws to use.
nsamples	Deprecated alias of <code>ndraws</code> .
missing	An optional numeric value or a named list of numeric values to use if a model does not contain a variable for which posterior draws should be averaged. Defaults to <code>NULL</code> , in which case only those variables can be averaged that are present in all of the models.
model_names	If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.
control	Optional list of further arguments passed to the function specified in <code>weights</code> .
seed	A single numeric value passed to set.seed to make results reproducible.

Details

Weights are computed with the [model_weights](#) method.

Value

A `data.frame` of posterior draws.

See Also

[model_weights](#), [pp_average](#)

Examples

```
## Not run:
# model with 'treat' as predictor
fit1 <- brm(rating ~ treat + period + carry, data = inhaler)
```

```

summary(fit1)

# model without 'treat' as predictor
fit2 <- brm(rating ~ period + carry, data = inhaler)
summary(fit2)

# compute model-averaged posteriors of overlapping parameters
posterior_average(fit1, fit2, weights = "waic")

## End(Not run)

```

posterior_epred.brmsfit

Draws from the Expected Value of the Posterior Predictive Distribution

Description

Compute posterior draws of the expected value of the posterior predictive distribution. Can be performed for the data used to fit the model (posterior predictive checks) or for new data. By definition, these predictions have smaller variance than the posterior predictions performed by the [posterior_predict.brmsfit](#) method. This is because only the uncertainty in the expected value of the posterior predictive distribution is incorporated in the draws computed by `posterior_epred` while the residual error is ignored there. However, the estimated means of both methods averaged across draws should be very similar.

Usage

```

## S3 method for class 'brmsfit'
posterior_epred(
  object,
  newdata = NULL,
  re_formula = NULL,
  re.form = NULL,
  resp = NULL,
  dpar = NULL,
  nlpar = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  ...
)

```

Arguments

object An object of class `brmsfit`.

newdata	An optional data.frame for which to evaluate predictions. If NULL (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
re_formula	formula containing group-level effects to be considered in the prediction. If NULL (default), include all group-level effects; if NA or ~0, include no group-level effects.
re.form	Alias of re_formula.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
dpar	Optional name of a predicted distributional parameter. If specified, expected predictions of this parameters are returned.
npar	Optional name of a predicted non-linear parameter. If specified, expected predictions of this parameters are returned.
ndraws	Positive integer indicating how many posterior draws should be used. If NULL (the default) all draws are used. Ignored if draw_ids is not NULL.
draw_ids	An integer vector specifying the posterior draws to be used. If NULL (the default), all draws are used.
sort	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (FALSE; default) or in the order of the time series (TRUE).
...	Further arguments passed to prepare_predictions that control several aspects of data validation and prediction.

Details

NA values within factors in newdata, are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding.

In multilevel models, it is possible to allow new levels of grouping factors to be used in the predictions. This can be controlled via argument `allow_new_levels`. New levels can be sampled in multiple ways, which can be controlled via argument `sample_new_levels`. Both of these arguments are documented in [prepare_predictions](#) along with several other useful arguments to control specific aspects of the predictions.

Value

An array of draws. For categorical and ordinal models, the output is an S x N x C array. Otherwise, the output is an S x N matrix, where S is the number of posterior draws, N is the number of observations, and C is the number of categories. In multivariate models, an additional dimension is added to the output which indexes along the different response variables.

Examples

```
## Not run:
## fit a model
```

```

fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler)

## compute expected predictions
ppe <- posterior_epred(fit)
str(ppe)

## End(Not run)

```

posterior_interval.brmsfit*Compute posterior uncertainty intervals***Description**

Compute posterior uncertainty intervals for `brmsfit` objects.

Usage

```
## S3 method for class 'brmsfit'
posterior_interval(object, pars = NA, variable = NULL, prob = 0.95, ...)
```

Arguments

- | | |
|-----------------------|---|
| <code>object</code> | An object of class <code>brmsfit</code> . |
| <code>pars</code> | Deprecated alias of <code>variable</code> . For reasons of backwards compatibility, <code>pars</code> is interpreted as a vector of regular expressions by default unless <code>fixed = TRUE</code> is specified. |
| <code>variable</code> | A character vector providing the variables to extract. By default, all variables are extracted. |
| <code>prob</code> | A value between 0 and 1 indicating the desired probability to be covered by the uncertainty intervals. The default is 0.95. |
| <code>...</code> | More arguments passed to as.matrix.brmsfit . |

Value

A `matrix` with lower and upper interval bounds as columns and as many rows as selected variables.

Examples

```

## Not run:
fit <- brm(count ~ zAge + zBase * Trt,
            data = epilepsy, family = negbinomial())
posterior_interval(fit)

## End(Not run)

```

posterior_linpred.brmsfit

Posterior Draws of the Linear Predictor

Description

Compute posterior draws of the linear predictor, that is draws before applying any link functions or other transformations. Can be performed for the data used to fit the model (posterior predictive checks) or for new data.

Usage

```
## S3 method for class 'brmsfit'
posterior_linpred(
  object,
  transform = FALSE,
  newdata = NULL,
  re_formula = NULL,
  re.form = NULL,
  resp = NULL,
  dpar = NULL,
  nlpar = NULL,
  incl_thres = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>transform</code>	Logical; if <code>FALSE</code> (the default), draws of the linear predictor are returned. If <code>TRUE</code> , draws of the transformed linear predictor, that is, after applying the inverse link function are returned.
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>re.form</code>	Alias of <code>re_formula</code> .
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.

dpar	Name of a predicted distributional parameter for which draws are to be returned. By default, draws of the main distributional parameter(s) "mu" are returned.
nlpar	Optional name of a predicted non-linear parameter. If specified, expected predictions of this parameters are returned.
incl_thres	Logical; only relevant for ordinal models when <code>transform</code> is FALSE, and ignored otherwise. Shall the thresholds and category-specific effects be included in the linear predictor? For backwards compatibility, the default is to not include them.
ndraws	Positive integer indicating how many posterior draws should be used. If NULL (the default) all draws are used. Ignored if <code>draw_ids</code> is not NULL.
draw_ids	An integer vector specifying the posterior draws to be used. If NULL (the default), all draws are used.
sort	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (FALSE; default) or in the order of the time series (TRUE).
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

See Also

[posterior_epred.brmsfit](#)

Examples

```
## Not run:
## fit a model
fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler)

## extract linear predictor values
pl <- posterior_linpred(fit)
str(pl)

## End(Not run)
```

posterior_predict.brmsfit

Draws from the Posterior Predictive Distribution

Description

Compute posterior draws of the posterior predictive distribution. Can be performed for the data used to fit the model (posterior predictive checks) or for new data. By definition, these draws have higher variance than draws of the expected value of the posterior predictive distribution computed by [posterior_epred.brmsfit](#). This is because the residual error is incorporated in `posterior_predict`. However, the estimated means of both methods averaged across draws should be very similar.

Usage

```
## S3 method for class 'brmsfit'
posterior_predict(
  object,
  newdata = NULL,
  re_formula = NULL,
  re.form = NULL,
  transform = NULL,
  resp = NULL,
  negative_rt = FALSE,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  ntrys = 5,
  cores = NULL,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>re.form</code>	Alias of <code>re_formula</code> .
<code>transform</code>	(Deprecated) A function or a character string naming a function to be applied on the predicted responses before summary statistics are computed.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>negative_rt</code>	Only relevant for Wiener diffusion models. A flag indicating whether response times of responses on the lower boundary should be returned as negative values. This allows to distinguish responses on the upper and lower boundary. Defaults to <code>FALSE</code> .
<code>ndraws</code>	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
<code>draw_ids</code>	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
<code>sort</code>	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (<code>FALSE</code> ; default) or in the order of the time series (<code>TRUE</code>).

ntrys	Parameter used in rejection sampling for truncated discrete models only (defaults to 5). See Details for more information.
cores	Number of cores (defaults to 1). On non-Windows systems, this argument can be set globally via the <code>mc.cores</code> option.
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Details

NA values within factors in `newdata`, are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding.

In multilevel models, it is possible to allow new levels of grouping factors to be used in the predictions. This can be controlled via argument `allow_new_levels`. New levels can be sampled in multiple ways, which can be controlled via argument `sample_new_levels`. Both of these arguments are documented in `prepare_predictions` along with several other useful arguments to control specific aspects of the predictions.

For truncated discrete models only: In the absence of any general algorithm to sample from truncated discrete distributions, rejection sampling is applied in this special case. This means that values are sampled until a value lies within the defined truncation boundaries. In practice, this procedure may be rather slow (especially in R). Thus, we try to do approximate rejection sampling by sampling each value `ntrys` times and then select a valid value. If all values are invalid, the closest boundary is used, instead. If there are more than a few of these pathological cases, a warning will occur suggesting to increase argument `ntrys`.

Value

An array of draws. In univariate models, the output is as an S x N matrix, where S is the number of posterior draws and N is the number of observations. In multivariate models, an additional dimension is added to the output which indexes along the different response variables.

Examples

```
## Not run:
## fit a model
fit <- brm(time | cens(censored) ~ age + sex + (1 + age || patient),
            data = kidney, family = "exponential", init = "0")

## predicted responses
pp <- posterior_predict(fit)
str(pp)

## predicted responses excluding the group-level effect of age
pp <- posterior_predict(fit, re_formula = ~ (1 | patient))
str(pp)

## predicted responses of patient 1 for new data
newdata <- data.frame(
  sex = factor(c("male", "female")),
  age = c(20, 50),
```

```

    patient = c(1, 1)
  )
pp <- posterior_predict(fit, newdata = newdata)
str(pp)

## End(Not run)

```

posterior_samples.brmsfit*(Deprecated) Extract Posterior Samples***Description**

Extract posterior samples of specified parameters. The `posterior_samples` method is deprecated. We recommend using the more modern and consistent `as_draws_*` extractor functions of the **posterior** package instead.

Usage

```

## S3 method for class 'brmsfit'
posterior_samples(
  x,
  pars = NA,
  fixed = FALSE,
  add_chain = FALSE,
  subset = NULL,
  as.matrix = FALSE,
  as.array = FALSE,
  ...
)
posterior_samples(x, pars = NA, ...)

```

Arguments

<code>x</code>	An R object typically of class <code>brmsfit</code>
<code>pars</code>	Names of parameters for which posterior samples should be returned, as given by a character vector or regular expressions. By default, all posterior samples of all parameters are extracted.
<code>fixed</code>	Indicates whether parameter names should be matched exactly (TRUE) or treated as regular expressions (FALSE). Default is FALSE.
<code>add_chain</code>	A flag indicating if the returned <code>data.frame</code> should contain two additional columns. The <code>chain</code> column indicates the chain in which each sample was generated, the <code>iter</code> column indicates the iteration number within each chain.
<code>subset</code>	A numeric vector indicating the rows (i.e., posterior samples) to be returned. If <code>NULL</code> (the default), all posterior samples are returned.

as.matrix	Should the output be a <code>matrix</code> instead of a <code>data.frame</code> ? Defaults to FALSE.
as.array	Should the output be an <code>array</code> instead of a <code>data.frame</code> ? Defaults to FALSE.
...	Arguments passed to individual methods (if applicable).

Value

A `data.frame` (matrix or array) containing the posterior samples.

See Also

[as_draws](#), [as.data.frame](#)

Examples

```
## Not run:
fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler, family = "cumulative")

# extract posterior samples of population-level effects
samples1 <- posterior_samples(fit, pars = "^b")
head(samples1)

# extract posterior samples of group-level standard deviations
samples2 <- posterior_samples(fit, pars = "^sd_")
head(samples2)

## End(Not run)
```

posterior_smooths.brmsfit

Posterior Predictions of Smooth Terms

Description

Compute posterior predictions of smooth s and t2 terms of models fitted with **brms**.

Usage

```
## S3 method for class 'brmsfit'
posterior_smooths(
  object,
  smooth,
  newdata = NULL,
  resp = NULL,
  dpar = NULL,
  npar = NULL,
  ndraws = NULL,
```

```

  draw_ids = NULL,
  ...
)
posterior_smooths(object, ...)

```

Arguments

object	An object of class <code>brmsfit</code> .
smooth	Name of a single smooth term for which predictions should be computed.
newdata	An optional <code>data.frame</code> for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. Only those variables appearing in the chosen smooth term are required.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
dpar	Optional name of a predicted distributional parameter. If specified, expected predictions of this parameters are returned.
npar	Optional name of a predicted non-linear parameter. If specified, expected predictions of this parameters are returned.
ndraws	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
draw_ids	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
...	Currently ignored.

Value

An S x N matrix, where S is the number of posterior draws and N is the number of observations.

Examples

```

## Not run:
set.seed(0)
dat <- mgcv:::gamSim(1, n = 200, scale = 2)
fit <- brm(y ~ s(x0) + s(x1) + s(x2) + s(x3), data = dat)
summary(fit)

newdata <- data.frame(x2 = seq(0, 1, 10))
str(posterior_smooths(fit, smooth = "s(x2)", newdata = newdata))

## End(Not run)

```

`posterior_summary` *Summarize Posterior draws*

Description

Summarizes posterior draws based on point estimates (mean or median), estimation errors (SD or MAD) and quantiles. This function mainly exists to retain backwards compatibility. It will eventually be replaced by functions of the **posterior** package (see examples below).

Usage

```
posterior_summary(x, ...)

## Default S3 method:
posterior_summary(x, probs = c(0.025, 0.975), robust = FALSE, ...)

## S3 method for class 'brmsfit'
posterior_summary(
  x,
  pars = NA,
  variable = NULL,
  probs = c(0.025, 0.975),
  robust = FALSE,
  ...
)
```

Arguments

<code>x</code>	An R object.
<code>...</code>	More arguments passed to or from other methods.
<code>probs</code>	The percentiles to be computed by the <code>quantile</code> function.
<code>robust</code>	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead.
<code>pars</code>	Deprecated alias of <code>variable</code> . For reasons of backwards compatibility, <code>pars</code> is interpreted as a vector of regular expressions by default unless <code>fixed = TRUE</code> is specified.
<code>variable</code>	A character vector providing the variables to extract. By default, all variables are extracted.

Value

A matrix where rows indicate variables and columns indicate the summary estimates.

See Also

[summarize_draws](#)

Examples

```
## Not run:
fit <- brm(time ~ age * sex, data = kidney)
posterior_summary(fit)

# recommended workflow using posterior
library(posterior)
draws <- as_draws_array(fit)
summarise_draws(draws, default_summary_measures())

## End(Not run)
```

posterior_table

Table Creation for Posterior Draws

Description

Create a table for unique values of posterior draws. This is usually only useful when summarizing predictions of ordinal models.

Usage

```
posterior_table(x, levels = NULL)
```

Arguments

- x A matrix of posterior draws where rows indicate draws and columns indicate parameters.
- levels Optional values of possible posterior values. Defaults to all unique values in x.

Value

A matrix where rows indicate parameters and columns indicate the unique values of posterior draws.

Examples

```
## Not run:
fit <- brm(rating ~ period + carry + treat,
           data = inhaler, family = cumulative())
pr <- predict(fit, summary = FALSE)
posterior_table(pr)

## End(Not run)
```

post_prob.brmsfit *Posterior Model Probabilities from Marginal Likelihoods*

Description

Compute posterior model probabilities from marginal likelihoods. The `brmsfit` method is just a thin wrapper around the corresponding method for `bridge` objects.

Usage

```
## S3 method for class 'brmsfit'
post_prob(x, ..., prior_prob = NULL, model_names = NULL)
```

Arguments

<code>x</code>	A <code>brmsfit</code> object.
<code>...</code>	More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments.
<code>prior_prob</code>	Numeric vector with prior model probabilities. If omitted, a uniform prior is used (i.e., all models are equally likely a priori). The default <code>NULL</code> corresponds to equal prior model weights.
<code>model_names</code>	If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Details

Computing the marginal likelihood requires samples of all variables defined in Stan's `parameters` block to be saved. Otherwise `post_prob` cannot be computed. Thus, please set `save_all_pars = TRUE` in the call to `brm`, if you are planning to apply `post_prob` to your models.

The computation of model probabilities based on bridge sampling requires a lot more posterior samples than usual. A good conservative rule of thumb is perhaps 10-fold more samples (read: the default of 4000 samples may not be enough in many cases). If not enough posterior samples are provided, the bridge sampling algorithm tends to be unstable leading to considerably different results each time it is run. We thus recommend running `post_prob` multiple times to check the stability of the results.

More details are provided under [bridgesampling::post_prob](#).

See Also

[bridge_sampler](#), [bayes_factor](#)

Examples

```

## Not run:
# model with the treatment effect
fit1 <- brm(
  count ~ zAge + zBase + Trt,
  data = epilepsy, family = negbinomial(),
  prior = prior(normal(0, 1), class = b),
  save_all_pars = TRUE
)
summary(fit1)

# model without the treatent effect
fit2 <- brm(
  count ~ zAge + zBase,
  data = epilepsy, family = negbinomial(),
  prior = prior(normal(0, 1), class = b),
  save_all_pars = TRUE
)
summary(fit2)

# compute the posterior model probabilities
post_prob(fit1, fit2)

# specify prior model probabilities
post_prob(fit1, fit2, prior_prob = c(0.8, 0.2))

## End(Not run)

```

`pp_average.brmsfit` *Posterior predictive draws averaged across models*

Description

Compute posterior predictive draws averaged across models. Weighting can be done in various ways, for instance using Akaike weights based on information criteria or marginal likelihoods.

Usage

```

## S3 method for class 'brmsfit'
pp_average(
  x,
  ...,
  weights = "stacking",
  method = "posterior_predict",
  ndraws = NULL,
  nsamples = NULL,
  summary = TRUE,
  probs = c(0.025, 0.975),

```

```

  robust = FALSE,
  model_names = NULL,
  control = list(),
  seed = NULL
)
pp_average(x, ...)

```

Arguments

<code>x</code>	A <code>brmsfit</code> object.
<code>...</code>	More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments.
<code>weights</code>	Name of the criterion to compute weights from. Should be one of "loo", "waic", "kfold", "stacking" (current default), "bma", or "pseudobma". For the former three options, Akaike weights will be computed based on the information criterion values returned by the respective methods. For "stacking" and "pseudobma", method loo_model_weights will be used to obtain weights. For "bma", method post_prob will be used to compute Bayesian model averaging weights based on log marginal likelihood values (make sure to specify reasonable priors in this case). For some methods, <code>weights</code> may also be a numeric vector of pre-specified weights.
<code>method</code>	Method used to obtain predictions to average over. Should be one of "posterior_predict" (default), "posterior_epred", "posterior_linpred" or "predictive_error".
<code>ndraws</code>	Total number of posterior draws to use.
<code>nsamples</code>	Deprecated alias of <code>ndraws</code> .
<code>summary</code>	Should summary statistics (i.e. means, sds, and 95% intervals) be returned instead of the raw values? Default is TRUE.
<code>probs</code>	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
<code>robust</code>	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
<code>model_names</code>	If NULL (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.
<code>control</code>	Optional list of further arguments passed to the function specified in <code>weights</code> .
<code>seed</code>	A single numeric value passed to set.seed to make results reproducible.

Details

Weights are computed with the [model_weights](#) method.

Value

Same as the output of the method specified in argument `method`.

See Also

[model_weights](#), [posterior_average](#)

Examples

```
## Not run:
# model with 'treat' as predictor
fit1 <- brm(rating ~ treat + period + carry, data = inhaler)
summary(fit1)

# model without 'treat' as predictor
fit2 <- brm(rating ~ period + carry, data = inhaler)
summary(fit2)

# compute model-averaged predicted values
(df <- unique(inhaler[, c("treat", "period", "carry")]))
pp_average(fit1, fit2, newdata = df)

# compute model-averaged fitted values
pp_average(fit1, fit2, method = "fitted", newdata = df)

## End(Not run)
```

Description

Perform posterior predictive checks with the help of the **bayesplot** package.

Usage

```
## S3 method for class 'brmsfit'
pp_check(
  object,
  type,
  ndraws = NULL,
  prefix = c("ppc", "ppd"),
  group = NULL,
  x = NULL,
  newdata = NULL,
  resp = NULL,
  draw_ids = NULL,
  nsamples = NULL,
  subset = NULL,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>type</code>	Type of the ppc plot as given by a character string. See PPC for an overview of currently supported types. You may also use an invalid type (e.g. <code>type = "xyz"</code>) to get a list of supported types in the resulting error message.
<code>ndraws</code>	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> all draws are used. If not specified, the number of posterior draws is chosen automatically. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
<code>prefix</code>	The prefix of the bayesplot function to be applied. Either <code>"ppc"</code> (posterior predictive check; the default) or <code>"ppd"</code> (posterior predictive distribution), the latter being the same as the former except that the observed data is not shown for <code>"ppd"</code> .
<code>group</code>	Optional name of a factor variable in the model by which to stratify the ppc plot. This argument is required for <code>ppc *_grouped</code> types and ignored otherwise.
<code>x</code>	Optional name of a variable in the model. Only used for ppc types having an <code>x</code> argument and ignored otherwise.
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>draw_ids</code>	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
<code>nsamples</code>	Deprecated alias of <code>ndraws</code> .
<code>subset</code>	Deprecated alias of <code>draw_ids</code> .
<code>...</code>	Further arguments passed to <code>predict.brmsfit</code> as well as to the PPC function specified in <code>type</code> .

Details

For a detailed explanation of each of the ppc functions, see the [PPC](#) documentation of the **bayesplot** package.

Value

A ggplot object that can be further customized using the **ggplot2** package.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt
            + (1|patient) + (1|obs),
            data = epilepsy, family = poisson())
```

```

pp_check(fit) # shows dens_overlay plot by default
pp_check(fit, type = "error_hist", ndraws = 11)
pp_check(fit, type = "scatter_avg", ndraws = 100)
pp_check(fit, type = "stat_2d")
pp_check(fit, type = "rootogram")
pp_check(fit, type = "loo_pit")

## get an overview of all valid types
pp_check(fit, type = "xyz")

## get a plot without the observed data
pp_check(fit, prefix = "ppd")

## End(Not run)

```

pp_mixture.brmsfit *Posterior Probabilities of Mixture Component Memberships*

Description

Compute the posterior probabilities of mixture component memberships for each observation including uncertainty estimates.

Usage

```

## S3 method for class 'brmsfit'
pp_mixture(
  x,
  newdata = NULL,
  re_formula = NULL,
  resp = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  log = FALSE,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)
pp_mixture(x, ...)

```

Arguments

x An R object usually of class brmsfit.

newdata	An optional data.frame for which to evaluate predictions. If NULL (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
re_formula	formula containing group-level effects to be considered in the prediction. If NULL (default), include all group-level effects; if NA or ~0, include no group-level effects.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
ndraws	Positive integer indicating how many posterior draws should be used. If NULL (the default) all draws are used. Ignored if draw_ids is not NULL.
draw_ids	An integer vector specifying the posterior draws to be used. If NULL (the default), all draws are used.
log	Logical; Indicates whether to return probabilities on the log-scale.
summary	Should summary statistics be returned instead of the raw values? Default is TRUE.
robust	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if summary is TRUE.
probs	The percentiles to be computed by the quantile function. Only used if summary is TRUE.
...	Further arguments passed to prepare_predictions that control several aspects of data validation and prediction.

Details

The returned probabilities can be written as $P(K_n = k|Y_n)$, that is the posterior probability that observation n originates from component k. They are computed using Bayes' Theorem

$$P(K_n = k|Y_n) = P(Y_n|K_n = k)P(K_n = k)/P(Y_n),$$

where $P(Y_n|K_n = k)$ is the (posterior) likelihood of observation n for component k, $P(K_n = k)$ is the (posterior) mixing probability of component k (i.e. parameter theta<k>), and

$$P(Y_n) = \sum_{k=1,\dots,K} P(Y_n|K_n = k)P(K_n = k)$$

is a normalizing constant.

Value

If summary = TRUE, an N x E x K array, where N is the number of observations, K is the number of mixture components, and E is equal to length(probs) + 2. If summary = FALSE, an S x N x K array, where S is the number of posterior draws.

Examples

```

## Not run:
## simulate some data
set.seed(1234)
dat <- data.frame(
  y = c(rnorm(100), rnorm(50, 2)),
  x = rnorm(150)
)
## fit a simple normal mixture model
mix <- mixture(gaussian, nmix = 2)
prior <- c(
  prior(normal(0, 5), Intercept, nlpar = mu1),
  prior(normal(0, 5), Intercept, nlpar = mu2),
  prior(dirichlet(2, 2), theta)
)
fit1 <- brm(bf(y ~ x), dat, family = mix,
            prior = prior, chains = 2, init = 0)
summary(fit1)

## compute the membership probabilities
ppm <- pp_mixture(fit1)
str(ppm)

## extract point estimates for each observation
head(ppm[, 1, ])

## classify every observation according to
## the most likely component
apply(ppm[, 1, ], 1, which.max)

## End(Not run)

```

predict.brmsfit *Draws from the Posterior Predictive Distribution*

Description

This method is an alias of [posterior_predict.brmsfit](#) with additional arguments for obtaining summaries of the computed draws.

Usage

```

## S3 method for class 'brmsfit'
predict(
  object,
  newdata = NULL,
  re_formula = NULL,
  transform = NULL,

```

```

  resp = NULL,
  negative_rt = FALSE,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  ntrys = 5,
  cores = NULL,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)

```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>transform</code>	(Deprecated) A function or a character string naming a function to be applied on the predicted responses before summary statistics are computed.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>negative_rt</code>	Only relevant for Wiener diffusion models. A flag indicating whether response times of responses on the lower boundary should be returned as negative values. This allows to distinguish responses on the upper and lower boundary. Defaults to <code>FALSE</code> .
<code>ndraws</code>	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
<code>draw_ids</code>	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
<code>sort</code>	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (<code>FALSE</code> ; default) or in the order of the time series (<code>TRUE</code>).
<code>ntrys</code>	Parameter used in rejection sampling for truncated discrete models only (defaults to 5). See Details for more information.
<code>cores</code>	Number of cores (defaults to 1). On non-Windows systems, this argument can be set globally via the <code>mc.cores</code> option.
<code>summary</code>	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .

robust	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Value

An array of predicted response values. If `summary` = FALSE the output resembles those of `posterior_predict.brmsfit`. If `summary` = TRUE the output depends on the family: For categorical and ordinal families, the output is an N x C matrix, where N is the number of observations, C is the number of categories, and the values are predicted category probabilities. For all other families, the output is a N x E matrix where E = 2 + `length(probs)` is the number of summary statistics: The Estimate column contains point estimates (either mean or median depending on argument `robust`), while the Est.Error column contains uncertainty estimates (either standard deviation or median absolute deviation depending on argument `robust`). The remaining columns starting with Q contain quantile estimates as specified via argument `probs`.

See Also

`posterior_predict.brmsfit`

Examples

```
## Not run:
## fit a model
fit <- brm(time | cens(censored) ~ age + sex + (1 + age || patient),
            data = kidney, family = "exponential", init = "0")

## predicted responses
pp <- predict(fit)
head(pp)

## predicted responses excluding the group-level effect of age
pp <- predict(fit, re_formula = ~ (1 | patient))
head(pp)

## predicted responses of patient 1 for new data
newdata <- data.frame(
  sex = factor(c("male", "female")),
  age = c(20, 50),
  patient = c(1, 1)
)
predict(fit, newdata = newdata)

## End(Not run)
```

`predictive_error.brmsfit`

Posterior Draws of Predictive Errors

Description

Compute posterior draws of predictive errors, that is, observed minus predicted responses. Can be performed for the data used to fit the model (posterior predictive checks) or for new data.

Usage

```
## S3 method for class 'brmsfit'
predictive_error(
  object,
  newdata = NULL,
  re_formula = NULL,
  re.form = NULL,
  method = "posterior_predict",
  resp = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>re.form</code>	Alias of <code>re_formula</code> .
<code>method</code>	Method used to obtain predictions. Can be set to <code>"posterior_predict"</code> (the default), <code>"posterior_epred"</code> , or <code>"posterior_linpred"</code> . For more details, see the respective function documentations.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>ndraws</code>	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .

draw_ids	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
sort	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (<code>FALSE</code> ; default) or in the order of the time series (<code>TRUE</code>).
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Value

An $S \times N$ array of predictive error draws, where S is the number of posterior draws and N is the number of observations.

Examples

```
## Not run:
## fit a model
fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler, cores = 2)

## extract predictive errors
pe <- predictive_error(fit)
str(pe)

## End(Not run)
```

predictive_interval.brmsfit
Predictive Intervals

Description

Compute intervals from the posterior predictive distribution.

Usage

```
## S3 method for class 'brmsfit'
predictive_interval(object, prob = 0.9, ...)
```

Arguments

object	An R object of class <code>brmsfit</code> .
prob	A number p ($0 < p < 1$) indicating the desired probability mass to include in the intervals. Defaults to <code>0.9</code> .
...	Further arguments passed to <code>posterior_predict</code> .

Value

A matrix with 2 columns for the lower and upper bounds of the intervals, respectively, and as many rows as observations being predicted.

Examples

```
## Not run:
fit <- brm(count ~ zBase, data = epilepsy, family = poisson())
predictive_interval(fit)

## End(Not run)
```

prepare_predictions.brmsfit
Prepare Predictions

Description

This method helps in preparing **brms** models for certain post-processing tasks most notably various forms of predictions. Unless you are a package developer, you will rarely need to call `prepare_predictions` directly.

Usage

```
## S3 method for class 'brmsfit'
prepare_predictions(
  x,
  newdata = NULL,
  re_formula = NULL,
  allow_new_levels = FALSE,
  sample_new_levels = "uncertainty",
  incl_autocor = TRUE,
  oos = NULL,
  resp = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  nsamples = NULL,
  subset = NULL,
  nug = NULL,
  smooths_only = FALSE,
  offset = TRUE,
  newdata2 = NULL,
  new_objects = NULL,
  point_estimate = NULL,
  ndraws_point_estimate = 1,
  ...
```

```
)
  prepare_predictions(x, ...)
```

Arguments

x	An R object typically of class 'brmsfit'.
newdata	An optional data.frame for which to evaluate predictions. If NULL (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
re_formula	formula containing group-level effects to be considered in the prediction. If NULL (default), include all group-level effects; if NA or ~0, include no group-level effects.
allow_new_levels	A flag indicating if new levels of group-level effects are allowed (defaults to FALSE). Only relevant if newdata is provided.
sample_new_levels	Indicates how to sample new levels for grouping factors specified in re_formula. This argument is only relevant if newdata is provided and allow_new_levels is set to TRUE. If "uncertainty" (default), each posterior sample for a new level is drawn from the posterior draws of a randomly chosen existing level. Each posterior sample for a new level may be drawn from a different existing level such that the resulting set of new posterior draws represents the variation across existing levels. If "gaussian", sample new levels from the (multivariate) normal distribution implied by the group-level standard deviations and correlations. This options may be useful for conducting Bayesian power analysis or predicting new levels in situations where relatively few levels were observed in the old_data. If "old_levels", directly sample new levels from the existing levels, where a new level is assigned all of the posterior draws of the same (randomly chosen) existing level.
incl_autocor	A flag indicating if correlation structures originally specified via autocor should be included in the predictions. Defaults to TRUE.
oos	Optional indices of observations for which to compute out-of-sample rather than in-sample predictions. Only required in models that make use of response values to make predictions, that is, currently only ARMA models.
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
ndraws	Positive integer indicating how many posterior draws should be used. If NULL (the default) all draws are used. Ignored if draw_ids is not NULL.
draw_ids	An integer vector specifying the posterior draws to be used. If NULL (the default), all draws are used.
nsamples	Deprecated alias of ndraws.
subset	Deprecated alias of draw_ids.

<code>nug</code>	Small positive number for Gaussian process terms only. For numerical reasons, the covariance matrix of a Gaussian process might not be positive definite. Adding a very small number to the matrix's diagonal often solves this problem. If <code>NULL</code> (the default), <code>nug</code> is chosen internally.
<code>smooths_only</code>	Logical; If <code>TRUE</code> only predictions related to smoothing splines (i.e., <code>s</code> or <code>t2</code>) will be computed. Defaults to <code>FALSE</code> .
<code>offset</code>	Logical; Indicates if offsets should be included in the predictions. Defaults to <code>TRUE</code> .
<code>newdata2</code>	A named list of objects containing new data, which cannot be passed via argument <code>newdata</code> . Required for some objects used in autocorrelation structures, or <code>stanvars</code> .
<code>new_objects</code>	Deprecated alias of <code>newdata2</code> .
<code>point_estimate</code>	Shall the returned object contain only point estimates of the parameters instead of their posterior draws? Defaults to <code>NULL</code> in which case no point estimate is computed. Alternatively, may be set to <code>"mean"</code> or <code>"median"</code> . This argument is primarily implemented to ensure compatibility with the <code>loo_subsample</code> method.
<code>ndraws_point_estimate</code>	Only used if <code>point_estimate</code> is not <code>NULL</code> . How often shall the point estimate's value be repeated? Defaults to 1.
<code>...</code>	Further arguments passed to <code>validate_newdata</code> .

Value

An object of class '`brmsprep`' or '`mvbrmsprep`', depending on whether a univariate or multivariate model is passed.

`print.brmsfit`

Print a summary for a fitted model represented by a `brmsfit` object

Description

Print a summary for a fitted model represented by a `brmsfit` object

Usage

```
## S3 method for class 'brmsfit'
print(x, digits = 2, short = getOption("brms.short_summary", FALSE), ...)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code>
<code>digits</code>	The number of significant digits for printing out the summary; defaults to 2. The effective sample size is always rounded to integers.

- short A flag indicating whether to provide a shorter summary with less informational text. Defaults to FALSE. Can be set globally for the current session via the `brms.short_summary` option.
- ... Additional arguments that would be passed to method `summary` of `brmsfit`.

See Also

[summary.brmsfit](#)

print.brmsprior *Print method for brmsprior objects*

Description

Print method for `brmsprior` objects

Usage

```
## S3 method for class 'brmsprior'  
print(x, show_df = NULL, ...)
```

Arguments

- x An object of class `brmsprior`.
- show_df Logical; Print priors as a single `data.frame` (TRUE) or as a sequence of sampling statements (FALSE)?
- ... Currently ignored.

prior_draws.brmsfit *Extract Prior Draws*

Description

Extract prior draws of specified parameters

Usage

```
## S3 method for class 'brmsfit'  
prior_draws(x, variable = NULL, pars = NULL, ...)  
  
prior_draws(x, ...)  
  
prior_samples(x, ...)
```

Arguments

- x An R object typically of class `brmsfit`.
- variable A character vector providing the variables to extract. By default, all variables are extracted.
- pars Deprecated alias of `variable`. For reasons of backwards compatibility, `pars` is interpreted as a vector of regular expressions by default unless `fixed = TRUE` is specified.
- ... Arguments passed to individual methods (if applicable).

Details

To make use of this function, the model must contain draws of prior distributions. This can be ensured by setting `sample_prior = TRUE` in function `brm`. Priors of certain parameters cannot be saved for technical reasons. For instance, this is the case for the population-level intercept, which is only computed after fitting the model by default. If you want to treat the intercept as part of all the other regression coefficients, so that sampling from its prior becomes possible, use `... ~ 0 + Intercept + ...` in the formulas.

Value

A `data.frame` containing the prior draws.

Examples

```
## Not run:
fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler, family = "cumulative",
            prior = set_prior("normal(0,2)", class = "b"),
            sample_prior = TRUE)

# extract all prior draws
draws1 <- prior_draws(fit)
head(draws1)

# extract prior draws for the coefficient of 'treat'
draws2 <- prior_draws(fit, "b_treat")
head(draws2)

## End(Not run)
```

prior_summary.brmsfit *Priors of brms models*

Description

Extract priors of models fitted with **brms**.

Usage

```
## S3 method for class 'brmsfit'
prior_summary(object, all = TRUE, ...)
```

Arguments

- object** An object of class `brmsfit`.
- all** Logical; Show all parameters in the model which may have priors (TRUE) or only those with proper priors (FALSE)?
- ...** Further arguments passed to or from other methods.

Value

An `brmsprior` object.

Examples

```
## Not run:
fit <- brm(
  count ~ zAge + zBase * Trt + (1|patient) + (1|obs),
  data = epilepsy, family = poisson(),
  prior = prior(student_t(5,0,10), class = b) +
    prior(cauchy(0,2), class = sd)
)

prior_summary(fit)
prior_summary(fit, all = FALSE)
print(prior_summary(fit, all = FALSE), show_df = FALSE)

## End(Not run)
```

Description

Implementation of Pareto smoothed importance sampling (PSIS), a method for stabilizing importance ratios. The version of PSIS implemented here corresponds to the algorithm presented in Vehtari, Simpson, Gelman, Yao, and Gabry (2024). For PSIS diagnostics see the [pareto-k-diagnostic](#) page.

Usage

```
## S3 method for class 'brmsfit'
psis(log_ratios, newdata = NULL, resp = NULL, model_name = NULL, ...)
```

Arguments

<code>log_ratios</code>	A fitted model object of class <code>brmsfit</code> . Argument is named "log_ratios" to match the argument name of the <code>loo::psis</code> generic function.
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If NULL (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>model_name</code>	Currently ignored.
<code>...</code>	Further arguments passed to <code>log_liik</code> and <code>loo::psis</code> .

Value

The `psis()` methods return an object of class "psis", which is a named list with the following components:

`log_weights` Vector or matrix of smoothed (and truncated) but *unnormalized* log weights. To get normalized weights use the `weights()` method provided for objects of class "psis".

`diagnostics` A named list containing two vectors:

- `pareto_k`: Estimates of the shape parameter k of the generalized Pareto distribution. See the [pareto-k-diagnostic](#) page for details.
- `n_eff`: PSIS effective sample size estimates.

Objects of class "psis" also have the following [attributes](#):

`norm_const_log` Vector of precomputed values of `colLogSumExps(log_weights)` that are used internally by the `weights` method to normalize the log weights.

`tail_len` Vector of tail lengths used for fitting the generalized Pareto distribution.

`r_eff` If specified, the user's `r_eff` argument.

`dims` Integer vector of length 2 containing S (posterior sample size) and N (number of observations).

`method` Method used for importance sampling, here `psis`.

References

Vehtari, A., Gelman, A., and Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27(5), 1413–1432. doi:10.1007/s11222-016-9696-4 ([journal version](#), [preprint arXiv:1507.04544](#)).

Vehtari, A., Simpson, D., Gelman, A., Yao, Y., and Gabry, J. (2024). Pareto smoothed importance sampling. *Journal of Machine Learning Research*, 25(72):1-58. [PDF](#)

Examples

```
## Not run:
fit <- brm(rating ~ treat + period + carry, data = inhaler)
psis(fit)

## End(Not run)
```

Description

Function used to set up R2D2(M2) priors in **brms**. The function does not evaluate its arguments – it exists purely to help set up the model.

Usage

```
R2D2(mean_R2 = 0.5, prec_R2 = 2, cons_D2 = 0.5, autoscale = TRUE, main = FALSE)
```

Arguments

<code>mean_R2</code>	Mean of the Beta prior on the coefficient of determination R^2 .
<code>prec_R2</code>	Precision of the Beta prior on the coefficient of determination R^2 .
<code>cons_D2</code>	Concentration vector of the Dirichlet prior on the variance decomposition parameters. Lower values imply more shrinkage.
<code>autoscale</code>	Logical; indicating whether the R2D2 prior should be scaled using the residual standard deviation <code>sigma</code> if possible and sensible (defaults to TRUE). Autoscaling is not applied for distributional parameters or when the model does not contain the parameter <code>sigma</code> .
<code>main</code>	Logical (defaults to FALSE); only relevant if the R2D2 prior spans multiple parameter classes. In this case, only arguments given in the single instance where <code>main</code> is TRUE will be used. Arguments given in other instances of the prior will be ignored. See the Examples section below.

Details

The prior does not account for scale differences of the terms it is applied on. Accordingly, please make sure that all these terms have a comparable scale to ensure that shrinkage is applied properly.

Currently, the following classes support the R2D2(M2) prior: `b` (overall regression coefficients), `sds` (SDs of smoothing splines), `sdgp` (SDs of Gaussian processes), `ar` (autoregressive coefficients), `ma` (moving average coefficients), `sderr` (SD of latent residuals), `sdcar` (SD of spatial CAR structures), `sd` (SD of varying coefficients).

When the prior is only applied to parameter class `b`, it is equivalent to the original R2D2 prior (with Gaussian kernel). When the prior is also applied to other parameter classes, it is equivalent to the R2D2M2 prior.

Even when the R2D2(M2) prior is applied to multiple parameter classes at once, the concentration vector (argument `cons_D2`) has to be provided jointly in the one instance of the prior where `main = TRUE`. The order in which the elements of concentration vector correspond to the classes' coefficients is the same as the order of the classes provided above.

References

Zhang, Y. D., Naughton, B. P., Bondell, H. D., & Reich, B. J. (2020). Bayesian regression using a prior on the model fit: The R2-D2 shrinkage prior. *Journal of the American Statistical Association*. <https://arxiv.org/pdf/1609.00046>

Aguilar J. E. & Bürkner P. C. (2022). Intuitive Joint Priors for Bayesian Linear Multilevel Models: The R2D2M2 prior. ArXiv preprint. <https://arxiv.org/pdf/2208.07132>

See Also

[set_prior](#)

Examples

```
set_prior(R2D2(mean_R2 = 0.8, prec_R2 = 10))

# specify the R2D2 prior across multiple parameter classes
set_prior(R2D2(mean_R2 = 0.8, prec_R2 = 10, main = TRUE), class = "b") +
  set_prior(R2D2(), class = "sd")
```

Description

Extract the group-level ('random') effects of each level from a `brmsfit` object.

Usage

```
## S3 method for class 'brmsfit'
ranef(
  object,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  pars = NULL,
  groups = NULL,
  ...
)
```

Arguments

object	An object of class <code>brmsfit</code> .
summary	Should summary statistics be returned instead of the raw values? Default is TRUE.
robust	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
pars	Optional names of coefficients to extract. By default, all coefficients are extracted.
groups	Optional names of grouping variables for which to extract effects.
...	Currently ignored.

Value

A list of 3D arrays (one per grouping factor). If `summary` is TRUE, the 1st dimension contains the factor levels, the 2nd dimension contains the summary statistics (see [posterior_summary](#)), and the 3rd dimension contains the group-level effects. If `summary` is FALSE, the 1st dimension contains the posterior draws, the 2nd dimension contains the factor levels, and the 3rd dimension contains the group-level effects.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt + (1+Trt|visit),
            data = epilepsy, family = gaussian(), chains = 2)
ranef(fit)

## End(Not run)
```

`read_csv_as_stanfit` *Read CmdStan CSV files as a brms-formatted stanfit object*

Description

`read_csv_as_stanfit` is used internally to read CmdStan CSV files into a `stanfit` object that is consistent with the structure of the `fit` slot of a `brmsfit` object.

Usage

```
read_csv_as_stanfit(
  files,
  variables = NULL,
  sampler_diagnostics = NULL,
  model = NULL,
  exclude = "",
  algorithm = "sampling"
)
```

Arguments

<code>files</code>	Character vector of CSV files names where draws are stored.
<code>variables</code>	Character vector of variables to extract from the CSV files.
<code>sampler_diagnostics</code>	Character vector of sampler diagnostics to extract.
<code>model</code>	A compiled cmdstanr model object (optional). Provide this argument if you want to allow updating the model without recompilation.
<code>exclude</code>	Character vector of variables to exclude from the stanfit. Only used when <code>variables</code> is also specified.
<code>algorithm</code>	The algorithm with which the model was fitted. See brm for details.

Value

A stanfit object consistent with the structure of the `fit` slot of a brmsfit object.

Examples

```
## Not run:
# fit a model manually via cmdstanr
scode <- stancode(count ~ Trt, data = epilepsy)
sdata <- standata(count ~ Trt, data = epilepsy)
mod <- cmdstanr::cmdstan_model(cmdstanr::write_stan_file(scode))
stanfit <- mod$sample(data = sdata)

# feed the Stan model back into brms
fit <- brm(count ~ Trt, data = epilepsy, empty = TRUE, backend = 'cmdstanr')
fit$fit <- read_csv_as_stanfit(stanfit$output_files(), model = mod)
fit <- rename_pars(fit)
summary(fit)

## End(Not run)
```

recompile_model	<i>Recompile Stan models in brmsfit objects</i>
-----------------	---

Description

Recompile the Stan model inside a `brmsfit` object, if necessary. This does not change the model, it simply recreates the executable so that sampling is possible again.

Usage

```
recompile_model(x, recompile = NULL)
```

Arguments

- | | |
|-----------|---|
| x | An object of class <code>brmsfit</code> . |
| recompile | Logical, indicating whether the Stan model should be recompiled. If <code>NULL</code> (the default), <code>recompile_model</code> tries to figure out internally, if recompilation is necessary. Setting it to <code>FALSE</code> will cause <code>recompile_model</code> to always return the <code>brmsfit</code> object unchanged. |

Value

A (possibly updated) `brmsfit` object.

reloo.brmsfit	<i>Compute exact cross-validation for problematic observations</i>
---------------	--

Description

Compute exact cross-validation for problematic observations for which approximate leave-one-out cross-validation may return incorrect results. Models for problematic observations can be run in parallel using the **future** package.

Usage

```
## S3 method for class 'brmsfit'
reloo(
  x,
  loo = NULL,
  k_threshold = 0.7,
  newdata = NULL,
  resp = NULL,
  check = TRUE,
  recompile = NULL,
  future_args = list(),
```

```

  ...
)

## S3 method for class 'loo'
reloo(x, fit, ...)

reloo(x, ...)

```

Arguments

<code>x</code>	An R object of class <code>brmsfit</code> or <code>loo</code> depending on the method.
<code>loo</code>	An R object of class <code>loo</code> . If <code>NULL</code> , <code>brms</code> will try to extract a precomputed <code>loo</code> object from the fitted model, added there via add_criterion .
<code>k_threshold</code>	The threshold at which Pareto k estimates are treated as problematic. Defaults to <code>0.7</code> . See pareto_k_ids for more details.
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>check</code>	Logical; If <code>TRUE</code> (the default), some checks check are performed if the <code>loo</code> object was generated from the <code>brmsfit</code> object passed to argument <code>fit</code> .
<code>recompile</code>	Logical, indicating whether the Stan model should be recompiled. This may be necessary if you are running <code>reloo</code> on another machine than the one used to fit the model.
<code>future_args</code>	A list of further arguments passed to future for additional control over parallel execution if activated.
<code>...</code>	Further arguments passed to update.brmsfit and log_lik.brmsfit .
<code>fit</code>	An R object of class <code>brmsfit</code> .

Details

Warnings about Pareto k estimates indicate observations for which the approximation to LOO is problematic (this is described in detail in Vehtari, Gelman, and Gabry (2017) and the `loo` package documentation). If there are J observations with k estimates above `k_threshold`, then `reloo` will refit the original model J times, each time leaving out one of the J problematic observations. The pointwise contributions of these observations to the total ELPD are then computed directly and substituted for the previous estimates from these J observations that are stored in the original `loo` object.

By default, this method uses `sample_new_levels = "gaussian"` to sample parameter values for new grouping-factor levels (see also [prepare_predictions](#)). This default will fail for models with non-Gaussian group-level effects. In this case, we recommend setting `sample_new_levels = "uncertainty"`.

Value

An object of the class `loo`.

Parallelization with multiple CPU cores

brms can make use of multiple CPU cores in parallel to speed up computations in various ways. For efficient use of the available resources it is recommended to only use parallelism to an extend such that the available physical CPUs are not oversubscribed. For example, when you have 8 CPU cores locally available, then you may consider to run 4 chains with 2 threads per chain for best performance if you happen to just run a single model. In case you run a simulation study which requires to run many times a given model, then neither chain nor within-chain parallelization is advisable as the computational resources are already exhausted by the simulation study and any further parallelization beyond the simulation study itself will in fact slow down the overall runtime. Please be aware that for historical reasons the nomenclature of the arguments is possibly confusing. The `cores` argument refers to running different chains in parallel and the within-chain parallelization will allocate for each chain as many threads as requested. The requested threads therefore increase the use of overall CPUs in a multiplicative way.

For more advanced parallelization (including beyond single model fits), **brms** also integrates with the **future** package. Importantly, this enables seamless integration with the **mirai** parallelization framework through the use of the `future.mirai` adapter. With **mirai** local and remote machines can be used in a fully transparent manner to the user. This includes the possibility to use large number of remote machines running in the context of a computer cluster, which are managed with queuing systems. Please refer to the section on distributed computing of [mirai::daemons](#).

See Also

[loo](#), [kfold](#)

Examples

```
## Not run:
fit1 <- brm(count ~ zAge + zBase * Trt + (1|patient),
             data = epilepsy, family = poisson())

# throws warning about some pareto k estimates being too high
(loo1 <- loo(fit1))

# no more warnings after reloo
(reloo1 <- reloo(fit1, loo = loo1, chains = 1))

## End(Not run)
```

Description

Rename parameters within the `stanfit` object after model fitting to ensure reasonable parameter names. This function is usually called automatically by `brm` and users will rarely be required to call it themselves.

Usage

```
rename_pars(x)
```

Arguments

x	A <code>brmsfit</code> object.
---	--------------------------------

Details

Function `rename_pars` is a deprecated alias of `rename_pars`.

Value

A `brmsfit` object with adjusted parameter names.

Examples

```
## Not run:
# fit a model manually via rstan
scode <- stancode(count ~ Trt, data = epilepsy)
sdata <- standata(count ~ Trt, data = epilepsy)
stanfit <- rstan::stan(model_code = scode, data = sdata)

# feed the Stan model back into brms
fit <- brm(count ~ Trt, data = epilepsy, empty = TRUE)
fit$fit <- stanfit
fit <- rename_pars(fit)
summary(fit)

## End(Not run)
```

Description

This method is an alias of `predictive_error.brmsfit` with additional arguments for obtaining summaries of the computed draws.

Usage

```
## S3 method for class 'brmsfit'
residuals(
  object,
  newdata = NULL,
  re_formula = NULL,
  method = "posterior_predict",
  type = c("ordinary", "pearson"),
  resp = NULL,
  ndraws = NULL,
  draw_ids = NULL,
  sort = FALSE,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If <code>NULL</code> (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>method</code>	Method used to obtain predictions. Can be set to <code>"posterior_predict"</code> (the default), <code>"posterior_epred"</code> , or <code>"posterior_linpred"</code> . For more details, see the respective function documentations.
<code>type</code>	The type of the residuals, either <code>"ordinary"</code> or <code>"pearson"</code> . More information is provided under 'Details'.
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>ndraws</code>	Positive integer indicating how many posterior draws should be used. If <code>NULL</code> (the default) all draws are used. Ignored if <code>draw_ids</code> is not <code>NULL</code> .
<code>draw_ids</code>	An integer vector specifying the posterior draws to be used. If <code>NULL</code> (the default), all draws are used.
<code>sort</code>	Logical. Only relevant for time series models. Indicating whether to return predicted values in the original order (<code>FALSE</code> ; default) or in the order of the time series (<code>TRUE</code>).
<code>summary</code>	Should summary statistics be returned instead of the raw values? Default is <code>TRUE</code> .

robust	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
probs	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
...	Further arguments passed to <code>prepare_predictions</code> that control several aspects of data validation and prediction.

Details

Residuals of type 'ordinary' are of the form $R = Y - Y_{rep}$, where Y is the observed and Y_{rep} is the predicted response. Residuals of type `pearson` are of the form $R = (Y - Y_{rep})/SD(Y_{rep})$, where $SD(Y_{rep})$ is an estimate of the standard deviation of Y_{rep} .

Value

An array of predictive error/residual draws. If `summary` = FALSE the output resembles those of `predictive_error.brmsfit`. If `summary` = TRUE the output is an N x E matrix, where N is the number of observations and E denotes the summary statistics computed from the draws.

Examples

```
## Not run:
## fit a model
fit <- brm(rating ~ treat + period + carry + (1|subject),
            data = inhaler, cores = 2)

## extract residuals/predictive errors
res <- residuals(fit)
head(res)

## End(Not run)
```

Description

`restructure` is a generic function used to restructure old R objects to work with newer versions of the package that generated them. Its original use is within the `brms` package, but new methods for use with objects from other packages can be registered to the same generic.

Usage

```
restructure(x, ...)
```

Arguments

- x An object to be restructured. The object's class will determine which method to apply
- ... Additional arguments to pass to the specific methods

Details

Usually the version of the package that generated the object will be stored somewhere in the object and this information will be used by the specific method to determine what transformations to apply. See [restructure.brmsfit](#) for the default method applied for **brms** models. You can view the available methods by typing: `methods(restructure)`

Value

An object of the same class as x compatible with the latest version of the package that generated it.

See Also

[restructure.brmsfit](#)

restructure.brmsfit *Restructure Old brmsfit Objects*

Description

Restructure old `brmsfit` objects to work with the latest **brms** version. This function is called internally when applying post-processing methods. However, in order to avoid unnecessary run time caused by the restructuring, I recommend explicitly calling `restructure` once per model after updating **brms**.

Usage

```
## S3 method for class 'brmsfit'  
restructure(x, ...)
```

Arguments

- x An object of class `brmsfit`.
- ... Currently ignored.

Details

If you are restructuring an old spline model (fitted with `brms < 2.19.3`) to avoid prediction inconsistencies between machines (see GitHub issue #1465), please make sure to `restructure` your model on the machine on which it was originally fitted.

Value

A `brmsfit` object compatible with the latest version of **brms**.

rows2labels	<i>Convert Rows to Labels</i>
-------------	-------------------------------

Description

Convert information in rows to labels for each row.

Usage

```
rows2labels(x, digits = 2, sep = " & ", incl_vars = TRUE, ...)
```

Arguments

<code>x</code>	A <code>data.frame</code> for which to extract labels.
<code>digits</code>	Minimal number of decimal places shown in the labels of numeric variables.
<code>sep</code>	A single character string defining the separator between variables used in the labels.
<code>incl_vars</code>	Indicates if variable names should be part of the labels. Defaults to TRUE.
<code>...</code>	Currently unused.

Value

A character vector of the same length as the number of rows of `x`.

See Also

[make_conditions](#), [conditional_effects](#)

s	<i>Defining smooths in brms formulas</i>
---	--

Description

Functions used in definition of smooth terms within a model formulas. The function does not evaluate a (spline) smooth - it exists purely to help set up a model using spline based smooths.

Usage

```
s(...)
```

```
t2(...)
```

Arguments

... Arguments passed to `mgcv::s` or `mgcv::t2`.

Details

The function defined here are just simple wrappers of the respective functions of the `mgcv` package. When using them, please cite the appropriate references obtained via `citation("mgcv")`.

brms uses the "random effects" parameterization of smoothing splines as explained in `mgcv::gamm`. A nice tutorial on this topic can be found in Pedersen et al. (2019). The answers provided in this [Stan discourse post](#) may also be helpful.

References

Pedersen, E. J., Miller, D. L., Simpson, G. L., & Ross, N. (2019). Hierarchical generalized additive models in ecology: an introduction with mgcv. *PeerJ*.

See Also

`brmsformula`, `mgcv::s`, `mgcv::t2`

Examples

```
## Not run:
# simulate some data
dat <- mgcv::gamSim(1, n = 200, scale = 2)

# fit univariate smooths for all predictors
fit1 <- brm(y ~ s(x0) + s(x1) + s(x2) + s(x3),
             data = dat, chains = 2)
summary(fit1)
plot(conditional_smooths(fit1), ask = FALSE)

# fit a more complicated smooth model
fit2 <- brm(y ~ t2(x0, x1) + s(x2, by = x3),
             data = dat, chains = 2)
summary(fit2)
plot(conditional_smooths(fit2), ask = FALSE)

## End(Not run)
```

Description

Set up an spatial simultaneous autoregressive (SAR) term in **brms**. The function does not evaluate its arguments – it exists purely to help set up a model with SAR terms.

Usage

```
sar(M, type = "lag")
```

Arguments

- M An object specifying the spatial weighting matrix. Can be either the spatial weight matrix itself or an object of class `listw` or `nb`, from which the spatial weighting matrix can be computed.
- type Type of the SAR structure. Either "lag" (for SAR of the response values) or "error" (for SAR of the residuals). More information is provided in the 'Details' section.

Details

The `lagsar` structure implements SAR of the response values:

$$y = \rho W y + \eta + e$$

The `errorsar` structure implements SAR of the residuals:

$$y = \eta + u, u = \rho W u + e$$

In the above equations, η is the predictor term and e are independent normally or t-distributed residuals. Currently, only families `gaussian` and `student` support SAR structures.

Value

An object of class '`sar_term`', which is a list of arguments to be interpreted by the formula parsing functions of `brms`.

See Also

[autocor-terms](#)

Examples

```
## Not run:
data(olddcol, package = "spdep")
fit1 <- brm(CRIME ~ INC + HOVAL + sar(COL.nb, type = "lag"),
             data = COL.OLD, data2 = list(COL.nb = COL.nb),
             chains = 2, cores = 2)
summary(fit1)
plot(fit1)

fit2 <- brm(CRIME ~ INC + HOVAL + sar(COL.nb, type = "error"),
             data = COL.OLD, data2 = list(COL.nb = COL.nb),
             chains = 2, cores = 2)
summary(fit2)
plot(fit2)

## End(Not run)
```

save_pars*Control Saving of Parameter Draws*

Description

Control which (draws of) parameters should be saved in a **brms** model. The output of this function is meant for usage in the `save_pars` argument of **brm**.

Usage

```
save_pars(group = TRUE, latent = FALSE, all = FALSE, manual = NULL)
```

Arguments

group	A flag to indicate if group-level coefficients for each level of the grouping factors should be saved (default is TRUE). Set to FALSE to save memory. Alternatively, <code>group</code> may also be a character vector naming the grouping factors for which to save draws of coefficients.
latent	A flag to indicate if draws of latent variables obtained by using <code>me</code> and <code>mi</code> terms should be saved (default is FALSE). Saving these draws allows to better use methods such as <code>posterior_predict</code> with the latent variables but leads to very large R objects even for models of moderate size and complexity. Alternatively, <code>latent</code> may also be a character vector naming the latent variables for which to save draws.
all	A flag to indicate if draws of all variables defined in Stan's <code>parameters</code> block should be saved (default is FALSE). Saving these draws is required in order to apply the certain methods such as <code>bridge_sampler</code> and <code>bayes_factor</code> .
manual	A character vector naming Stan variable names which should be saved. These names should match the variable names inside the Stan code before renaming. This feature is meant for power users only and will rarely be useful outside of very special cases.

Value

A list of class "save_pars".

Examples

```
## Not run:
# don't store group-level coefficients
fit <- brm(count ~ zAge + zBase * Trt + (1|patient),
            data = epilepsy, family = poisson(),
            save_pars = save_pars(group = FALSE))
variables(fit)

## End(Not run)
```

set_prior*Prior Definitions for brms Models*

Description

Define priors for specific parameters or classes of parameters.

Usage

```
set_prior(
  prior,
  class = "b",
  coef = "",
  group = "",
  resp = "",
  dpar = "",
  npar = "",
  lb = NA,
  ub = NA,
  tag = "",
  check = TRUE
)

prior(prior, ...)

prior_(prior, ...)

prior_string(prior, ...)

empty_prior()
```

Arguments

prior	A character string defining a distribution in Stan language
class	The parameter class. Defaults to "b" (i.e. population-level effects). See 'Details' for other valid parameter classes.
coef	Name of the coefficient within the parameter class.
group	Grouping factor for group-level parameters.
resp	Name of the response variable. Only used in multivariate models.
dpar	Name of a distributional parameter. Only used in distributional models.
npar	Name of a non-linear parameter. Only used in non-linear models.
lb	Lower bound for parameter restriction. Currently only allowed for classes "b". Defaults to NULL, that is no restriction.
ub	Upper bound for parameter restriction. Currently only allowed for classes "b". Defaults to NULL, that is no restriction.

tag	Character to append to the lprior variable in the Stan code. Used for selectively checking sensitivity of priors in priorsense.
check	Logical; Indicates whether priors should be checked for validity (as far as possible). Defaults to TRUE. If FALSE, prior is passed to the Stan code as is, and all other arguments are ignored.
...	Arguments passed to set_prior.

Details

`set_prior` is used to define prior distributions for parameters in **brms** models. The functions `prior`, `prior_`, and `prior_string` are aliases of `set_prior` each allowing for a different kind of argument specification. `prior` allows specifying arguments as expression without quotation marks using non-standard evaluation. `prior_` allows specifying arguments as one-sided formulas or wrapped in quote. `prior_string` allows specifying arguments as strings just as `set_prior` itself.

Below, we explain its usage and list some common prior distributions for parameters. A complete overview on possible prior distributions is given in the Stan Reference Manual available at <https://mc-stan.org/>.

To combine multiple priors, use `c(...)` or the `+` operator (see 'Examples'). **brms** does not check if the priors are written in correct **Stan** language. Instead, **Stan** will check their syntactical correctness when the model is parsed to C++ and returns an error if they are not. This, however, does not imply that priors are always meaningful if they are accepted by **Stan**. Although **brms** tries to find common problems (e.g., setting bounded priors on unbounded parameters), there is no guarantee that the defined priors are reasonable for the model. Below, we list the types of parameters in **brms** models, for which the user can specify prior distributions.

Below, we provide details for the individual parameter classes that you can set priors on. Often, it may not be immediately clear, which parameters are present in the model. To get a full list of parameters and parameter classes for which priors can be specified (depending on the model) use function `default_prior`.

1. Population-level ('fixed') effects

Every Population-level effect has its own regression parameter represents the name of the corresponding population-level effect. Suppose, for instance, that `y` is predicted by `x1` and `x2` (i.e., `y ~ x1 + x2` in formula syntax). Then, `x1` and `x2` have regression parameters `b_x1` and `b_x2` respectively. The default prior for population-level effects (including monotonic and category specific effects) is an improper flat prior over the reals. Other common options are normal priors or student-t priors. If we want to have a normal prior with mean 0 and standard deviation 5 for `x1`, and a unit student-t prior with 10 degrees of freedom for `x2`, we can specify this via `set_prior("normal(0, 5)", class = "b", coef = "x1")` and `set_prior("student_t(10, 0, 1)", class = "b", coef = "x2")`. To put the same prior on all population-level effects at once, we may write as a shortcut `set_prior("<prior>", class = "b")`. This also leads to faster sampling, because priors can be vectorized in this case. Both ways of defining priors can be combined using for instance `set_prior("normal(0, 2)", class = "b")` and `set_prior("normal(0, 10)", class = "b", coef = "x1")` at the same time. This will set a `normal(0, 10)` prior on the effect of `x1` and a `normal(0, 2)` prior on all other population-level effects. However, this will break vectorization and may slow down the sampling procedure a bit.

In case of the default intercept parameterization (discussed in the 'Details' section of **brmsformula**), general priors on class "b" will *not* affect the intercept. Instead, the intercept has its own parameter

class named "Intercept" and priors can thus be specified via `set_prior("<prior>", class = "Intercept")`. Setting a prior on the intercept will not break vectorization of the other population-level effects. Note that technically, this prior is set on an intercept that results when internally centering all population-level predictors around zero to improve sampling efficiency. On this centered intercept, specifying a prior is actually much easier and intuitive than on the original intercept, since the former represents the expected response value when all predictors are at their means. To treat the intercept as an ordinary population-level effect and avoid the centering parameterization, use `0 + Intercept` on the right-hand side of the model formula.

In non-linear models, population-level effects are defined separately for each non-linear parameter. Accordingly, it is necessary to specify the non-linear parameter in `set_prior` so that priors we can be assigned correctly. If, for instance, `alpha` is the parameter and `x` the predictor for which we want to define the prior, we can write `set_prior("<prior>", coef = "x", npar = "alpha")`. As a shortcut we can use `set_prior("<prior>", npar = "alpha")` to set the same prior on all population-level effects of `alpha` at once.

The same goes for specifying priors for specific distributional parameters in the context of distributional regression, for example, `set_prior("<prior>", coef = "x", dpar = "sigma")`. For most other parameter classes (see below), you need to indicate non-linear and distributional parameters in the same way as shown here.

If desired, population-level effects can be restricted to fall only within a certain interval using the `lb` and `ub` arguments of `set_prior`. This is often required when defining priors that are not defined everywhere on the real line, such as uniform or gamma priors. When defining a `uniform(2, 4)` prior, you should write `set_prior("uniform(2, 4)", lb = 2, ub = 4)`. When using a prior that is defined on the positive reals only (such as a gamma prior) set `lb = 0`. In most situations, it is not useful to restrict population-level parameters through bounded priors (non-linear models are an important exception), but if you really want to this is the way to go.

2. Group-level ('random') effects

Each group-level effect of each grouping factor has a standard deviation named `sd_<group>_<coef>`. Consider, for instance, the formula `y ~ x1 + x2 + (1 + x1 | g)`. We see that the intercept as well as `x1` are group-level effects nested in the grouping factor `g`. The corresponding standard deviation parameters are named as `sd_g_Intercept` and `sd_g_x1` respectively. These parameters are restricted to be non-negative and, by default, have a half student-t prior with 3 degrees of freedom and a scale parameter that depends on the standard deviation of the response after applying the link function. Minimally, the scale parameter is 2.5. This prior is used (a) to be only weakly informative in order to influence results as few as possible, while (b) providing at least some regularization to considerably improve convergence and sampling efficiency. To define a prior distribution only for standard deviations of a specific grouping factor, use

`set_prior("<prior>", class = "sd", group = "<group>")`. To define a prior distribution only for a specific standard deviation of a specific grouping factor, you may write
`set_prior("<prior>", class = "sd", group = "<group>", coef = "<coef>")`.

If there is more than one group-level effect per grouping factor, the correlations between those effects have to be estimated. The prior `lkj_corr_cholesky(eta)` or in short `lkj(eta)` with `eta > 0` is essentially the only prior for (Cholesky factors) of correlation matrices. If `eta = 1` (the default) all correlations matrices are equally likely a priori. If `eta > 1`, extreme correlations become less likely, whereas `0 < eta < 1` results in higher probabilities for extreme correlations. Correlation matrix parameters in `brms` models are named as `cor_<group>`, (e.g., `cor_g` if `g` is the grouping factor). To set the same prior on every correlation matrix, use for instance `set_prior("lkj(2)", class = "cor")`. Internally, the priors are transformed to be put on the Cholesky factors of the correlation

matrices to improve efficiency and numerical stability. The corresponding parameter class of the Cholesky factors is `L`, but it is not recommended to specify priors for this parameter class directly.

4. Smoothing Splines

Smoothing splines are implemented in **brms** using the 'random effects' formulation as explained in `gamm`). Thus, each spline has its corresponding standard deviations modeling the variability within this term. In **brms**, this parameter class is called `sds` and priors can be specified via `set_prior("<prior>", class = "sds", coef = "<term label>")`. The default prior is the same as for standard deviations of group-level effects.

5. Gaussian processes

Gaussian processes as currently implemented in **brms** have two parameters, the standard deviation parameter `sdgp`, and characteristic length-scale parameter `lscale` (see `gp` for more details). The default prior of `sdgp` is the same as for standard deviations of group-level effects. The default prior of `lscale` is an informative inverse-gamma prior specifically tuned to the covariates of the Gaussian process (for more details see https://betanalpha.github.io/assets/case_studies/gp_part3/part3.html). This tuned prior may be overly informative in some cases, so please consider other priors as well to make sure inference is robust to the prior specification. If tuning fails, a half-normal prior is used instead.

6. Autocorrelation parameters

The autocorrelation parameters currently implemented are named `ar` (autoregression), `ma` (moving average), `sderr` (standard deviation of latent residuals in latent ARMA models), `cosy` (compound symmetry correlation), `car` (spatial conditional autoregression), as well as `lagsar` and `errorsar` (spatial simultaneous autoregression).

Priors can be defined by `set_prior("<prior>", class = "ar")` for `ar` and similar for other autocorrelation parameters. By default, `ar` and `ma` are bounded between -1 and 1; `cosy`, `car`, `lagsar`, and `errorsar` are bounded between 0 and 1. The default priors are flat over the respective definition areas.

7. Parameters of measurement error terms

Latent variables induced via measurement error `me` terms require both mean and standard deviation parameters, whose prior classes are named "`meanme`" and "`sdme`", respectively. If multiple latent variables are induced this way, their correlation matrix will be modeled as well and corresponding priors can be specified via the "`corme`" class. All of the above parameters have flat priors over their respective definition spaces by default.

8. Distance parameters of monotonic effects

As explained in the details section of `brm`, monotonic effects make use of a special parameter vector to estimate the 'normalized distances' between consecutive predictor categories. This is realized in **Stan** using the simplex parameter type. This class is named "`simo`" (short for simplex monotonic) in **brms**. The only valid prior for simplex parameters is the dirichlet prior, which accepts a vector of length $K - 1$ (K = number of predictor categories) as input defining the 'concentration' of the distribution. Explaining the dirichlet prior is beyond the scope of this documentation, but we want to describe how to define this prior syntactically correct. If a predictor `x` with K categories is modeled as monotonic, we can define a prior on its corresponding simplex via `prior(dirichlet(<vector>), class = simo, coef = mox1)`. The 1 in the end of `coef` indicates that this is the first simplex in this term. If interactions between multiple monotonic variables are modeled, multiple simplexes per term are required. For `<vector>`, we can put in any R expression

defining a vector of length $K - 1$. The default is a uniform prior (i.e. `<vector> = rep(1, K-1)`) over all simplexes of the respective dimension.

9. Parameters for specific families

Some families need additional parameters to be estimated. Families `gaussian`, `student`, `skew_normal`, `lognormal`, and `gen_extreme_value` need the parameter `sigma` to account for the residual standard deviation. By default, `sigma` has a half student-t prior that scales in the same way as the group-level standard deviations. Further, family `student` needs the parameter `nu` representing the degrees of freedom of Student-t distribution. By default, `nu` has prior `gamma(2, 0.1)`, which is close to a penalized complexity prior (see Stan prior choice Wiki), and a fixed lower bound of 1. Family `negbinomial` needs a shape parameter that has by default `inv_gamma(0.4, 0.3)` prior which is close to a penalized complexity prior (see Stan prior choice Wiki). Families `gamma`, `weibull`, and `inverse.gaussian`, need a shape parameter that has a `gamma(0.01, 0.01)` prior by default. For families `cumulative`, `cratio`, `sratio`, and `acat`, and only if `threshold = "equidistant"`, the parameter `delta` is used to model the distance between two adjacent thresholds. By default, `delta` has an improper flat prior over the reals. The `von_mises` family needs the parameter `kappa`, representing the concentration parameter. By default, `kappa` has prior `gamma(2, 0.01)`.

Every family specific parameter has its own prior class, so that `set_prior("<prior>", class = "<parameter>")` is the right way to go. All of these priors are chosen to be weakly informative, having only minimal influence on the estimations, while improving convergence and sampling efficiency.

10. Shrinkage priors

To reduce the danger of overfitting in models with many predictor terms fit on comparably sparse data, `brms` supports special shrinkage priors, namely the (regularized) `horseshoe` and the `R2D2` prior. These priors can be applied on many parameter classes, either directly on the coefficient classes (e.g., class `b`), if directly setting priors on them is supported, or on the corresponding standard deviation hyperparameters (e.g., class `sd`) otherwise. Currently, the following classes support shrinkage priors: `b` (overall regression coefficients), `sds` (SDs of smoothing splines), `sdgp` (SDs of Gaussian processes), `ar` (autoregressive coefficients), `ma` (moving average coefficients), `sder` (SD of latent residuals), `sdcar` (SD of spatial CAR structures), `sd` (SD of varying coefficients).

11. Fixing parameters to constants

Fixing parameters to constants is possible by using the `constant` function, for example, `constant(1)` to fix a parameter to 1. Broadcasting to vectors and matrices is done automatically.

Value

An object of class `brmsprior` to be used in the `prior` argument of `brm`.

Functions

- `prior()`: Alias of `set_prior` allowing to specify arguments as expressions without quotation marks.
- `prior_()`: Alias of `set_prior` allowing to specify arguments as one-sided formulas or wrapped in quote.
- `prior_string()`: Alias of `set_prior` allowing to specify arguments as strings.
- `empty_prior()`: Create an empty `brmsprior` object.

See Also

[default_prior](#)

Examples

```

## use alias functions
(prior1 <- prior(cauchy(0, 1), class = "sd"))
(prior2 <- prior_(-cauchy(0, 1), class = "sd"))
(prior3 <- prior_string("cauchy(0, 1)", class = "sd"))
identical(prior1, prior2)
identical(prior1, prior3)

# check which parameters can have priors
default_prior(rating ~ treat + period + carry + (1|subject),
               data = inhaler, family = cumulative())

# define some priors
bprior <- c(prior_string("normal(0,10)", class = "b"),
            prior(normal(1,2), class = "b", coef = "treat"),
            prior_(-cauchy(0,2), class = "sd",
                   group = -subject, coef = -Intercept))

# verify that the priors indeed found their way into Stan's model code
stancode(rating ~ treat + period + carry + (1|subject),
          data = inhaler, family = cumulative(),
          prior = bprior)

# use the horseshoe prior to model sparsity in regression coefficients
stancode(count ~ zAge + zBase * Trt,
         data = epilepsy, family = poisson(),
         prior = set_prior("horseshoe(3)"))

# fix certain priors to constants
bprior <- prior(constant(1), class = "b") +
  prior(constant(2), class = "b", coef = "zBase") +
  prior(constant(0.5), class = "sd")
stancode(count ~ zAge + zBase + (1 | patient),
         data = epilepsy, prior = bprior)

# pass priors to Stan without checking
prior <- prior_string("target += normal_lpdf(b[1] | 0, 1)", check = FALSE)
stancode(count ~ Trt, data = epilepsy, prior = prior)

# define priors in a vectorized manner
# useful in particular for categorical or multivariate models
set_prior("normal(0, 2)", dpar = c("muX", "muY", "muZ"))

# specify tags for different priors for sensitivity analysis using priorsense
# It is then possible to check the sensitivity when changing the priors with
# the same tag while leaving others the same
prior_cov <- prior(normal(0, 10), class = "b", tag = "covariates")
prior_trt <- prior(normal(0, 1), class = "b", coef = "Trt1", tag = "treatment")

```

```
stancode(count ~ Trt + zAge + zBase + (1 | patient),
         data = epilepsy, prior = c(prior_cov, prior_trt))
```

Shifted_Lognormal *The Shifted Log Normal Distribution*

Description

Density, distribution function, quantile function and random generation for the shifted log normal distribution with mean `meanlog`, standard deviation `sdlog`, and shift parameter `shift`.

Usage

```
dshifted_lnorm(x, meanlog = 0, sdlog = 1, shift = 0, log = FALSE)

pshifted_lnorm(
  q,
  meanlog = 0,
  sdlog = 1,
  shift = 0,
  lower.tail = TRUE,
  log.p = FALSE
)

qshifted_lnorm(
  p,
  meanlog = 0,
  sdlog = 1,
  shift = 0,
  lower.tail = TRUE,
  log.p = FALSE
)

rshifted_lnorm(n, meanlog = 0, sdlog = 1, shift = 0)
```

Arguments

<code>x, q</code>	Vector of quantiles.
<code>meanlog</code>	Vector of means.
<code>sdlog</code>	Vector of standard deviations.
<code>shift</code>	Vector of shifts.
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If TRUE, values are returned on the log scale.
<code>p</code>	Vector of probabilities.
<code>n</code>	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

SkewNormal

The Skew-Normal Distribution

Description

Density, distribution function, and random generation for the skew-normal distribution with mean `mu`, standard deviation `sigma`, and skewness `alpha`.

Usage

```
dskew_normal(  
  x,  
  mu = 0,  
  sigma = 1,  
  alpha = 0,  
  xi = NULL,  
  omega = NULL,  
  log = FALSE  
)  
  
pskew_normal(  
  q,  
  mu = 0,  
  sigma = 1,  
  alpha = 0,  
  xi = NULL,  
  omega = NULL,  
  lower.tail = TRUE,  
  log.p = FALSE  
)  
  
qskew_normal(  
  p,  
  mu = 0,  
  sigma = 1,  
  alpha = 0,  
  xi = NULL,  
  omega = NULL,  
  lower.tail = TRUE,  
  log.p = FALSE,  
  tol = 1e-08  
)  
  
rskew_normal(n, mu = 0, sigma = 1, alpha = 0, xi = NULL, omega = NULL)
```

Arguments

<code>x, q</code>	Vector of quantiles.
<code>mu</code>	Vector of mean values.
<code>sigma</code>	Vector of standard deviation values.
<code>alpha</code>	Vector of skewness values.
<code>xi</code>	Optional vector of location values. If <code>NULL</code> (the default), will be computed internally.
<code>omega</code>	Optional vector of scale values. If <code>NULL</code> (the default), will be computed internally.
<code>log</code>	Logical; If <code>TRUE</code> , values are returned on the log scale.
<code>lower.tail</code>	Logical; If <code>TRUE</code> (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If <code>TRUE</code> , values are returned on the log scale.
<code>p</code>	Vector of probabilities.
<code>tol</code>	Tolerance of the approximation used in the computation of quantiles.
<code>n</code>	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

Description

`stancode` is a generic function that can be used to generate Stan code for Bayesian models. Its original use is within the **brms** package, but new methods for use with objects from other packages can be registered to the same generic.

Usage

```
stancode(object, ...)
make_stancode(formula, ...)
```

Arguments

<code>object</code>	An object whose class will determine which method to apply. Usually, it will be some kind of symbolic description of the model form which Stan code should be generated.
<code>...</code>	Further arguments passed to the specific method.
<code>formula</code>	Synonym of <code>object</code> for use in <code>make_stancode</code> .

Details

See [stancode.default](#) for the default method applied for **brms** models. You can view the available methods by typing: `methods(stancode)` The `make_stancode` function is an alias of `stancode`.

Value

Usually, a character string containing the generated Stan code. For pretty printing, we recommend the returned object to be of class `c("character", "brmsmodel")`.

See Also

[stancode.default](#), [stancode.brmsfit](#)

Examples

```
stancode(rating ~ treat + period + carry + (1|subject),
          data = inhaler, family = "cumulative")
```

`stancode.brmsfit` *Extract Stan code from brmsfit objects*

Description

Extract Stan code from a fitted **brms** model.

Usage

```
## S3 method for class 'brmsfit'
stancode(
  object,
  version = TRUE,
  regenerate = NULL,
  threads = NULL,
  backend = NULL,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>version</code>	Logical; indicates if the first line containing the brms version number should be included. Defaults to <code>TRUE</code> .
<code>regenerate</code>	Logical; indicates if the Stan code should be regenerated with the current brms version. By default, <code>regenerate</code> will be <code>FALSE</code> unless required to be <code>TRUE</code> by other arguments.
<code>threads</code>	Controls whether the Stan code should be threaded. See threading for details.

- backend Controls the Stan backend. See [brm](#) for details.
 ... Further arguments passed to [stancode](#) if the Stan code is regenerated.

Value

Stan code for further processing.

stancode.default *Stan Code for brms Models*

Description

Generate Stan code for **brms** models

Usage

```
## Default S3 method:
stancode(
  object,
  data,
  family = gaussian(),
  prior = NULL,
  autocor = NULL,
  data2 = NULL,
  cov_ranef = NULL,
  sparse = NULL,
  sample_prior = "no",
  stanvars = NULL,
  stan_funs = NULL,
  knots = NULL,
  drop_unused_levels = TRUE,
  threads = getOption("brms.threads", NULL),
  normalize = getOption("brms.normalize", TRUE),
  save_model = NULL,
  ...
)
```

Arguments

- object An object of class [formula](#), [brmsformula](#), or [mvbrmsformula](#) (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in [brmsformula](#).
- data An object of class [data.frame](#) (or one that can be coerced to that class) containing data of all variables used in the model.

family	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see brmsfamily . By default, a linear gaussian model is applied. In multivariate models, family might also be a list of families.
prior	One or more <code>brmsprior</code> objects created by set_prior or related functions and combined using the <code>c</code> method or the <code>+</code> operator. See also default_prior for more help.
autocor	(Deprecated) An optional <code>cor_brms</code> object describing the correlation structure within the response variable (i.e., the 'autocorrelation'). See the documentation of <code>cor_brms</code> for a description of the available correlation structures. Defaults to <code>NULL</code> , corresponding to no correlations. In multivariate models, <code>autocor</code> might also be a list of autocorrelation structures. It is now recommend to specify autocorrelation terms directly within <code>formula</code> . See brmsformula for more details.
data2	A named list of objects containing data, which cannot be passed via argument <code>data</code> . Required for some objects used in autocorrelation structures to specify dependency structures as well as for within-group covariance matrices.
cov_ranef	(Deprecated) A list of matrices that are proportional to the (within) covariance structure of the group-level effects. The names of the matrices should correspond to columns in <code>data</code> that are used as grouping factors. All levels of the grouping factor should appear as rownames of the corresponding matrix. This argument can be used, among others to model pedigrees and phylogenetic effects. It is now recommended to specify those matrices in the formula interface using the <code>gr</code> and related functions. See <code>vignette("brms_phylogenetics")</code> for more details.
sparse	(Deprecated) Logical; indicates whether the population-level design matrices should be treated as sparse (defaults to <code>FALSE</code>). For design matrices with many zeros, this can considerably reduce required memory. Sampling speed is currently not improved or even slightly decreased. It is now recommended to use the <code>sparse</code> argument of brmsformula and related functions.
sample_prior	Indicate if draws from priors should be drawn additionally to the posterior draws. Options are "no" (the default), "yes", and "only". Among others, these draws can be used to calculate Bayes factors for point hypotheses via hypothesis . Please note that improper priors are not sampled, including the default improper priors used by <code>brm</code> . See set_prior on how to set (proper) priors. Please also note that prior draws for the overall intercept are not obtained by default for technical reasons. See brmsformula how to obtain prior draws for the intercept. If <code>sample_prior</code> is set to "only", draws are drawn solely from the priors ignoring the likelihood, which allows among others to generate draws from the prior predictive distribution. In this case, all parameters must have proper priors.
stanvars	An optional <code>stanvars</code> object generated by function stanvar to define additional variables for use in Stan's program blocks.
stan_funs	(Deprecated) An optional character string containing self-defined Stan functions, which will be included in the functions block of the generated Stan code. It is now recommended to use the <code>stanvars</code> argument for this purpose instead.

<code>knots</code>	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See gamm for more details.
<code>drop_unused_levels</code>	Should unused factors levels in the data be dropped? Defaults to TRUE.
<code>threads</code>	Number of threads to use in within-chain parallelization. For more control over the threading process, <code>threads</code> may also be a <code>brmsthreads</code> object created by threading . Within-chain parallelization is experimental! We recommend its use only if you are experienced with Stan's <code>reduce_sum</code> function and have a slow running model that cannot be sped up by any other means. Can be set globally for the current R session via the "brms.threads" option (see options).
<code>normalize</code>	Logical. Indicates whether normalization constants should be included in the Stan code (defaults to TRUE). Setting it to FALSE requires Stan version >= 2.25 to work. If FALSE, sampling efficiency may be increased but some post processing functions such as <code>bridge_sampler</code> will not be available. Can be controlled globally for the current R session via the 'brms.normalize' option.
<code>save_model</code>	Either NULL or a character string. In the latter case, the model's Stan code is saved via <code>cat</code> in a text file named after the string supplied in <code>save_model</code> .
<code>...</code>	Other arguments for internal usage only.

Value

A character string containing the fully commented **Stan** code to fit a **brms** model. It is of class `c("character", "brmsmodel")` to facilitate pretty printing.

Examples

```
stancode(rating ~ treat + period + carry + (1|subject),
         data = inhaler, family = "cumulative")

stancode(count ~ zAge + zBase * Trt + (1|patient),
         data = epilepsy, family = "poisson")
```

Description

`standata` is a generic function that can be used to generate data for Bayesian models to be passed to Stan. Its original use is within the **brms** package, but new methods for use with objects from other packages can be registered to the same generic.

Usage

```
standata(object, ...)
make_standata(formula, ...)
```

Arguments

- object A formula object whose class will determine which method will be used. A symbolic description of the model to be fitted.
- ... Further arguments passed to the specific method.
- formula Synonym of object for use in make_standata.

Details

See [standata.default](#) for the default method applied for **brms** models. You can view the available methods by typing `methods(standata)`. The `make_standata` function is an alias of `standata`.

Value

A named list of objects containing the required data to fit a Bayesian model with **Stan**.

See Also

[standata.default](#), [standata.brmsfit](#)

Examples

```
sdata1 <- standata(rating ~ treat + period + carry + (1|subject),  
                     data = inhaler, family = "cumulative")  
str(sdata1)
```

standata.brmsfit *Extract data passed to Stan from brmsfit objects*

Description

Extract all data that was used by Stan to fit a **brms** model.

Usage

```
## S3 method for class 'brmsfit'  
standata(  
  object,  
  newdata = NULL,  
  re_formula = NULL,  
  newdata2 = NULL,  
  new_objects = NULL,  
  incl_autocor = TRUE,  
  ...  
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>newdata</code>	An optional data.frame for which to evaluate predictions. If NULL (default), the original data of the model is used. NA values within factors (excluding grouping variables) are interpreted as if all dummy variables of this factor are zero. This allows, for instance, to make predictions of the grand mean when using sum coding. NA values within grouping variables are treated as a new level.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If NULL (default), include all group-level effects; if NA or <code>~0</code> , include no group-level effects.
<code>newdata2</code>	A named list of objects containing new data, which cannot be passed via argument <code>newdata</code> . Required for some objects used in autocorrelation structures, or <code>stanvars</code> .
<code>new_objects</code>	Deprecated alias of <code>newdata2</code> .
<code>incl_autocor</code>	A flag indicating if correlation structures originally specified via <code>autocor</code> should be included in the predictions. Defaults to TRUE.
<code>...</code>	More arguments passed to <code>standata.default</code> and <code>validate_newdata</code> .

Value

A named list containing the data passed to Stan.

`standata.default` *Data for brms Models*

Description

Generate data for **brms** models to be passed to **Stan**.

Usage

```
## Default S3 method:
standata(
  object,
  data,
  family = gaussian(),
  prior = NULL,
  autocor = NULL,
  data2 = NULL,
  cov_ranef = NULL,
  sample_prior = "no",
  stanvars = NULL,
  threads = getOption("brms.threads", NULL),
  knots = NULL,
  drop_unused_levels = TRUE,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>formula</code> , <code>brmsformula</code> , or <code>mvbrmsformula</code> (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in <code>brmsformula</code> .
<code>data</code>	An object of class <code>data.frame</code> (or one that can be coerced to that class) containing data of all variables used in the model.
<code>family</code>	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see <code>brmsfamily</code> . By default, a linear gaussian model is applied. In multivariate models, <code>family</code> might also be a list of families.
<code>prior</code>	One or more <code>brmsprior</code> objects created by <code>set_prior</code> or related functions and combined using the <code>c</code> method or the <code>+</code> operator. See also <code>default_prior</code> for more help.
<code>autocor</code>	(Deprecated) An optional <code>cor_brms</code> object describing the correlation structure within the response variable (i.e., the 'autocorrelation'). See the documentation of <code>cor_brms</code> for a description of the available correlation structures. Defaults to <code>NULL</code> , corresponding to no correlations. In multivariate models, <code>autocor</code> might also be a list of autocorrelation structures. It is now recommend to specify autocorrelation terms directly within <code>formula</code> . See <code>brmsformula</code> for more details.
<code>data2</code>	A named list of objects containing data, which cannot be passed via argument <code>data</code> . Required for some objects used in autocorrelation structures to specify dependency structures as well as for within-group covariance matrices.
<code>cov_ranef</code>	(Deprecated) A list of matrices that are proportional to the (within) covariance structure of the group-level effects. The names of the matrices should correspond to columns in <code>data</code> that are used as grouping factors. All levels of the grouping factor should appear as rownames of the corresponding matrix. This argument can be used, among others to model pedigrees and phylogenetic effects. It is now recommended to specify those matrices in the <code>formula</code> interface using the <code>gr</code> and related functions. See <code>vignette("brms_phylogenetics")</code> for more details.
<code>sample_prior</code>	Indicate if draws from priors should be drawn additionally to the posterior draws. Options are "no" (the default), "yes", and "only". Among others, these draws can be used to calculate Bayes factors for point hypotheses via <code>hypothesis</code> . Please note that improper priors are not sampled, including the default improper priors used by <code>brm</code> . See <code>set_prior</code> on how to set (proper) priors. Please also note that prior draws for the overall intercept are not obtained by default for technical reasons. See <code>brmsformula</code> how to obtain prior draws for the intercept. If <code>sample_prior</code> is set to "only", draws are drawn solely from the priors ignoring the likelihood, which allows among others to generate draws from the prior predictive distribution. In this case, all parameters must have proper priors.
<code>stanvars</code>	An optional <code>stanvars</code> object generated by function <code>stanvar</code> to define additional variables for use in Stan's program blocks.

<code>threads</code>	Number of threads to use in within-chain parallelization. For more control over the threading process, <code>threads</code> may also be a <code>brmstthreads</code> object created by threading . Within-chain parallelization is experimental! We recommend its use only if you are experienced with Stan's <code>reduce_sum</code> function and have a slow running model that cannot be sped up by any other means. Can be set globally for the current R session via the " <code>brms.threads</code> " option (see options).
<code>knots</code>	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See gamm for more details.
<code>drop_unused_levels</code>	Should unused factors levels in the data be dropped? Defaults to TRUE.
<code>...</code>	Other arguments for internal use.

Value

A named list of objects containing the required data to fit a **brms** model with **Stan**.

Examples

```
sdata1 <- standata(rating ~ treat + period + carry + (1|subject),
                     data = inhaler, family = "cumulative")
str(sdata1)

sdata2 <- standata(count ~ zAge + zBase * Trt + (1|patient),
                     data = epilepsy, family = "poisson")
str(sdata2)
```

Description

Prepare user-defined variables to be passed to one of Stan's program blocks. This is primarily useful for defining more complex priors, for refitting models without recompilation despite changing priors, or for defining custom Stan functions.

Usage

```
stanvar(
  x = NULL,
  name = NULL,
  scode = NULL,
  block = "data",
  position = "start",
  pll_args = NULL
)
```

Arguments

x	An R object containing data to be passed to Stan. Only required if block = 'data' and ignored otherwise.
name	Optional character string providing the desired variable name of the object in x. If NULL (the default) the variable name is directly inferred from x.
scode	Line of Stan code to define the variable in Stan language. If block = 'data', the Stan code is inferred based on the class of x by default.
block	Name of one of Stan's program blocks in which the variable should be defined. Can be 'data', 'tdata' (transformed data), 'parameters', 'tparameters' (transformed parameters), 'model', 'likelihood' (part of the model block where the likelihood is given), 'genquant' (generated quantities) or 'functions'.
position	Name of the position within the block where the Stan code should be placed. Currently allowed are 'start' (the default) and 'end' of the block.
pll_args	Optional Stan code to be put into the header of partial_log_lik functions. This ensures that the variables specified in scode can be used in the likelihood even when within-chain parallelization is activated via threading .

Details

The `stanvar` function is not vectorized. Instead, multiple `stanvars` objects can be added together via `+` (see Examples).

Special attention is necessary when using `stanvars` to inject code into the 'likelihood' block while having [threading](#) activated. In this case, your custom Stan code may need adjustments to ensure correct observation indexing. Please investigate the generated Stan code via `stancode` to see which adjustments are necessary in your case.

Value

An object of class `stanvars`.

Examples

```

bprior <- prior(normal(mean_intercept, 10), class = "Intercept")
stanvars <- stanvar(5, name = "mean_intercept")
stancode(count ~ Trt, epilepsy, prior = bprior,
          stanvars = stanvars)

# define a multi-normal prior with known covariance matrix
bprior <- prior(multi_normal(M, V), class = "b")
stanvars <- stanvar(rep(0, 2), "M", scode = "  vector[K] M;") +
  stanvar(diag(2), "V", scode = "  matrix[K, K] V;")
stancode(count ~ Trt + zBase, epilepsy,
          prior = bprior, stanvars = stanvars)

# define a hierarchical prior on the regression coefficients
bprior <- set_prior("normal(0, tau)", class = "b") +
  set_prior("target += normal_lpdf(tau | 0, 10)", check = FALSE)
stanvars <- stanvar(scode = "real<lower=0> tau;",

```

```

            block = "parameters")
stancode(count ~ Trt + zBase, epilepsy,
         prior = bprior, stanvars = stanvars)

# ensure that 'tau' is passed to the likelihood of a threaded model
# not necessary for this example but may be necessary in other cases
stanvars <- stanvar(scode = "real<lower=0> tau;",
                     block = "parameters", pll_args = "real tau")
stancode(count ~ Trt + zBase, epilepsy,
         stanvars = stanvars, threads = threading(2))

```

StudentT*The Student-t Distribution***Description**

Density, distribution function, quantile function and random generation for the Student-t distribution with location `mu`, scale `sigma`, and degrees of freedom `df`.

Usage

```

dstudent_t(x, df, mu = 0, sigma = 1, log = FALSE)
pstUDENT_t(q, df, mu = 0, sigma = 1, lower.tail = TRUE, log.p = FALSE)
qstudent_t(p, df, mu = 0, sigma = 1, lower.tail = TRUE, log.p = FALSE)
rstUDENT_t(n, df, mu = 0, sigma = 1)

```

Arguments

<code>x</code>	Vector of quantiles.
<code>df</code>	Vector of degrees of freedom.
<code>mu</code>	Vector of location values.
<code>sigma</code>	Vector of scale values.
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>q</code>	Vector of quantiles.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If TRUE, values are returned on the log scale.
<code>p</code>	Vector of probabilities.
<code>n</code>	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

See Also[TDist](#)

summary.brmsfit	<i>Create a summary of a fitted model represented by a brmsfit object</i>
-----------------	---

Description

Create a summary of a fitted model represented by a `brmsfit` object

Usage

```
## S3 method for class 'brmsfit'
summary(
  object,
  priors = FALSE,
  prob = 0.95,
  robust = FALSE,
  mc_se = FALSE,
  ...
)
```

Arguments

<code>object</code>	An object of class <code>brmsfit</code> .
<code>priors</code>	Logical; Indicating if priors should be included in the summary. Default is FALSE.
<code>prob</code>	A value between 0 and 1 indicating the desired probability to be covered by the uncertainty intervals. The default is 0.95.
<code>robust</code>	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead.
<code>mc_se</code>	Logical; Indicating if the uncertainty in Estimate caused by the MCMC sampling should be shown in the summary. Defaults to FALSE.
<code>...</code>	Other potential arguments

Details

The convergence diagnostics `Rhat`, `Bulk_ESS`, and `Tail_ESS` are described in detail in Vehtari et al. (2020).

References

Aki Vehtari, Andrew Gelman, Daniel Simpson, Bob Carpenter, and Paul-Christian Bürkner (2020). Rank-normalization, folding, and localization: An improved R-hat for assessing convergence of MCMC. **Bayesian Analysis**. 1–28. doi:10.1214/20-BA1221

theme_black*(Deprecated) Black Theme for ggplot2 Graphics*

Description

A black theme for ggplot graphics inspired by a blog post of Jon Lefcheck (<https://jonlefcheck.net/2013/03/11/black-theme-for-ggplot2-2/>).

Usage

```
theme_black(base_size = 12, base_family = "")
```

Arguments

base_size	base font size
base_family	base font family

Details

When using theme_black in plots powered by the **bayesplot** package such as pp_check or stanplot, I recommend using the "viridisC" color scheme (see examples).

Value

A theme object used in **ggplot2** graphics.

Examples

```
## Not run:
# change default ggplot theme
ggplot2::theme_set(theme_black())

# change default bayesplot color scheme
bayesplot::color_scheme_set("viridisC")

# fit a simple model
fit <- brm(count ~ zAge + zBase * Trt + (1|patient),
            data = epilepsy, family = poisson(), chains = 2)
summary(fit)

# create various plots
plot(marginal_effects(fit), ask = FALSE)
pp_check(fit)
mcmc_hex(fit, type = "hex", variable = c("b_Intercept", "b_Trt1"))

## End(Not run)
```

theme_default	<i>Default bayesplot Theme for ggplot2 Graphics</i>
---------------	---

Description

This theme is imported from the **bayesplot** package. See [theme_default](#) for a complete documentation.

Arguments

base_size	base font size
base_family	base font family

Value

A theme object used in **ggplot2** graphics.

threading	<i>Threading in Stan</i>
-----------	--------------------------

Description

Use threads for within-chain parallelization in **Stan** via the **brms** interface. Within-chain parallelization is experimental! We recommend its use only if you are experienced with Stan's `reduce_sum` function and have a slow running model that cannot be sped up by any other means.

Usage

```
threading(threads = NULL, grainsize = NULL, static = FALSE, force = FALSE)
```

Arguments

threads	Number of threads to use in within-chain parallelization.
grainsize	Number of observations evaluated together in one chunk on one of the CPUs used for threading. If <code>NULL</code> (the default), <code>grainsize</code> is currently chosen as <code>max(100, N / (2 * threads))</code> , where <code>N</code> is the number of observations in the data. This default is experimental and may change in the future without prior notice.
static	Logical. Apply the static (non-adaptive) version of <code>reduce_sum?</code> Defaults to <code>FALSE</code> . Setting it to <code>TRUE</code> is required to achieve exact reproducibility of the model results (if the random seed is set as well).
force	Logical. Defaults to <code>FALSE</code> . If <code>TRUE</code> , this will force the Stan model to compile with threading enabled without altering the Stan code generated by <code>brms</code> . This can be useful if your own custom Stan functions use threading internally.

Details

The adaptive scheduling procedure used by `reduce_sum` will prevent the results to be exactly reproducible even if you set the random seed. If you need exact reproducibility, you have to set argument `static = TRUE` which may reduce efficiency a bit.

To ensure that chunks (whose size is defined by `grainsize`) require roughly the same amount of computing time, we recommend storing observations in random order in the data. At least, please avoid sorting observations after the response values. This is because the latter often cause variations in the computing time of the pointwise log-likelihood, which makes up a big part of the parallelized code.

Value

A `brmstthreads` object which can be passed to the `threads` argument of `brm` and related functions.

Examples

```
## Not run:
# this model just serves as an illustration
# threading may not actually speed things up here
fit <- brm(count ~ zAge + zBase * Trt + (1|patient),
            data = epilepsy, family = negbinomial(),
            chains = 1, threads = threading(2, grainsize = 100),
            backend = "cmdstanr")
summary(fit)

## End(Not run)
```

unstr

Set up UNSTR correlation structures

Description

Set up an unstructured (UNSTR) correlation term in `brms`. The function does not evaluate its arguments – it exists purely to help set up a model with UNSTR terms.

Usage

```
unstr(time, gr)
```

Arguments

<code>time</code>	An optional time variable specifying the time ordering of the observations. By default, the existing order of the observations in the data is used.
<code>gr</code>	An optional grouping variable. If specified, the correlation structure is assumed to apply only to observations within the same grouping level.

Value

An object of class 'unstr_term', which is a list of arguments to be interpreted by the formula parsing functions of **brms**.

See Also

[autocor-terms](#)

Examples

```
## Not run:  
# add an unstructured correlation matrix for visits within the same patient  
fit <- brm(count ~ Trt + unstr(visit, patient), data = epilepsy)  
summary(fit)  
  
## End(Not run)
```

update.brmsfit *Update brms models*

Description

This method allows to update an existing **brmsfit** object.

Usage

```
## S3 method for class 'brmsfit'  
update(object, formula., newdata = NULL, recompile = NULL, ...)
```

Arguments

object	An object of class brmsfit .
formula.	Changes to the formula; for details see update.formula and brmsformula .
newdata	Optional <code>data.frame</code> to update the model with new data. Data-dependent default priors will not be updated automatically.
recompile	Logical, indicating whether the Stan model should be recompiled. If <code>NULL</code> (the default), <code>update</code> tries to figure out internally, if recompilation is necessary. Setting it to <code>FALSE</code> will cause all Stan code changing arguments to be ignored.
...	Other arguments passed to brm .

Details

When updating a **brmsfit** created with the **cmdstanr** backend in a different R session, a recompilation will be triggered because by default, **cmdstanr** writes the model executable to a temporary directory. To avoid that, set option "`cmdstanr_write_stan_file_dir`" to a nontemporary path of your choice before creating the original **brmsfit** (see section 'Examples' below).

Examples

```

## Not run:
fit1 <- brm(time | cens(censored) ~ age * sex + disease + (1|patient),
             data = kidney, family = gaussian("log"))
summary(fit1)

## remove effects of 'disease'
fit2 <- update(fit1, formula. = ~ . - disease)
summary(fit2)

## remove the group specific term of 'patient' and
## change the data (just take a subset in this example)
fit3 <- update(fit1, formula. = ~ . - (1|patient),
               newdata = kidney[1:38, ])
summary(fit3)

## use another family and add population-level priors
fit4 <- update(fit1, family = weibull(), init = "0",
               prior = set_prior("normal(0,5)"))
summary(fit4)

## to avoid a recompilation when updating a 'cmdstanr'-backend fit in a fresh
## R session, set option 'cmdstanr_write_stan_file_dir' before creating the
## initial 'brmsfit'
## CAUTION: the following code creates some files in the current working
## directory: two 'model_<hash>.stan' files, one 'model_<hash>(.exe)'
## executable, and one 'fit_cmdstanr_<some_number>.rds' file
set.seed(7)
fname <- paste0("fit_cmdstanr_", sample.int(.Machine$integer.max, 1))
options(cmdstanr_write_stan_file_dir = getwd())
fit_cmdstanr <- brm(rate ~ conc + state,
                     data = Puromycin,
                     backend = "cmdstanr",
                     file = fname)
# now restart the R session and run the following (after attaching 'brms')
set.seed(7)
fname <- paste0("fit_cmdstanr_", sample.int(.Machine$integer.max, 1))
fit_cmdstanr <- brm(rate ~ conc + state,
                     data = Puromycin,
                     backend = "cmdstanr",
                     file = fname)
upd_cmdstanr <- update(fit_cmdstanr,
                       formula. = rate ~ conc)

## End(Not run)

```

Description

This method allows to update an existing `brmsfit_multiple` object.

Usage

```
## S3 method for class 'brmsfit_multiple'
update(object, formula., newdata = NULL, data2 = NULL, ...)
```

Arguments

<code>object</code>	An object of class <code>brmsfit_multiple</code> .
<code>formula.</code>	Changes to the formula; for details see update.formula and brmsformula .
<code>newdata</code>	List of <code>data.frames</code> to update the model with new data. Currently required even if the original data should be used.
<code>data2</code>	A <i>list</i> of named lists each of which will be used to fit a separate model. Each of the named lists contains objects representing data which cannot be passed via argument <code>data</code> (see brm for examples). The length of the outer list should match the length of the list passed to the <code>data</code> argument. Currently required even if the original data should be used.
<code>...</code>	Other arguments passed to update.brmsfit and brm_multiple .

Examples

```
## Not run:
library(mice)
imp <- mice(nhanes2)

# initially fit the model
fit_imp1 <- brm_multiple(bmi ~ age + hyp + chl, data = imp, chains = 1)
summary(fit_imp1)

# update the model using fewer predictors
fit_imp2 <- update(fit_imp1, formula. = . ~ hyp + chl, newdata = imp)
summary(fit_imp2)

## End(Not run)
```

Description

Update additions terms used in formulas of **brms**. See [addition-terms](#) for details.

Usage

```
update_adterms(formula, adform, action = c("update", "replace"))
```

Arguments

- formula Two-sided formula to be updated.
- adform One-sided formula containing addition terms to update formula with.
- action Indicates what should happen to the existing addition terms in formula. If "update" (the default), old addition terms that have no corresponding term in adform will be kept. If "replace", all old addition terms will be removed.

Value

An object of class `formula`.

Examples

```
form <- y | trials(size) ~ x
update_adterms(form, ~ trials(10))
update_adterms(form, ~ weights(w))
update_adterms(form, ~ weights(w), action = "replace")
update_adterms(y ~ x, ~ trials(10))
```

`validate_newdata` *Validate New Data*

Description

Validate new data passed to post-processing methods of `brms`. Unless you are a package developer, you will rarely need to call `validate_newdata` directly.

Usage

```
validate_newdata(
  newdata,
  object,
  re_formula = NULL,
  allow_new_levels = FALSE,
  newdata2 = NULL,
  resp = NULL,
  check_response = TRUE,
  incl_autocor = TRUE,
  group_vars = NULL,
  req_vars = NULL,
  ...
)
```

Arguments

<code>newdata</code>	A <code>data.frame</code> containing new data to be validated.
<code>object</code>	A <code>brmsfit</code> object.
<code>re_formula</code>	formula containing group-level effects to be considered in the prediction. If <code>NULL</code> (default), include all group-level effects; if <code>NA</code> or <code>~0</code> , include no group-level effects.
<code>allow_new_levels</code>	A flag indicating if new levels of group-level effects are allowed (defaults to <code>FALSE</code>). Only relevant if <code>newdata</code> is provided.
<code>newdata2</code>	A named list of objects containing new data, which cannot be passed via argument <code>newdata</code> . Required for some objects used in autocorrelation structures, or <code>stanvars</code> .
<code>resp</code>	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
<code>check_response</code>	Logical; Indicates if response variables should be checked as well. Defaults to <code>TRUE</code> .
<code>incl_autocor</code>	A flag indicating if correlation structures originally specified via <code>autocor</code> should be included in the predictions. Defaults to <code>TRUE</code> .
<code>group_vars</code>	Optional names of grouping variables to be validated. Defaults to all grouping variables in the model.
<code>req_vars</code>	Optional names of variables required in <code>newdata</code> . If <code>NULL</code> (the default), all variables in the original data are required (unless ignored for some other reason).
<code>...</code>	Currently ignored.

Value

A validated '`data.frame`' based on `newdata`.

`validate_prior`

Validate Prior for brms Models

Description

Validate priors supplied by the user. Return a complete set of priors for the given model, including default priors.

Usage

```
validate_prior(
  prior,
  formula,
  data,
  family = gaussian(),
  sample_prior = "no",
```

```

  data2 = NULL,
  knots = NULL,
  drop_unused_levels = TRUE,
  ...
)

```

Arguments

<code>prior</code>	One or more <code>brmsprior</code> objects created by <code>set_prior</code> or related functions and combined using the <code>c</code> method or the <code>+</code> operator. See also <code>default_prior</code> for more help.
<code>formula</code>	An object of class <code>formula</code> , <code>brmsformula</code> , or <code>mvbrmsformula</code> (or one that can be coerced to that classes): A symbolic description of the model to be fitted. The details of model specification are explained in <code>brmsformula</code> .
<code>data</code>	An object of class <code>data.frame</code> (or one that can be coerced to that class) containing data of all variables used in the model.
<code>family</code>	A description of the response distribution and link function to be used in the model. This can be a family function, a call to a family function or a character string naming the family. Every family function has a <code>link</code> argument allowing to specify the link function to be applied on the response variable. If not specified, default links are used. For details of supported families see <code>brmsfamily</code> . By default, a linear gaussian model is applied. In multivariate models, <code>family</code> might also be a list of families.
<code>sample_prior</code>	Indicate if draws from priors should be drawn additionally to the posterior draws. Options are "no" (the default), "yes", and "only". Among others, these draws can be used to calculate Bayes factors for point hypotheses via <code>hypothesis</code> . Please note that improper priors are not sampled, including the default improper priors used by <code>brm</code> . See <code>set_prior</code> on how to set (proper) priors. Please also note that prior draws for the overall intercept are not obtained by default for technical reasons. See <code>brmsformula</code> how to obtain prior draws for the intercept. If <code>sample_prior</code> is set to "only", draws are drawn solely from the priors ignoring the likelihood, which allows among others to generate draws from the prior predictive distribution. In this case, all parameters must have proper priors.
<code>data2</code>	A named list of objects containing data, which cannot be passed via argument <code>data</code> . Required for some objects used in autocorrelation structures to specify dependency structures as well as for within-group covariance matrices.
<code>knots</code>	Optional list containing user specified knot values to be used for basis construction of smoothing terms. See <code>gamm</code> for more details.
<code>drop_unused_levels</code>	Should unused factors levels in the data be dropped? Defaults to <code>TRUE</code> .
<code>...</code>	Other arguments for internal usage only.

Value

An object of class `brmsprior`.

See Also

[default_prior](#), [set_prior](#).

Examples

```
prior1 <- prior(normal(0,10), class = b) +
  prior(cauchy(0,2), class = sd)
validate_prior(prior1, count ~ zAge + zBase * Trt + (1|patient),
  data = epilepsy, family = poisson())
```

VarCorr.brmsfit

Extract Variance and Correlation Components

Description

This function calculates the estimated standard deviations, correlations and covariances of the group-level terms in a multilevel model of class `brmsfit`. For linear models, the residual standard deviations, correlations and covariances are also returned.

Usage

```
## S3 method for class 'brmsfit'
VarCorr(
  x,
  sigma = 1,
  summary = TRUE,
  robust = FALSE,
  probs = c(0.025, 0.975),
  ...
)
```

Arguments

<code>x</code>	An object of class <code>brmsfit</code> .
<code>sigma</code>	Ignored (included for compatibility with VarCorr).
<code>summary</code>	Should summary statistics be returned instead of the raw values? Default is TRUE.
<code>robust</code>	If FALSE (the default) the mean is used as the measure of central tendency and the standard deviation as the measure of variability. If TRUE, the median and the median absolute deviation (MAD) are applied instead. Only used if <code>summary</code> is TRUE.
<code>probs</code>	The percentiles to be computed by the <code>quantile</code> function. Only used if <code>summary</code> is TRUE.
<code>...</code>	Currently ignored.

Value

A list of lists (one per grouping factor), each with three elements: a matrix containing the standard deviations, an array containing the correlation matrix, and an array containing the covariance matrix with variances on the diagonal.

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt + (1+Trt|visit),
            data = epilepsy, family = gaussian(), chains = 2)
VarCorr(fit)

## End(Not run)
```

vcov.brmsfit

*Covariance and Correlation Matrix of Population-Level Effects***Description**

Get a point estimate of the covariance or correlation matrix of population-level parameters

Usage

```
## S3 method for class 'brmsfit'
vcov(object, correlation = FALSE, pars = NULL, ...)
```

Arguments

- | | |
|-------------|--|
| object | An object of class <code>brmsfit</code> . |
| correlation | Logical; if FALSE (the default), compute the covariance matrix, if TRUE, compute the correlation matrix. |
| pars | Optional names of coefficients to extract. By default, all coefficients are extracted. |
| ... | Currently ignored. |

Details

Estimates are obtained by calculating the maximum likelihood covariances (correlations) of the posterior draws.

Value

covariance or correlation matrix of population-level parameters

Examples

```
## Not run:
fit <- brm(count ~ zAge + zBase * Trt + (1+Trt|visit),
            data = epilepsy, family = gaussian(), chains = 2)
vcov(fit)

## End(Not run)
```

Description

Density, distribution function, and random generation for the von Mises distribution with location μ , and precision κ .

Usage

```
dvon_mises(x, mu, kappa, log = FALSE)

pvon_mises(q, mu, kappa, lower.tail = TRUE, log.p = FALSE, acc = 1e-20)

rvon_mises(n, mu, kappa)
```

Arguments

<code>x, q</code>	Vector of quantiles between $-\pi$ and π .
<code>mu</code>	Vector of location values.
<code>kappa</code>	Vector of precision values.
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>lower.tail</code>	Logical; If TRUE (default), return $P(X \leq x)$. Else, return $P(X > x)$.
<code>log.p</code>	Logical; If TRUE, values are returned on the log scale.
<code>acc</code>	Accuracy of numerical approximations.
<code>n</code>	Number of draws to sample from the distribution.

Details

See `vignette("brms_families")` for details on the parameterization.

waic.brmsfit*Widely Applicable Information Criterion (WAIC)*

Description

Compute the widely applicable information criterion (WAIC) based on the posterior likelihood using the [loo](#) package. For more details see [waic](#).

Usage

```
## S3 method for class 'brmsfit'
waic(
  x,
  ...,
  compare = TRUE,
  resp = NULL,
  pointwise = FALSE,
  model_names = NULL
)
```

Arguments

x	A <code>brmsfit</code> object.
...	More <code>brmsfit</code> objects or further arguments passed to the underlying post-processing functions. In particular, see prepare_predictions for further supported arguments.
compare	A flag indicating if the information criteria of the models should be compared to each other via loo_compare .
resp	Optional names of response variables. If specified, predictions are performed only for the specified response variables.
pointwise	A flag indicating whether to compute the full log-likelihood matrix at once or separately for each observation. The latter approach is usually considerably slower but requires much less working memory. Accordingly, if one runs into memory issues, <code>pointwise = TRUE</code> is the way to go.
model_names	If <code>NULL</code> (the default) will use model names derived from deparsing the call. Otherwise will use the passed values as model names.

Details

See [loo_compare](#) for details on model comparisons. For `brmsfit` objects, `WAIC` is an alias of `waic`. Use method [add_criterion](#) to store information criteria in the fitted model object for later usage.

Value

If just one object is provided, an object of class `loo`. If multiple objects are provided, an object of class `loolist`.

References

- Vehtari, A., Gelman, A., & Gabry J. (2016). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. In *Statistics and Computing*, doi:10.1007/s11222-016-9696-4. arXiv preprint arXiv:1507.04544.
- Gelman, A., Hwang, J., & Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing*, 24, 997-1016.
- Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *The Journal of Machine Learning Research*, 11, 3571-3594.

Examples

```
## Not run:
# model with population-level effects only
fit1 <- brm(rating ~ treat + period + carry,
             data = inhaler)
(waic1 <- waic(fit1))

# model with an additional varying intercept for subjects
fit2 <- brm(rating ~ treat + period + carry + (1|subject),
             data = inhaler)
(waic2 <- waic(fit2))

# compare both models
loo_compare(waic1, waic2)

## End(Not run)
```

Description

Density function and random generation for the Wiener diffusion model distribution with boundary separation alpha, non-decision time tau, bias beta and drift rate delta.

Usage

```
dwiener(
  x,
  alpha,
  tau,
  beta,
  delta,
  resp = 1,
  log = FALSE,
```

```

backend = getOption("wiener_backend", "Rwiener")
)

rwiener(
  n,
  alpha,
  tau,
  beta,
  delta,
  types = c("q", "resp"),
  backend = getOption("wiener_backend", "Rwiener")
)

```

Arguments

<code>x</code>	Vector of quantiles.
<code>alpha</code>	Boundary separation parameter.
<code>tau</code>	Non-decision time parameter.
<code>beta</code>	Bias parameter.
<code>delta</code>	Drift rate parameter.
<code>resp</code>	Response: "upper" or "lower". If no character vector, it is coerced to logical where TRUE indicates "upper" and FALSE indicates "lower".
<code>log</code>	Logical; If TRUE, values are returned on the log scale.
<code>backend</code>	Name of the package to use as backend for the computations. Either "Rwiener" (the default) or "rtdists". Can be set globally for the current R session via the "wiener_backend" option (see options).
<code>n</code>	Number of draws to sample from the distribution.
<code>types</code>	Which types of responses to return? By default, return both the response times "q" and the dichotomous responses "resp". If either "q" or "resp", return only one of the two types.

Details

These are wrappers around functions of the **RWiener** or **rtdists** package (depending on the chosen backend). See `vignette("brms_families")` for details on the parameterization.

See Also

[wienerdist](#), [Diffusion](#)

ZeroInflated*Zero-Inflated Distributions*

Description

Density and distribution functions for zero-inflated distributions.

Usage

```
dzero_inflated_poisson(x, lambda, zi, log = FALSE)

pzero_inflated_poisson(q, lambda, zi, lower.tail = TRUE, log.p = FALSE)

dzero_inflated_negbinomial(x, mu, shape, zi, log = FALSE)

pzero_inflated_negbinomial(q, mu, shape, zi, lower.tail = TRUE, log.p = FALSE)

dzero_inflated_binomial(x, size, prob, zi, log = FALSE)

pzero_inflated_binomial(q, size, prob, zi, lower.tail = TRUE, log.p = FALSE)

dzero_inflated_beta_binomial(x, size, mu, phi, zi, log = FALSE)

pzero_inflated_beta_binomial(
  q,
  size,
  mu,
  phi,
  zi,
  lower.tail = TRUE,
  log.p = FALSE
)
dzero_inflated_beta(x, shape1, shape2, zi, log = FALSE)

pzero_inflated_beta(q, shape1, shape2, zi, lower.tail = TRUE, log.p = FALSE)
```

Arguments

x	Vector of quantiles.
zi	zero-inflation probability
log	Logical; If TRUE, values are returned on the log scale.
q	Vector of quantiles.
lower.tail	Logical; If TRUE (default), return P(X <= x). Else, return P(X > x) .
log.p	Logical; If TRUE, values are returned on the log scale.

mu, lambda	location parameter
shape, shape1, shape2	shape parameter
size	number of trials
prob	probability of success on each trial
phi	precision parameter

Details

The density of a zero-inflated distribution can be specified as follows. If $x = 0$ set $f(x) = \theta + (1 - \theta) * g(0)$. Else set $f(x) = (1 - \theta) * g(x)$, where $g(x)$ is the density of the non-zero-inflated part.

Index

* datasets
 epilepsy, 96
 inhaler, 118
 kidney, 130
 loss, 146

acat (brmsfamily), 34
acformula, 19
acformula (brmsformula-helpers), 51
add_criterion, 11, 12, 64, 137, 138, 140, 206, 248
add_ic, 64
add_ic (add_loo), 12
add_ic<- (add_loo), 12
add_loo, 12
add_rstan_model, 13
add_waic (add_loo), 12
addition-terms, 8
and (draws-index-brms), 93
ar, 13, 15, 19, 73, 148
arma, 14, 14, 19, 74, 148
as.array.brmsfit
 (as.data.frame.brmsfit), 16
as.brmsprior, 15
as.data.frame, 16, 178
as.data.frame.brmsfit, 16
as.matrix.brmsfit, 172
as.matrix.brmsfit
 (as.data.frame.brmsfit), 16
as.mcmc (as.mcmc.brmsfit), 17
as.mcmc.brmsfit, 17
as_draws, 17, 178
as_draws (draws-brms), 92
as_draws_*, 17, 177
as_draws_array (draws-brms), 92
as_draws_df (draws-brms), 92
as_draws_list (draws-brms), 92
as_draws_matrix (draws-brms), 92
as_draws_rvars (draws-brms), 92
asym_laplace (brmsfamily), 34

AsymLaplace, 18
attributes, 200
autocor (autocor.brmsfit), 20
autocor-terms, 19
autocor.brmsfit, 20

bayes_factor, 7, 24, 117, 182
bayes_factor (bayes_factor.brmsfit), 20
bayes_factor.brmsfit, 20
bayes_R2 (bayes_R2.brmsfit), 21
bayes_R2.brmsfit, 21
bayesplot, 7, 186
bernoulli (brmsfamily), 34
Beta (brmsfamily), 34
beta_binomial (brmsfamily), 34
BetaBinomial, 23
bf (brmsformula), 41
bf-helpers (brmsformula-helpers), 51
bridge_sampler, 20, 21, 29, 182, 228
bridge_sampler
 (bridge_sampler.brmsfit), 24
bridge_sampler.brmsfit, 24
bridge_sampler.stanfit, 24
bridgesampling::bayes_factor, 21
bridgesampling::bridge_sampler, 24
bridgesampling::post_prob, 182
brm, 7, 8, 10, 25, 40–42, 55, 57, 59, 80, 113, 125, 137, 141, 154, 204, 208, 215, 219, 220, 226, 239, 241
brm_multiple, 44, 56, 241
brms, 32, 40, 41
brms (brms-package), 6
brms-package, 6
brmsfamily, 8, 26, 30, 32, 34, 41, 48, 57, 83, 85, 88, 154, 227, 231, 244
brmsfit, 8, 32
brmsfit(brmsfit-class), 40
brmsfit-class, 40
brmsfit_needs_refit, 30, 58

brmsformula, 7, 8, 10, 19, 26, 27, 30, 32, 39–41, 41, 51, 52, 55–57, 83–85, 88, 110, 111, 151, 152, 154, 156, 159, 162, 163, 213, 217, 226, 227, 231, 239, 241, 244
brmsformula-helpers, 51
brmshypothesis, 53, 117
brmsprior, 40
brmsprior(set_prior), 216
brmsprior-class(set_prior), 216
brmsterms, 54, 123, 124
car, 19, 60, 76
cat, 30, 228
cat(addition-terms), 8
categorical(brmsfamily), 34
cbind, 157, 162
cens(addition-terms), 8
chains,(draws-index-brms), 93
coef.brmsfit, 61, 116
combine_models, 58, 63
compare_ic, 63
conditional_effects, 7, 71, 148, 149, 212
conditional_effects
 (**conditional_effects.brmsfit**), 64
conditional_effects.brmsfit, 64
conditional_smooths
 (**conditional_smooths.brmsfit**), 69
conditional_smooths.brmsfit, 69
constant, 72
control_params, 73
cor_ar, 73, 75, 76
cor_arma, 74, 74, 76, 79
cor_arma-class(cor_arma), 74
cor_brms, 19, 27, 41, 42, 57, 75, 88, 227, 231
cor_brms-class(cor_brms), 75
cor_car, 76, 76
cor_cosy, 77
cor_cosy-class(cor_cosy), 77
cor_errorvars(cor_sar), 80
cor_fixed, 76, 78
cor_icar(cor_car), 76
cor_lagsar(cor_sar), 80
cor_ma, 75, 76, 79
cor_sar, 76, 80
cosy, 19, 77, 81
cov_fixed(cor_fixed), 78
cox(brmsfamily), 34
cratio(brmsfamily), 34
create_priorsense_data.brmsfit, 82
cs, 83
cse(cs), 83
cumulative(brmsfamily), 34
custom_family, 34, 83
customfamily, 40
customfamily(custom_family), 83
cv_varsel, 106
dasym_laplace(AsymLaplace), 18
dbeta_binomial(BetaBinomial), 23
ddirichlet(Dirichlet), 91
dec(addition-terms), 8
default_prior, 27, 31, 57, 86, 89, 217, 221, 227, 231, 244, 245
default_prior.default, 87, 87
density, 89
density_ratio, 89
dexgaussian(ExGaussian), 97
dfrechet(Frechet), 103
dgen_extreme_value(GenExtremeValue), 104
dhurdle_gamma(Hurdle), 114
dhurdle_lognormal(Hurdle), 114
dhurdle_negbinomial(Hurdle), 114
dhurdle_poisson(Hurdle), 114
diagnostic-quantities, 90
Diffusion, 250
dinv_gaussian(InvGaussian), 120
Dirichlet, 91
dirichlet(brmsfamily), 34
dirichlet_multinomial(brmsfamily), 34
dlogistic_normal(LogisticNormal), 133
dmulti_normal(MultiNormal), 160
dmulti_student_t(MultiStudentT), 161
draws, 93
draws-brms, 92
draws-index-brms, 93
draws.(draws-index-brms), 93
dshifted_lnorm(Shifted_Lognormal), 222
dskew_normal(SkewNormal), 223
dstudent_t(StudentT), 234
dvon_mises(VonMises), 247
dwiener(Wiener), 249
dzero_inflated_beta(ZeroInflated), 251
dzero_inflated_beta_binomial(ZeroInflated), 251

dzero_inflated_binomial (ZeroInflated),
 251
dzero_inflated_negbinomial
 (ZeroInflated), 251
dzero_inflated_poisson (ZeroInflated),
 251

E_loo, 142
emm_basis.brmsfit
 (emmeans-brms-helpers), 94
emmeans-brms-helpers, 94
empty_prior (set_prior), 216
environment, 85
epilepsy, 96
ExGaussian, 97
exgaussian (brmsfamily), 34
exponential (brmsfamily), 34
expose_functions
 (expose_functions.brmsfit), 98
expose_functions.brmsfit, 98
expose_stan_functions, 98
expp1, 98
extend_family, 107
extract_draws
 (prepare_predictions.brmsfit),
 194

facet_wrap, 68
family, 34, 39, 40
family.brmsfit, 99
fcor, 19, 78, 99
fitted.brmsfit, 100
fixef(fixef.brmsfit), 102
fixef.brmsfit, 62, 102
formula, 26, 55, 56, 88, 226, 231, 244
Frechet, 103
frechet (brmsfamily), 34
future, 29, 126, 206

gam, 43
gamm, 27, 43, 57, 88, 219, 228, 232, 244
Gamma, 39
gen_extreme_value (brmsfamily), 34
GenExtremeValue, 104
geom_contour, 68
geom_errorbar, 68
geom_jitter, 67, 68
geom_point, 67
geom_raster, 68

geom_rug, 67, 68
geom_smooth, 67, 68
geometric (brmsfamily), 34
get_dpar, 105
get_prior (default_prior), 86
get_refmodel, 107
get_refmodel.brmsfit, 106
ggplot, 68, 150
ggtheme, 54, 68, 167
gp, 43, 108, 219
gr, 27, 43, 57, 110, 227, 231, 231
gtable, 168

horseshoe, 112, 131, 220
Hurdle, 114
hurdle_cumulative (brmsfamily), 34
hurdle_gamma (brmsfamily), 34
hurdle_lognormal (brmsfamily), 34
hurdle_negbinomial (brmsfamily), 34
hurdle_poisson (brmsfamily), 34
hypothesis, 27, 53, 54, 57, 227, 231, 244
hypothesis(hypothesis.brmsfit), 115
hypothesis.brmsfit, 115

Index (draws-index-brms), 93
index (addition-terms), 8
inhaler, 118
init_refmodel, 107
inits (inits.brmsfit), 119
inits.brmsfit, 119
inv_logit_scaled, 121
InvGaussian, 120
is.brmsfit, 121
is.brmsfit_multiple, 122
is.brmsformula, 122
is.brmsprior, 122
is.brmsterms, 123
is.cor_arma (is.cor_brms), 123
is.cor_brms, 123
is.cor_car (is.cor_brms), 123
is.cor_cosy (is.cor_brms), 123
is.cor_fixed (is.cor_brms), 123
is.cor_sar (is.cor_brms), 123
is.mvbrmsformula, 124
is.mvbrmsterms, 124
iterations, (draws-index-brms), 93

kfold, 128, 129, 207
kfold (kfold.brmsfit), 125

kfold_helpers, 127
 kfold.brmsfit, 107, 125
 kfold_predict, 128
 kidney, 130

 lasso, 131
 launch_shinystan, 132, 150
 launch_shinystan
 (launch_shinystan.brmsfit), 132
 launch_shinystan.brmsfit, 132
 lf(brmsformula-helpers), 51
 log_lik, 82, 85, 125, 143, 144, 200
 log_lik(log_lik.brmsfit), 134
 log_lik.brmsfit, 134, 206
 log_posterior(diagnostic-quantities),
 90
 log_prob, 13
 logistic_normal(brmsfamily), 34
 LogisticNormal, 133
 logit_scaled, 133
 logLik.brmsfit(log_lik.brmsfit), 134
 logm1, 134
 lognormal(brmsfamily), 34
 L0O(loo.brmsfit), 136
 loo, 7, 63, 64, 127, 135–137, 141, 206, 207
 loo(loo.brmsfit), 136
 L0O.brmsfit(loo.brmsfit), 136
 loo.brmsfit, 136
 loo::kfold_split_grouped, 126
 loo::kfold_split_stratified, 126
 loo::loo_model_weights, 139
 loo::psis, 200
 loo_compare, 63, 64, 126, 136–138, 146, 248
 loo_compare(loo_compare.brmsfit), 138
 loo_compare.brmsfit, 138
 loo_epred(loo_predict.brmsfit), 142
 loo_lnpred(loo_predict.brmsfit), 142
 loo_model_weights, 160, 169, 184
 loo_model_weights
 (loo_model_weights.brmsfit),
 139
 loo_model_weights.brmsfit, 139
 loo_moment_match, 137, 140
 loo_moment_match
 (loo_moment_match.brmsfit), 140
 loo_moment_match.brmsfit, 137, 140
 loo_predict(loo_predict.brmsfit), 142
 loo_predict.brmsfit, 142

 loo_predictive_interval
 (loo_predict.brmsfit), 142
 loo_R2(loo_R2.brmsfit), 144
 loo_R2.brmsfit, 144
 loo_subsample, 135, 146, 196
 loo_subsample(loo_subsample.brmsfit),
 145
 loo_subsample.brmsfit, 145
 loss, 146

 ma, 14, 15, 19, 79, 147
 make_conditions, 66, 148, 212
 make_stancode(stancode), 224
 make_standata(standata), 228
 marginal_effects
 (conditional_effects.brmsfit),
 64
 marginal_smooths
 (conditional_smooths.brmsfit),
 69
 MCMC, 167
 mcmc_combo, 168
 mcmc_pairs, 165
 mcmc_plot(mcmc_plot.brmsfit), 149
 mcmc_plot.brmsfit, 149
 me, 52, 151, 219
 mgcv::gamm, 213
 mgcv::s, 213
 mgcv::t2, 213
 mi, 44, 46, 151, 152
 mirai::daemons, 32, 59, 127, 207
 mixture, 48, 153
 mm, 43, 155, 157
 mmc, 156, 157
 mo, 158
 model_weights, 169, 184, 185
 model_weights(model_weights.brmsfit),
 159
 model_weights.brmsfit, 159
 multinomial(brmsfamily), 34
 MultiNormal, 160
 MultiStudentT, 161
 mvbf, 48
 mvbf(mvbrmsformula), 162
 mvbind, 162
 mvbrmsformula, 26, 49, 52, 55, 56, 88, 162,
 162, 226, 231, 244

 nchains(draws-index-brms), 93

ndraws (draws-index-brms), 93
neff_ratio (diagnostic-quantities), 90
negbinomial (brmsfamily), 34
ngrps (ngrps.brmsfit), 163
ngrps.brmsfit, 163
niterations (draws-index-brms), 93
nlf (brmsformula-helpers), 51
nsamples (nsamples.brmsfit), 164
nsamples.brmsfit, 164
nuts_params (diagnostic-quantities), 90
nvariables (draws-index-brms), 93

opencl, 29, 40, 164
options, 28–30, 58, 228, 232, 250

pairs, 165
pairs.brmsfit, 165
pareto-k-diagnostic, 199, 200
pareto_k_ids, 137, 140, 206
parnames, 166
parse_bf (brmsterms), 54
pasym_laplace (AsymLaplace), 18
pbeta_binomial (BetaBinomial), 23
pexgaussian (ExGaussian), 97
pfrechet (Frechet), 103
pgen_extreme_value (GenExtremeValue),
 104
phurdle_gamma (Hurdle), 114
phurdle_lognormal (Hurdle), 114
phurdle_negbinomial (Hurdle), 114
phurdle_poisson (Hurdle), 114
pinv_gaussian (InvGaussian), 120
plan, 29
plot.brms_conditional_effects
 (conditional_effects.brmsfit),
 64
plot.brmsfit, 166
plot.brmhypothesis (brmhypothesis), 53
post_prob, 21, 24, 160, 169, 184
post_prob (post_prob.brmsfit), 182
post_prob.brmsfit, 182
posterior_average, 185
posterior_average
 (posterior_average.brmsfit),
 168
posterior_average.brmsfit, 168
posterior_epred, 22, 67, 85, 143, 144
posterior_epred
 (posterior_epred.brmsfit), 170

posterior_epred.brmsfit, 95, 100–102,
 170, 174
posterior_interval
 (posterior_interval.brmsfit),
 172
posterior_interval.brmsfit, 172
posterior_linpred, 143
posterior_linpred
 (posterior_linpred.brmsfit),
 173
posterior_linpred.brmsfit, 95, 173
posterior_predict, 67, 85, 143, 193
posterior_predict
 (posterior_predict.brmsfit),
 174
posterior_predict.brmsfit, 170, 174, 189,
 191
posterior_samples
 (posterior_samples.brmsfit),
 177
posterior_samples.brmsfit, 177
posterior_smooths
 (posterior_smooths.brmsfit),
 178
posterior_smooths.brmsfit, 178
posterior_summary, 62, 103, 180, 203
posterior_table, 181
powerscale, 82
pp_average, 169
pp_average (pp_average.brmsfit), 183
pp_average.brmsfit, 183
pp_check, 7
pp_check (pp_check.brmsfit), 185
pp_check.brmsfit, 185
pp_expect (posterior_epred.brmsfit), 170
pp_mixture (pp_mixture.brmsfit), 187
pp_mixture.brmsfit, 187
PPC, 186
predict.brmsfit, 186, 189
predict.refmodel, 107
predictive_error
 (predictive_error.brmsfit), 192
predictive_error.brmsfit, 192, 208, 210
predictive_interval
 (predictive_interval.brmsfit),
 193
predictive_interval.brmsfit, 193
prepare_predictions, 101, 105, 127, 129,

135, 136, 139, 145, 160, 169, 171,
 174, 176, 182, 184, 188, 191, 193,
 206, 210, 248
 prepare_predictions
 (prepare_predictions.brmsfit),
 194
 prepare_predictions.brmsfit, 194
 print.brmsfit, 196
 print.brmshypothesis (brmshypothesis),
 53
 print.brmsprior, 197
 print.brmssummary (print.brmsfit), 196
 print.default, 54
 prior (set_prior), 216
 prior_(set_prior), 216
 prior_draws (prior_draws.brmsfit), 197
 prior_draws.brmsfit, 197
 prior_samples (prior_draws.brmsfit), 197
 prior_string (set_prior), 216
 prior_summary (prior_summary.brmsfit),
 198
 prior_summary.brmsfit, 198
 proj_linpred, 107
 proj_predict, 107
 pshifted_lnorm (Shifted_Lognormal), 222
 psis, 143
 psis (psis.brmsfit), 199
 psis.brmsfit, 199
 pskew_normal (SkewNormal), 223
 pstudent_t (StudentT), 234
 pvon_mises (VonMises), 247
 pzero_inflated_beta (ZeroInflated), 251
 pzero_inflated_beta_binomial
 (ZeroInflated), 251
 pzero_inflated_binomial (ZeroInflated),
 251
 pzero_inflated_negbinomial
 (ZeroInflated), 251
 pzero_inflated_poisson (ZeroInflated),
 251

 qasym_laplace (AsymLaplace), 18
 qfrechet (Frechet), 103
 qgen_extreme_value (GenExtremeValue),
 104
 qshifted_lnorm (Shifted_Lognormal), 222
 qskew_normal (SkewNormal), 223
 qstudent_t (StudentT), 234
 quantile, 180

 R2D2, 131, 201, 220
 ranef (ranef.brmsfit), 202
 ranef.brmsfit, 62, 116, 202
 rasym_laplace (AsymLaplace), 18
 rate (addition-terms), 8
 rbeta_binomial (BetaBinomial), 23
 rdirichlet (Dirichlet), 91
 read_csv_as_stanfit, 203
 recompile_model, 205
 recover_data.brmsfit
 (emmeans-brms-helpers), 94
 reloo, 127, 137
 reloo (reloo.brmsfit), 205
 reloo.brmsfit, 205
 rename_pars, 207
 residuals.brmsfit, 208
 resp_bhaz (addition-terms), 8
 resp_cat (addition-terms), 8
 resp_cens (addition-terms), 8
 resp_dec (addition-terms), 8
 resp_index (addition-terms), 8
 resp_mi, 152
 resp_mi (addition-terms), 8
 resp_rate (addition-terms), 8
 resp_se (addition-terms), 8
 resp_subset (addition-terms), 8
 resp_thres (addition-terms), 8
 resp_trials (addition-terms), 8
 resp_trunc (addition-terms), 8
 resp_vint (addition-terms), 8
 resp_vreal (addition-terms), 8
 resp_weights (addition-terms), 8
 restructure, 210
 restructure.brmsfit, 211, 211
 rexgaussian (ExGaussian), 97
 rfrechet (Frechet), 103
 rgen_extreme_value (GenExtremeValue),
 104
 rhat (diagnostic-quantities), 90
 rinv_gaussian (InvGaussian), 120
 rlogistic_normal (LogisticNormal), 133
 rmulti_normal (MultiNormal), 160
 rmulti_student_t (MultiStudentT), 161
 rows2labels, 149, 212
 rshifted_lnorm (Shifted_Lognormal), 222
 rskew_normal (SkewNormal), 223
 rstan::stan_model, 30
 rstudent_t (StudentT), 234

runApp, 132
rvon_mises (VonMises), 247
rwiener (Wiener), 249

s, 43, 212
sampling, 30
sar, 19, 80, 213
save_pars, 28, 215
saveRDS, 11, 30, 58
scale_colour_gradient, 68
scale_colour_grey, 68
se (addition-terms), 8
set.seed, 107, 116, 169, 184
set_mecor (brmsformula-helpers), 51
set_nl (brmsformula-helpers), 51
set_prior, 27, 31, 46, 47, 57, 72, 87, 89, 114, 131, 154, 202, 216, 227, 231, 244, 245
set_rescor (brmsformula-helpers), 51
Shifted_Lognormal, 222
shifted_lognormal (brmsfamily), 34
skew_normal (brmsfamily), 34
SkewNormal, 223
sratio (brmsfamily), 34
Stan, 7
stan, 29, 31, 73
stancode, 7, 224, 226, 233
stancode.brmsfit, 225, 225
stancode.default, 225, 226
standata, 7, 228
standata.brmsfit, 229, 229
standata.default, 229, 230, 230
stanfit, 41
stanmodel, 13
stanplot, 7
stanplot (mcmc_plot.brmsfit), 149
stanvar, 27, 57, 84, 85, 227, 231, 232
stanvars, 40, 196, 230, 243
stanvars (stanvar), 232
student (brmsfamily), 34
StudentT, 234
subset (addition-terms), 8
subset_draws, 17, 92, 93
summarize_draws, 180
summary, 7
summary.brmsfit, 197, 235

t2, 43
t2(s), 212

TDist, 235
theme, 54, 68, 167
theme_black, 236
theme_default, 54, 68, 167, 237, 237
threading, 29, 40, 225, 228, 232, 233, 237
thres (addition-terms), 8
trials (addition-terms), 8
trunc (addition-terms), 8

unstr, 19, 238
update, 28, 58
update.brmsfit, 206, 239, 241
update.brmsfit_multiple, 240
update.formula, 239, 241
update_adterms, 241

validate_newdata, 196, 230, 242
validate_prior, 243
VarCorr, 245
VarCorr (VarCorr.brmsfit), 245
VarCorr.brmsfit, 245
variables, 116
variables (draws-index-brms), 93
variables, (draws-index-brms), 93
variables.brmsfit (draws-index-brms), 93
varsel, 106
vb, 30
vcov.brmsfit, 246
Vectorize, 98
vint (addition-terms), 8
von_mises (brmsfamily), 34
VonMises, 247
vreal (addition-terms), 8

WAIC (waic.brmsfit), 248
waic, 7, 63, 64, 135, 248
waic (waic.brmsfit), 248
WAIC.brmsfit (waic.brmsfit), 248
waic.brmsfit, 248
weibull (brmsfamily), 34
weights (addition-terms), 8
weights(), 200
Wiener, 249
wiener (brmsfamily), 34
wienerdist, 250

xbeta (brmsfamily), 34

zero_inflated_beta (brmsfamily), 34

zero_inflated_beta_binomial
 (brmsfamily), [34](#)
zero_inflated_binomial (brmsfamily), [34](#)
zero_inflated_negbinomial (brmsfamily),
 [34](#)
zero_inflated_poisson (brmsfamily), [34](#)
zero_one_inflated_beta (brmsfamily), [34](#)
ZeroInflated, [251](#)